

1

# Deformation-Driven Element 2 Packings

3 by

4 Reza Adhitya Saputra

5 A thesis  
6 presented to the University of Waterloo  
7 in fulfillment of the  
8 thesis requirement for the degree of  
9 Doctor of Philosophy  
10 in  
11 Computer Science

12 Waterloo, Ontario, Canada, 2020

13 © Reza Adhitya Saputra 2020  
14 Last edited: Wednesday 29<sup>th</sup> July, 2020 06:26

## Examining Committee Membership

<sup>16</sup> The following served on the Examining Committee for this thesis. The decision of the  
<sup>17</sup> Examining Committee is by majority vote.

External Examiner: YYY YYY  
Professor, Dept. of YYY, University of YYY

 Supervisor(s): Craig S. Kaplan  
Professor, Dept. of Computer Science, University of Waterloo

Internal Member: Christopher Batty  
Professor, Dept. of Computer Science, University of Waterloo

Internal-External Member: YYY YYY  
Professor, Dept. of YYY, University of Waterloo

Other Member(s): Daniel Vogel  
Professor, Dept. of Computer Science, University of Waterloo

## Author's Declaration

19 This thesis consists of materials which I authored or co-authored: see Statement of  
20 Contributions. This is a true copy of the thesis, including any required final revisions, as  
21 accepted by my examiners

22 I understand that my thesis may be made electronically available to the public.

23

## Statement of Contribution

24 This thesis was built from previous publications authored by myself, Dr. Craig S.  
25 Kaplan, and Dr. Paul Asente.

26 FLOWPAK in Chapter 3 was initially conducted with Dr. Paul Asente, and Dr.  
27 Radomír Měch at Adobe Research in San Jose, CA. FLOWPAK is also discussed in  
28 greater detail in our publication “FLOWPAK: Flow-Based Ornamental Element Pack-  
29 ing” [SKAM17] and in our patent “Computerized Generation of Ornamental Designs by  
30 Placing Instances of Simple Shapes in Accordance With a Direction Guide” [AKMS20].

31 RepulsionPak in Chapter 4 and Quantitative Metrics in Chapter 6 were taken from  
32 “RepulsionPak: Deformation-Driven Element Packing with Repulsion Forces” [SKA18]  
33 and “Improved Deformation-Driven Element Packing with RepulsionPak” [SKA19].

34 AnimationPak in Chapter 5 was taken from “AnimationPak: Packing Elements with  
35 Scripted Animations” [SKA20].

## Abstract

We present three methods to fill a container shape with deformable instances of geometric elements, creating an artistic composition called a packing. The first method is FLOWPAK that deforms elements to flow along a vector field interpolated from user-supplied strokes, giving a sense of visual flow to the final composition. The second method is RepulsionPak that utilizes repulsion forces to pack elements, each represented as a mass-spring system, allowing them to deform to achieve a better fit with their neighbors and the container. The last method is AnimationPak that creates animated packings by arranging animated two-dimensional elements inside a static container. We represent animated elements in a three-dimensional spacetime domain, and view the animated packing problem as a three-dimensional packing in that domain. Finally, we propose statistical methods for measuring the evenness of 2D element distributions, which provide quantitative means of evaluating and comparing packing algorithms.

## Acknowledgements

I would like to thank all people who have helped me during my PhD study:

- Craig S. Kaplan and Paul Asente, for being the best teachers.
  - Radomír Měch, for the internship and the help with FLOWPAK paper.
  - Daichi Ito, for the discussion about design principles and using RepulsionPak to create a packing.
  - Dietrich Stoyan, for the discussion about spatial statistics.
  - Danny Kaufman, for the discussion about physical simulations.
  - Steve Mann, for giving me valuable lessons.
  - Bill Cowan, for his advice.
  - Christopher Batty and Dan Vogel, for being my PhD committee members.
  - Matt Thorne, Alex Pytel, and Tyler Nowicki, for all the lunches together.
  - All my running friends in Kitchener-Waterloo and ultra running community in Ontario, who teach me on how to keep grinding.
  - Russ Jones, for permission to use the fish art in Figure 3.2d.
  - Yusuf Umar, for designing ornamental elements and the box bird in Figure 3.17.

Dedication

65

For my Mom, Dad, and sister.

# <sup>67</sup> Table of Contents

<sup>68</sup> <b>Abstract</b>	<sup>v</sup>
<sup>69</sup> <b>List of Tables</b>	<sup>xi</sup>
<sup>70</sup> <b>List of Figures</b>	<sup>xii</sup>
<sup>71</sup> <b>1 Introduction</b>	<sup>1</sup>
<sup>72</sup> <b>2 Related Work</b>	<sup>5</sup>
<sup>73</sup> 2.1 2D Packings . . . . .	<sup>5</sup>
<sup>74</sup> 2.2 Tilings . . . . .	<sup>9</sup>
<sup>75</sup> 2.3 Discrete Texture Synthesis . . . . .	<sup>9</sup>
<sup>76</sup> 2.4 3D Packings . . . . .	<sup>10</sup>
<sup>77</sup> 2.5 Packings on Surfaces . . . . .	<sup>11</sup>
<sup>78</sup> <b>3 FLOWPAK: Flow-Based Ornamental Element Packing</b>	<sup>13</sup>
<sup>79</sup> 3.1 Introduction . . . . .	<sup>14</sup>
<sup>80</sup> 3.2 Related Work . . . . .	<sup>16</sup>
<sup>81</sup> 3.3 Problem Formulation . . . . .	<sup>17</sup>
<sup>82</sup> 3.4 Approach . . . . .	<sup>18</sup>
<sup>83</sup> 3.4.1 Target Containers . . . . .	<sup>18</sup>
<sup>84</sup> 3.4.2 Ornamental Elements and LR functions . . . . .	<sup>19</sup>

85	3.4.3	Creating Vector Fields and Tracing Streamlines . . . . .	21
86	3.4.4	Sub-Region Blobs . . . . .	23
87	3.4.5	Shape Matching and Deformation . . . . .	23
88	3.4.6	Iterative Refinement . . . . .	24
89	3.5	Implementation and Results . . . . .	28
90	3.6	Conclusions . . . . .	30
91	<b>4</b>	<b>RepulsionPak: Deformation-Driven Element Packing with Repulsion Forces</b>	<b>34</b>
93	4.1	Introduction . . . . .	35
94	4.2	Related Work . . . . .	35
95	4.3	System Overview . . . . .	37
96	4.4	Preprocessing . . . . .	39
97	4.5	Simulation . . . . .	41
98	4.5.1	Element Growth . . . . .	45
99	4.5.2	Stopping Criteria . . . . .	46
100	4.6	Secondary Elements . . . . .	46
101	4.7	Shape Matching . . . . .	47
102	4.8	Implementation and Results . . . . .	50
103	4.9	Conclusions . . . . .	53
104	<b>5</b>	<b>AnimationPak: Packing Elements with Scripted Animations</b>	<b>55</b>
105	5.1	Introduction . . . . .	56
106	5.2	Related Work . . . . .	56
107	5.3	Animated Elements . . . . .	58
108	5.3.1	Spacetime Extrusion . . . . .	59
109	5.3.2	Animation . . . . .	60
110	5.4	Initial Configuration . . . . .	61

111	5.5 Simulation . . . . .	62
112	5.5.1 Spatial Queries . . . . .	66
113	5.5.2 Slice Constraints . . . . .	67
114	5.5.3 Element Growth and Stopping Criteria . . . . .	69
115	5.6 Rendering . . . . .	70
116	5.7 Implementation and Results . . . . .	71
117	5.8 Conclusions . . . . .	75
118	<b>6 Quantitative Metrics for 2D Packings</b>	<b>79</b>
119	6.1 Introduction . . . . .	79
120	6.2 Related Work . . . . .	79
121	6.3 Quantitative Metrics . . . . .	80
122	6.3.1 Spherical Contact Probability . . . . .	81
123	6.3.2 Histograms of the Distance Transform . . . . .	81
124	6.3.3 The Overlap Function . . . . .	84
125	6.4 Comparisons . . . . .	85
126	6.5 Conclusions . . . . .	89
127	<b>7 Conclusions and Future Work</b>	<b>90</b>
128	7.1 Conclusions . . . . .	90
129	7.2 Future Work . . . . .	90
130	7.2.1 FLOWPAK . . . . .	91
131	7.2.2 RepulsionPak . . . . .	92
132	7.2.3 AnimationPak . . . . .	93
133	7.2.4 Packing Evaluation . . . . .	94
134	<b>References</b>	<b>96</b>

# <sup>135</sup> List of Tables

<sup>136</sup>	4.1 Data and statistics for the RepulsionPak results . . . . .	51
<sup>137</sup>	5.1 Data and statistics for the AnimationPak results . . . . .	72

# <sup>138</sup> List of Figures

<sup>139</sup>	1.1 Unilever packing and its negative space . . . . .	2
<sup>140</sup>	1.2 Packings in graphic design . . . . .	2
<sup>141</sup>	1.3 A packing in land art . . . . .	3
<sup>142</sup>	2.1 An illustration of Lloyd’s method. . . . .	6
<sup>143</sup>	2.2 Decorative mosaics using Lloyd’s method . . . . .	6
<sup>144</sup>	2.3 Packings generated by JIM and PAD . . . . .	8
<sup>145</sup>	2.4 A calligraphy packing of an “elephant” . . . . .	8
<sup>146</sup>	2.5 A tiling of cats . . . . .	10
<sup>147</sup>	2.6 A 3D packing and a packing on a surface . . . . .	12
<sup>148</sup>	3.1 Packings of lion and unicorn . . . . .	13
<sup>149</sup>	3.2 Examples of flow visual style . . . . .	15
<sup>150</sup>	3.3 FLOWPAK pipeline . . . . .	19
<sup>151</sup>	3.4 Ornamental elements and their LR functions . . . . .	20
<sup>152</sup>	3.5 The streamline tracing process . . . . .	22
<sup>153</sup>	3.6 Shape deformation . . . . .	24
<sup>154</sup>	3.7 Shifting a streamline . . . . .	25
<sup>155</sup>	3.8 Tracing a shortest path . . . . .	27
<sup>156</sup>	3.9 Stretching an element . . . . .	29
<sup>157</sup>	3.10 Rotating an element . . . . .	29

158	3.11 Reflecting an element . . . . .	29
159	3.12 A packing of a rhinoceros . . . . .	31
160	3.13 A packing of a bear with leaf elements . . . . .	31
161	3.14 A packing of a cat . . . . .	32
162	3.15 A packing of a dog . . . . .	32
163	3.16 A packing of a bear with elements created from negative space . . . . .	33
164	3.17 A packing of a bird . . . . .	33
165	4.1 A packing of a cat . . . . .	34
166	4.2 Primary and secondary elements . . . . .	36
167	4.3 RepulsionPak pipeline . . . . .	38
168	4.4 Element deformation . . . . .	39
169	4.5 Element discretization . . . . .	39
170	4.6 Illustrations of the forces in a RepulsionPak simulation . . . . .	44
171	4.7 A local shape descriptor for shape matching . . . . .	48
172	4.8 A demonstration of shape matching to place elements . . . . .	49
173	4.9 Three packings created using RepulsionPak: Birds, Bats, and Butterflies .	52
174	4.10 A packing that demonstrates torsional forces . . . . .	53
175	4.11 A paisley-inspired toroidal packing . . . . .	54
176	4.12 Two examples of how extreme parameter values can lead to low-quality results . . . . .	54
178	5.1 An animated packing of aquatic fauna . . . . .	55
179	5.2 An animated packing of copies of Lisa Simpson . . . . .	57
180	5.3 The creation of a discretized spacetime element . . . . .	60
181	5.4 A spacetime element with a scripted animation . . . . .	61
182	5.5 A 2D illustration of a guided element . . . . .	63
183	5.6 An illustration of the AnimationPak simulation process . . . . .	64
184	5.7 An illustration of the repulsion forces . . . . .	65

185	5.8 An illustration of the temporal forces . . . . .	66
186	5.9 An envelope and a collision grid . . . . .	67
187	5.10 Constraints . . . . .	68
188	5.11 Element growths . . . . .	70
189	5.12 An illustration of rendering spacetime elements . . . . .	71
190	5.13 An animated packing of a snake chases a bird . . . . .	73
191	5.14 An animated packing of the penguins-to-giraffes illusion . . . . .	74
192	5.15 A traveling bird passing through a packing of birds . . . . .	75
193	5.16 A packing of lion's mane . . . . .	76
194	5.17 A comparison between RepulsionPak and AnimationPak . . . . .	76
195	5.18 A comparison between CAVD, The spectral approach, and AnimationPak .	77
196	5.19 The effect of adjusting time spring stiffness . . . . .	77
197	5.20 A failure case for AnimationPak . . . . .	78
198	6.1 Spherical contact probabilities and distance histograms for reference packings . . . . .	82
199	6.2 An example of distance transform of negative space . . . . .	83
200	6.3 An example of a medial axis of negative space . . . . .	83
201	6.4 An illustration of offsetting elements outward . . . . .	84
202	6.5 A comparison between a PAD packing and a RepulsionPak packing . . . . .	86
203	6.6 A comparison between the artist-made packing and a RepulsionPak packing . . . . .	87
204	6.7 A demonstration of the effect of deformation on the evenness of negative space . . . . .	88
205	7.1 Element threading and branching . . . . .	92

<sup>209</sup> **Chapter 1**

<sup>210</sup> **Introduction**

<sup>211</sup> A *packing* is an arrangement of geometric *elements* within a *container* region in the plane.  
<sup>212</sup> As an artistic composition, a packing can communicate a relationship between a whole and  
<sup>213</sup> the parts that make it up. Elements work together to communicate the overall container  
<sup>214</sup> shape, but each is large enough to be appreciated individually. Elements are shapes like  
<sup>215</sup> animals, plants, ~~abstract, geometric~~, or man-made objects. Elements can also be inter-  
<sup>216</sup> preted as *positive space*. Packings are popular in logo design, graphic design, and art.  
<sup>217</sup> Figure 1.1a shows the Unilever logo, which is a 2D packing of elements arranged inside a  
<sup>218</sup> “U” container.

 <sup>219</sup> The subset of the container that does not belong to any element is called *negative*  
<sup>220</sup> *space* (Figure 1.1b). We can also interpret negative space as separation and gaps between  
<sup>221</sup> elements. The evenness of negative space plays an important role in packings. The artist  
<sup>222</sup> should arrange elements in a way that their boundaries interlock with each other, causing  
<sup>223</sup> the separation between neighboring elements to become roughly the same everywhere.  
 <sup>224</sup> After all, element interlocking is imperfect. It is not uncommon that the artist artfully  
<sup>225</sup> places small elements such as triangles or circles to reduce remaining large gaps, see two  
<sup>226</sup> examples in Figure 1.2.

<sup>227</sup> Elements can also be oriented to follow certain directions for aesthetic reasons. Fig-  
<sup>228</sup> ure 1.2b is an ornamental packing of long thin elements  that are bent and oriented outward  
<sup>229</sup> from the center of the torso to create a flow visual style. Another example is shown in  
<sup>230</sup> Figure 1.3, which is a packing of leaves, each is oriented toward the center of the container.  
<sup>231</sup> On both examples, the flow visual style gives an impression of progression and movements.

<sup>232</sup> There has been a moderate amount of past research in computer graphics, particu-  
<sup>233</sup> larly in the field of non-photorealistic rendering, on the generation of packings, or mosaics.

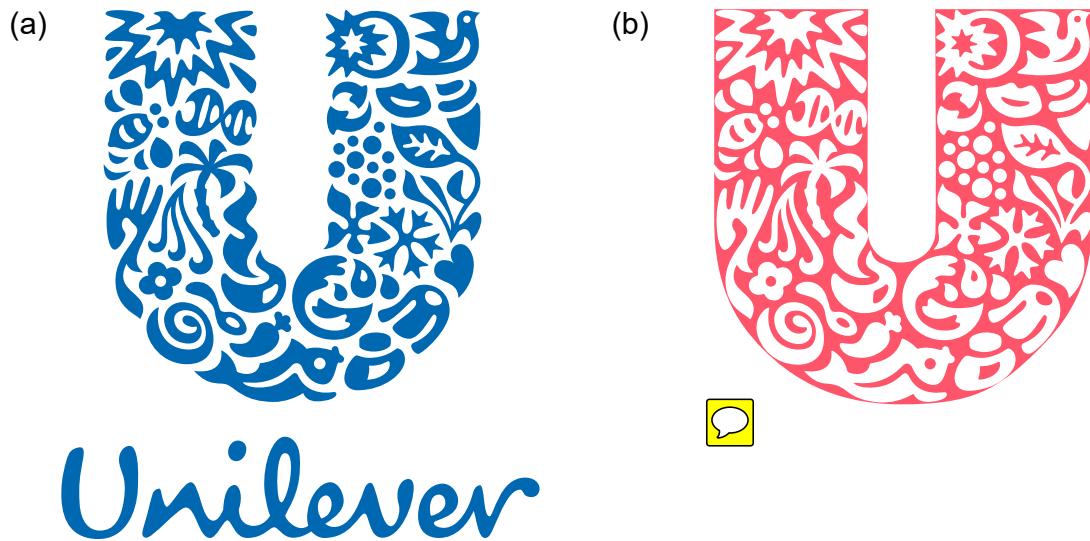


Figure 1.1: (a) The Unilever logo packing that consists of 25 elements arranged inside a “U” container. (b) The red region is called negative space, which is the remaining area that is not claimed by the elements.

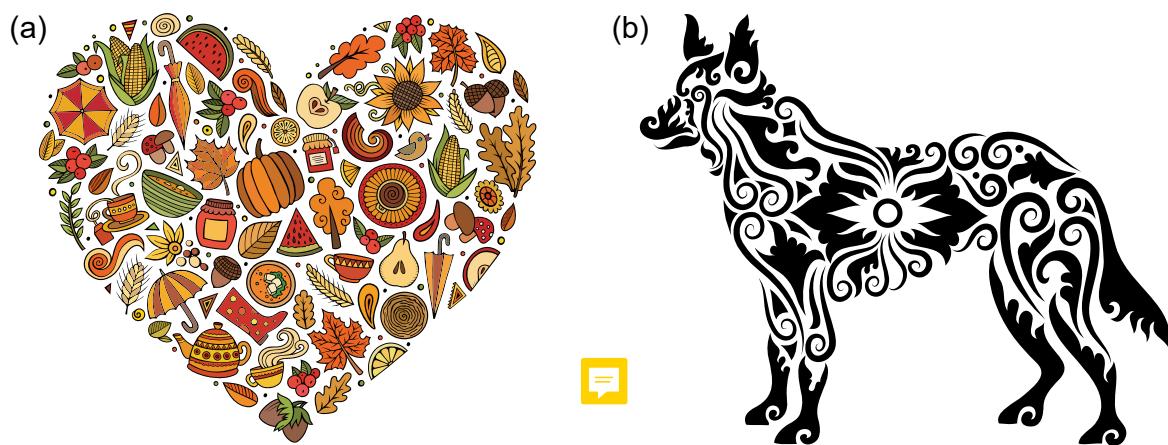


Figure 1.2: Two packing examples in graphic design. (a) A packing of autumn-themed elements, created by Balabolka. (b) A packing of abstract-shaped elements that create a flow visual style, created by ComicVector703.



Figure 1.3: A land art packing that is composed of leaves, created by James Brunt. Each leaf is oriented toward the center of the circle.

234 See Chapter 2 for specific examples. However, most techniques pack elements via rigid  
235 transformations, leading to uneven element distribution and overlaps. Jigsaw Image Mo-  
236 saics [KP02] and collages based on the Pyramid of Arclength Descriptor [KSH<sup>+</sup>16] are  
237 *data-driven*. These techniques rely on a large library of elements, so that given an area to  
238 fill in a partial composition, there is likely to be an element in the library with a compati-  
239 ble shape. The challenge is *finding compatible elements*, which requires designing a shape  
240 matching technique that is fast and robust. Additionally, they cannot guarantee to find a  
241 compatible element at every iteration, and elements typically do not fit perfectly with each  
242 other or the container boundary. The remedy here is providing more data with increased  
243 computation time. However, a large library may not be feasible or artistically desirable.  
244 If an artist wants a packing of hand-drawn cats, and a data-driven approach cannot find  
245 a good result with ten cat shapes, the artist may not want to draw 100 or 1000 cats to  
246 ensure a better fit.

247 In this thesis, we propose *deformation-driven* methods. Instead of finding compatible  
248 elements, we intent to *create compatible elements* through deformation that can adapt  
249 to the shapes of neighboring elements and the container boundary. We allow elements to  
250 deform in a controlled way, to trade off between the evenness of the element distribution and  
251 the deformations of the individual elements. By building an algorithm with a controllable  
252 deformation model at its core, we achieve a more even distribution of negative space, and

253 we only require a small library of element shapes.

254 Deformation-driven methods also allow us to work toward a design principle called  
255 *uniformity amidst variety* [Hut29, Gom84]. *Uniformity* aims for an overall unity of design  
256 and *variety* seeks to break up the monotony of pure repetition. We can achieve a degree  
257 of uniformity by using repeated copies of a small library of elements, but balance that  
258 uniformity with variety by deforming those elements. We believe that there is a value in  
259 deformation that can generate plausible families of related elements from a single input  
260 shape.

261 The perceived packing quality is closely tied to the evenness of negative space. The more  
262 the elements interlock, the more even the negative space. We are interested in quantitative  
263 measurements of the evenness in order to evaluate and compare packing algorithms. In  
264 this thesis we discuss several possible measurements of evenness based on methods from  
265 spatial statistics.

266 In this thesis, our contribution is developing three deformation driven packing methods  
267 and 2D packing quantitative metrics:

- 268 1. FLOWPAK is a packing method that deforms long thin elements to follow a user-  
269 defined vector field (Chapter 3).
- 270 2. RepulsionPak is a packing method that utilizes repulsion forces to distribute and  
271 deform elements, each is represented as a mass-spring system (Chapter 4).
- 272 3. AnimationPak is a method to pack 2D animated elements inside a static 2D container.  
273 Each element is an extruded 3D shape in a spacetime domain and we view the  
274 animated packing problem as a 3D packing in that domain. (Chapter 5).
- 275 4. Quantitative metrics for measuring the evenness of negative space: spherical con-  
276 tact probabilities, histograms of the distance transforms, and the overlap functions  
277 (Chapter 6).

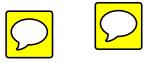


278

# Chapter 2

279

## Related Work



280

### 2.1 2D Packings

281 **Lloyd's Method:** An approach to generate a packing 282 is to start with an initial configuration and iteratively refine it using Lloyd's method 283 to obtain an even distribution 284 of elements. Figure 2.1 shows an illustration of point-based Lloyd's method. Given  $N$  285 point elements, we first compute a *Voronoi diagram* that partitions the plane into  $N$  regions such that all points inside a Voronoi cell are closest to its associated point element. 286 Lloyd's method moves every point element to the centroid of its Voronoi cell, then the 287 Voronoi diagram is recomputed. The process is repeated until the distribution is even, 288 that means all the point elements are located at the centroids of the Voronoi cells. The 289 final structure is called a *centroidal Voronoi diagram* (CVD).

290 Secord [Sec02] computed a weighted CVD so that the resulting point element distribution 291 resembles a stippling artwork. His method incorporates the gradient of an input 292 image into the centroid calculation, attracting point elements to low-intensity regions.

293 A standard CVD is generated using Euclidean distance metric, but it can be replaced to 294 manipulate the shapes of Voronoi cells. Hausner [Hau01] used Manhattan distance metric 295 so that Voronoi cells resemble squares instead of hexagons. Each Voronoi cell is then 296 replaced with a square element and the resulting arrangement resembles the appearance 297 of traditional mosaics (Figure 2.2). However, Hausner's approach can only position the 298 square elements, and it must incorporate a vector field to modify the orientation of the 299 metric and rotate the square elements. Additionally, the approach is only suitable for 300 square elements as more complicated shapes such as long rectangles would have severe

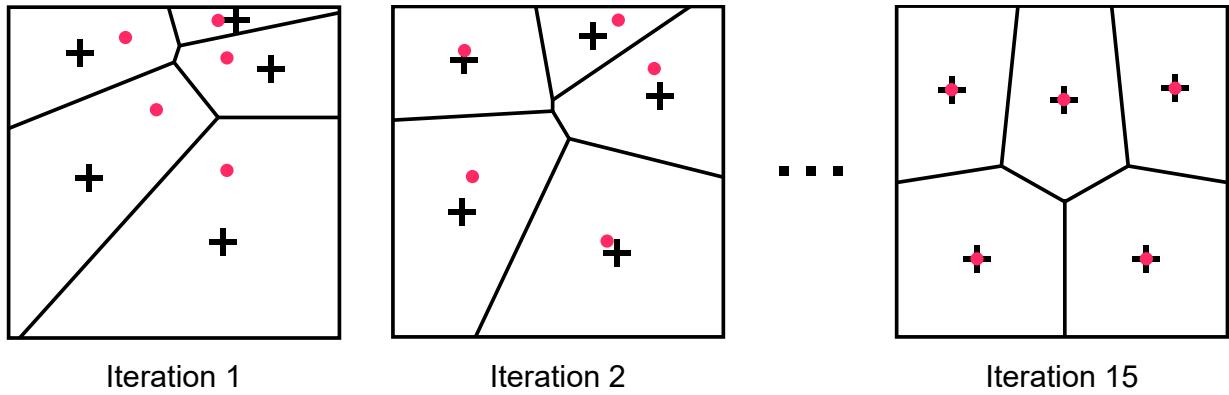


Figure 2.1: An illustration of point-based Lloyd's method. A Voronoi diagram is generated from the input point elements (drawn as red dots). We then move the point elements to the centroids of Voronoi cells (drawn as plus signs). The process is repeated until convergence, producing a centroidal Voronoi diagram. Figure source is Wikipedia, drawn by Dominik Moritz under CC0 1.0.

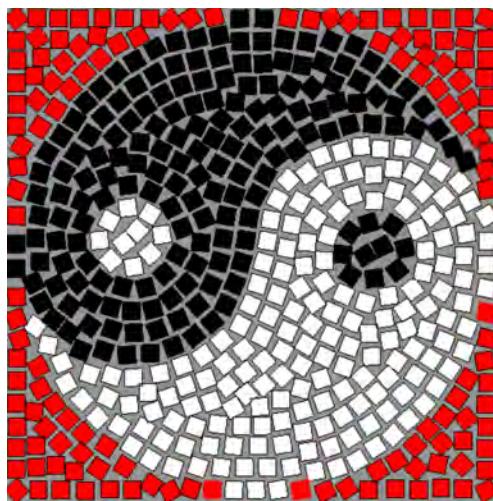


Figure 2.2: Mosaics of square elements generated using Lloyd's method with Manhattan distance [Hau01]

301 overlaps. More recently, Javid et al. [JDM19] constructed a metric that incorporates a  
302 spatial distance, a color distance, and an elongation factor. Their resulting Voronoi cells  
303 have elongated shapes, resembling pebble-like mosaics.

304 Hiller et al. [HHD03] improved upon the Lloyd's method to generate *centroidal area*  
305 *Voronoi diagram* (CAVD), a variant of CVD that is a distribution of polygonal elements.  
306 This new extension computes the main inertia axis for each Voronoi cell so that its element  
307 can be rotated to get better alignment with the Voronoi cell boundary. In follow up  
308 work, Smith et al. [SLK05] generated temporally coherent animated mosaics by utilizing  
309 CAVD. Dalal et al. [DKLS06] used an FFT-based image correlation to reposition and rotate  
310 elements so that they achieve the best alignments with their Voronoi cell boundaries, which  
311 could be seen as making more effective use of negative space, and permitting non-convex  
312 elements to interlock more than they did in earlier methods.

313 **Data-Driven Method:** A different approach to generate packings is to use a shape  
314 matching algorithm (Figure 2.3). This is usually done by placing an element one by  
315 one until the container area is filled. For each step, a shape matching algorithm selects  
316 an element that has the best compatibility with the previously placed elements or the  
317 container boundary. Jigsaw Image Mosaics [KP02] uses a geometric hashing to find a  
318 compatible element. JIM places elements using a greedy approach, but is able to backtrack  
319 if a previous configuration is more optimal. Pyramid of Arclength Descriptor [KSH<sup>+</sup>16] is  
320 a partial shape matching algorithm that can match subsets of element boundaries. 

321 It is also possible to initially partition the container into smaller segments then inde-  
322 pendently replace each segment with a matching element. This is desirable if the container  
323 has colors or salient parts that need to be emphasized. Huang et al. [HZZ11] produced  
324 Arcimboldo-like collages by arranging cutout images collected from the internet. Their  
325 container is a larger cutout image which is partitioned into parts using an image segmen-  
326 tation algorithm. Each part is then replaced with a smaller cutout image that has a similar  
327 shape and color. 

328 All these data-driven methods require an element library that is big enough, the more  
329 elements in the library will increase their chance to find a compatible element bound-  
330 ary. However, bigger database means increased computational time and collecting a large  
331 number of elements is also not always feasible.

332 **Deformation-Driven Method:** Instead of finding compatible elements, a deformation-  
333 driven method allows elements to create compatibilities. As an advantage, a deformation-  
334 driven method can rely on a small size element library. Xu and Kaplan [XK07] and Zou et  
335 al. [ZCR<sup>+</sup>16] constructed *calligrams* by filling a container with a small number of deformed  
336 letters composing one or two words (Figure 2.4). Their goal was to balance between filling  
  




Figure 2.3: Data-driven methods: (a) Jigsaw Image Mosaics [KP02]. (b) Pyramid of Arc length Descriptor [KSH<sup>+</sup>16].

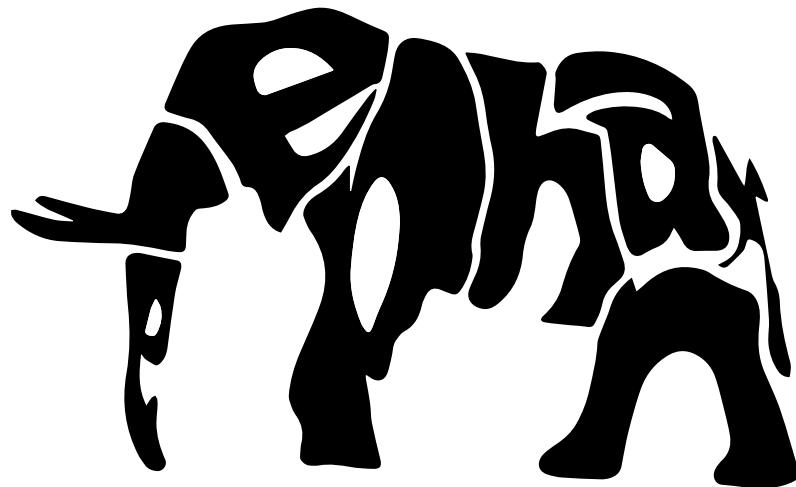


Figure 2.4: A deformation-driven packing of a calligraphic “elephant” [XK07].

337 the container and preserving readability. Peng et al. [PYW14] computed layouts by pack-  
338 ing and deforming simple polygons and polyominoes, although their method cannot handle  
339 more complicated shapes. Abdrashitov et al. [AGYS14] developed a sketch-based interface  
340 where an artist can draw curves where square-like elements are placed along it to create a  
341 mosaic arrangement. After the square elements are frozen in place, they are deformed to  
342 eliminate overlap, to create shape variations, and to even out the negative space.  

## 343 2.2 Tilings

344 As elements reach perfect compatibility, a packing turns to a tiling, see Figure 2.5 for  
345 an example. The element boundaries interlock, leaving no negative space. Kaplan and  
346 Salesin [KS00, KS04] deformed one or two input elements into ones that can tile the plane.  
347 Given an input element  $E$ , they performed a search in a parameterized tiling space to  
348 find its deformed variant  $T$ . Due to high frequency deformation,  $T$  is often perceptually  
349 unrecognizable from its silhouette unless an interior picture is added. Their method can  
350 fail to find  $T$  if  $E$  has deep concavity, for example, a “G” shape. The tile  $T$  also cannot  
351 fill a container due to its highly structured boundary.  

352 Lin et al. [LML<sup>+</sup>18] generated ground reversal compositions that resemble Escher’s Sky   
353 and Water I, where elements serve as both positive and negative space depending on the  
354 viewer’s perception. Given two elements  $S_1$  and  $S_2$ , each is placed on the opposite side  
355 of a rectangular canvas. They are then copied, arranged, and deformed to interlock each  
356 other. The farther a copy from the original element, the more deformed it is, creating a  
357 “fading effect”,  $S_1$  and  $S_2$  spatially fade to nothingness.  This work is data-driven, the user  
358 needs to provide  $S_1$ , and the method tries to find a compatible element  $S_2$  from a library. 

## 359 2.3 Discrete Texture Synthesis

360 Some past work has sought to adapt example-based texture synthesis methods from raster  
361 images to vector graphics, producing distributions of rigidly transformed elements that  
362 mimic the statistics of an *exemplar*. An element is represented as a single point, which is  
363 adequate for small near-convex elements. Barla et al. [BBT<sup>+</sup>06] and Ijiri et al. [IMIM08]  
364 used a growth model that copies small neighborhoods from the exemplar into a larger  
365 output texture. Hurtut et al. [HLT<sup>+</sup>09] developed a statistical sampling method based on  
366 multitype point processes. AlMeraj et al. [AKA13a] stamped out copies of the exemplar  
367 and discarded overlapping elements. Loi et al. [LHVT17] developed a texture synthesis 

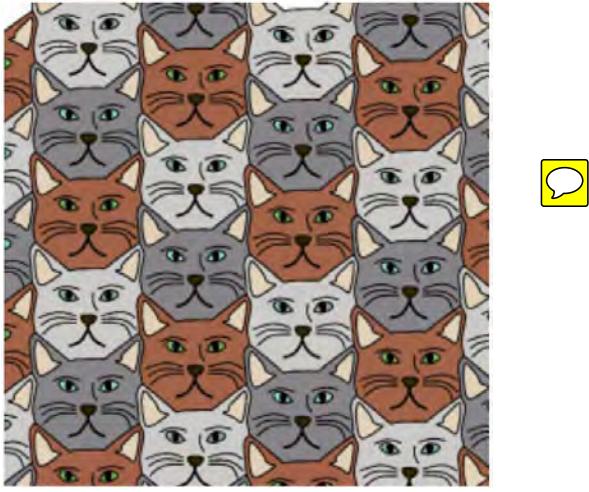


Figure 2.5: A tiling of cats [KS00], elements perfectly interlock each other and the arrangement does not have negative space.

368 method that can specify global arrangements, local arrangements, or a blend of multiple  
 369 arrangements. These techniques are all concerned with replicating the uneven element  
 370 distribution in the exemplar, without regard for negative space. 

371 For larger concave elements, a single point representation is not enough. Ma et al.  
 372 [MWT11] used sample-based representation that is created by generating a sparse set of  
 373 points inside an element. They later distributed these points using a neighborhood metric  
 374 and an iterative optimization. Unlike previous work, they were able to synthesize textures  
 375 with long deformable elements, for example, spaghetti. Ma et al. later extended their  
 376 work to accept animated elements [MWLT13], where each sample point has a spatial and  
 377 time position, turning the problem to spacetime texture synthesis. More recently, Hsu et  
 378 al. [HWYZ20] adapted the sample-based representation into an interactive app, where an  
 379 artist can initially distribute elements by drawing strokes which are then optimized using  
 380 a Lloyd-like optimization.

## 381 2.4 3D Packings

382 Collages are arrangements of overlapping elements, similar to portrait paintings by Giuseppe  
 383 Arcimboldo. Gal et al. [GSP<sup>+</sup>07] presented a method for constructing 3D collages (Figure 2.6a). They filled a 3D container with overlapping 3D elements using a greedy approach  
 384

385 and a partial shape matching algorithm. Theobalt et al. [TRdAS07] developed a method  
386 to generate animated 3D collages. They segmented an animated container into smaller  
387 rigid parts, each is replaced with a matching element. The final result of animated collage  
388 consists of rigid motions of elements. Huang et al. [HWFL14] designed a method to  
389 generate mechanical collages, such as giant robots. All these collage methods require a 3D  
390 shape database so they are considered as data-driven.

391 The cutting and packing problem (C&P) is defined as cutting a large object into smaller  
392 parts which are then packed inside a container. C&P is popular in manufacturing and 3D  
393 printing because objects can be produced with less waste material and packed into a smaller  
394 box. A good cutting process is critical in C&P if it can decompose the input object into  
395 simpler parts, then the packing process can be easier. Chernov et al. [CSR10] proposed a  
396 method to decompose a large object into smaller phi objects which can be packed more  
397 efficiently. A phi object is defined as a shape whose surface boundaries are flat, spherical,  
398 cylindrical, or conical. Vanek et al. [VGB<sup>+</sup>14] introduced PackMerger, a method to pack  
399 thin shells which can be assembled together into a larger watertight object. In follow up  
400 work, Chen et al. [CZL<sup>+</sup>15] introduced Dapper, a method to cut and pack volumetric  
401 printed objects.

402 In engineering, packings are useful for a number of applications, such as product pack-  
403 aging, circuit designs, or mechanical layouts. This requires elements to be packed without  
404 any overlap. Cagan et al. [CSY02] compiled a survey of 3D packing approaches such as  
405 gradient methods, simulated annealing, and genetic algorithms. Byholm et al. [BTW09]  
406 developed a method to pack voxelized elements, which are computationally easier for col-  
407 lision detection. Ma et al. [MCHW18] proposed a heuristic method to pack triangular  
408 meshes.

## 409 2.5 Packings on Surfaces

410 Some recent work has explored the elaboration of ornamental patterns on surfaces, under  
411 constraints imposed by fabrication, such as connectivity. Chen et al. [CZX<sup>+</sup>16] developed  
412 a method to generate a filigree pattern, which is an arrangement of decorative thin rod  
413 elements on a surface. Given an initial random configuration of overlapping elements,  
414 their method removes overlaps by either deforming or trimming the rods. In similar work,  
415 Zehnder et al. [ZCT16] proposed a semi-automated tool for deforming ornamental curves  
416 to cover a surface (Figure 2.6b). They started with an initial configuration of scaled down  
417 elements that has no overlaps. They later grew the elements and avoided overlaps using  
418 curve deformation. Bian et al. [BWL18] used Wang tiles made of element parts to generate

419 filigree patterns. Martínez et al. [MSS<sup>+</sup>19] developed a CVD-based method to generate  
420 star-shaped tiling patterns that are printed onto tracery sheets. Their method works by  
421 constructing a star-shaped metric to manipulate the Voronoi cell shapes.

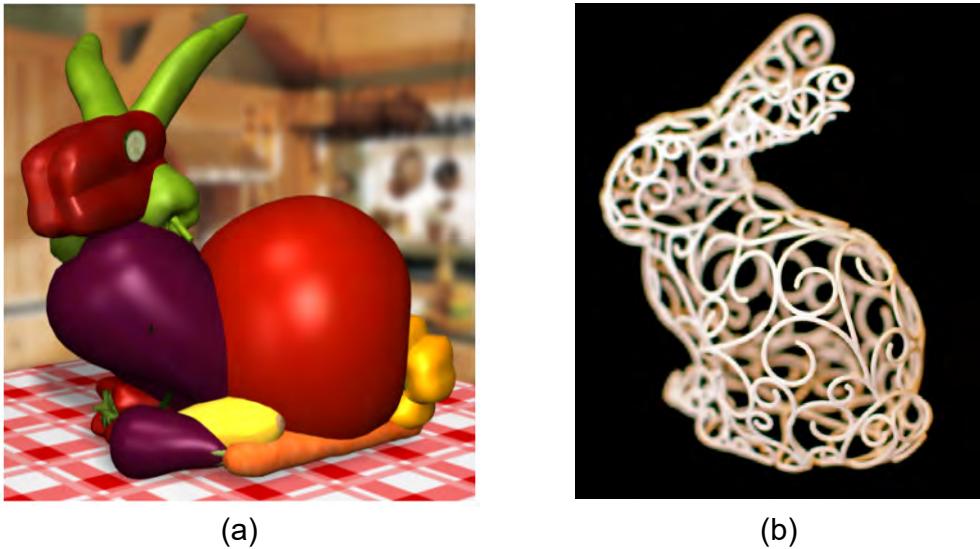


Figure 2.6: (a) An Arcimboldo-like 3D packing [GSP<sup>+</sup>07]. (b) Decorative elements that fill a surface [ZCT16].

422 Chapter 3

423 **FLOWPAK: Flow-Based Ornamental Element Packing**  
424



Figure 3.1: Ornamental packings of a lion and a unicorn. The diagram next to each animal shows a set of four ornamental elements used in the packing (top) and the annotated container regions (bottom). Each ornamental element has a red spine that is used to deform it along a streamline. In the containers, black curves represent boundaries, red curves with arrows represent directional guides, and green curves are fixed elements copied into the final design. The colors in the final rendering were added manually.

## 425 3.1 Introduction

426 FLOWPAK [SKAM17] is a technique for filling a container region with elements that are  
427 deformed to communicate a sense of directionality or flow. Elements have designs of simple  
428 geometric forms, often stylized flora, spirals, or other abstract shapes. Figure 3.2 shows  
429 four examples of these sorts of compositions, which we refer to as ornamental packings. 

430 In studying designs like those in Figure 3.2, we have identified five high-level principles  
431 that are important to ornamental packing construction: 

- 432 • **Balance.** A composition does not exhibit too much variation in local amounts of  
433 positive and negative space. Typically, this goal is accomplished by limiting variation  
434 in the diameters of elements (controlling the variation in positive space), and in ensuring  
435 that elements are spaced evenly (controlling negative space).
- 436 • **Flow.** In local parts of a composition, the elements are oriented to communicate a  
437 sense of directionality or flow. All of the examples in Figure 3.2 exhibit some amount  
438 of flow. In the dog, many elements appear to flow outward from the flower in the  
439 center of the torso, and then up the neck and down into the legs. The scales and other  
440 elements on the fish flow along the length of its body. In the lion and skull, elements  
441 flow horizontally outward from a central axis of symmetry, suggesting fur in the case of  
442 the lion. Flow adds visual interest to a composition, engaging the viewer by providing  
443 a sense of progression and movement through elements.
- 444 • **Uniformity Amidst Variety.** Repeated elements must balance between two opposing  
445 forces. *Uniformity* aims for an overall unity of design; *variety* seeks to break up the  
446 monotony of pure repetition. Elements should be permitted to vary in shape, but in a  
447 controlled way. We refer to this principle as *uniformity amidst variety*, a term borrowed  
448 from philosopher Francis Hutcheson [Hut29]. Gombrich also writes eloquently on the  
449 role of variation in design [Gom84]. In our examples, the dog’s spirals and the fish’s  
450 scales both obey this principle. The lion and skull do as well, except that half of the  
451 elements are reflected copies of the other half, across a vertical line through the center  
452 of the composition. This repetition emphasizes the bilateral symmetry in the design.
- 453 • **Fixed Elements.** Compositions use a small number of fixed elements to solve specific  
454 design problems or provide focal points. In any figurative drawing, eyes serve as a  
455 powerful focal point; every example in Figure 3.2 has eyes drawn in as unique elements  
456 (the dog’s eye is expressed via a carefully placed spiral). Other situations that call for  
457 specialized shapes include the dog’s paws, the fish’s teeth and fins, the lion’s eyes and  
458 nose, and the skull’s teeth. Sharp variation in the balance of positive and negative

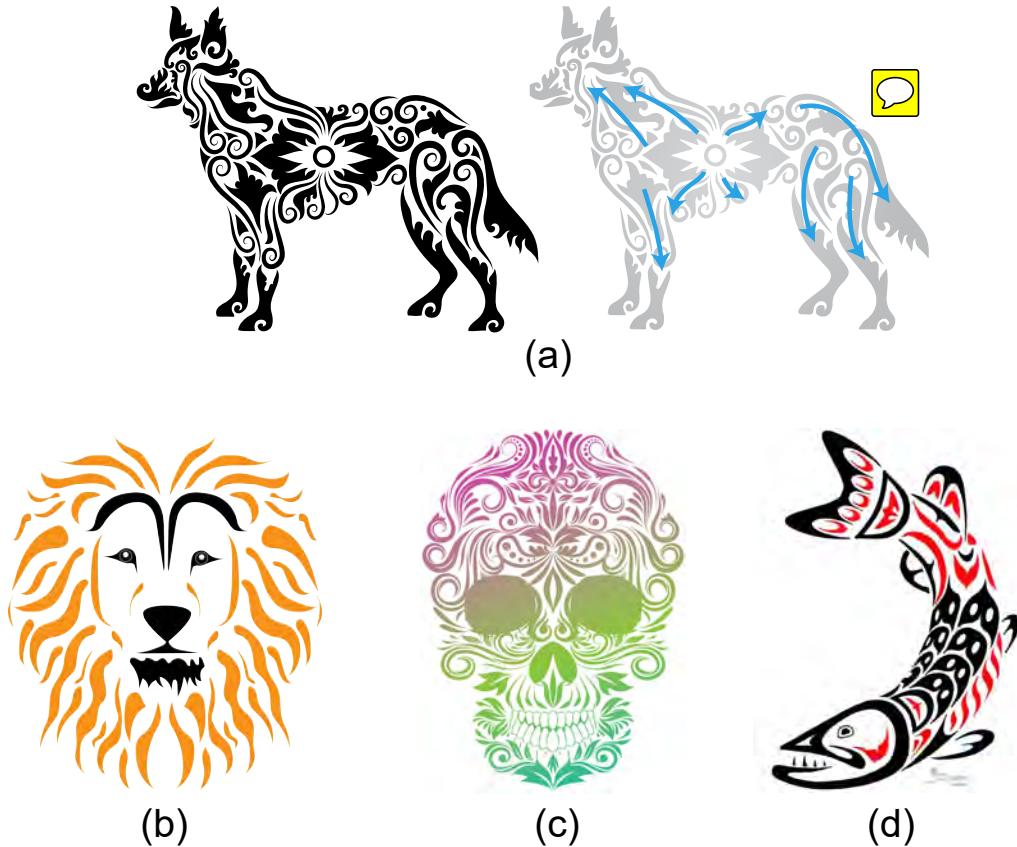


Figure 3.2: Examples of ornamental packings: (a) Dog (by ComicVector703 on Shutterstock), including a visualization showing the flow directions of elements; (b) Lion (from StockUnlimited); (c) Skull (by alitdesign on Shutterstock). (d) Fish in the style of Haida art (by Russ Jones, used with permission);

459 space can also be used to emphasize a focal point, as in the fish’s head and the lion’s  
 460 face that contain considerable amount of empty space.

461 • **Boundaries.** In many ornamental packings, elements are carefully arranged to conform  
 462 to and emphasize container boundaries. The fish demonstrates this principle most  
 463 clearly: we can easily fill in the gaps between elements to form a mental image of a  
 464 continuous outline. The dog’s elements are also well aligned to indicate the container  
 465 shape. However, this rule is not universal. The lion artfully subverts it with elements  
 466 that flow outward to an indistinct boundary, helping to convey the appearance of fur.

467 In FLOWPAK, elements can be oriented in the local direction of flow, but can also be

468 deformed to capture changes in flow direction. We express the user’s desired flow by placing  
469 evenly spaced streamlines inside the container region. Each streamline is then replaced by  
470 an element chosen from a pre-drawn set. The element is bent along the streamline to  
471 communicate flow, and also deformed to balance the usage of negative space with elements  
472 placed on adjacent streamlines. The final designs, such as the lion and unicorn shown in  
473 Figure 3.1, aim to obey the design principles articulated above.

## 474 3.2 Related Work

475 **Packings:** Chapter 2 discusses methods on the generation of packings, or mosaics. How-  
476 ever, past work is not appropriate for creating designs like those of Figure 3.2. Most  
477 techniques pack elements via rigid transformations, leading to high uniformity but insuf-  
478 ficient variety. We design FLOWPAK to be a deformation-driven approach so that it can  
479 generate plausible families of related decorative elements from a single input shape.

480 **Decorative Ornaments:** A distinct category of past research seeks to develop explicit  
481 procedural models for authoring decorative patterns. Wong et al. [WZS98] articulated  
482 a set of design principles for decorative art: repetition, balance, and conformation to  
483 geometric constraints. They went on to describe a grammar-like system for laying out floral  
484 ornament. Beneš et al. [BvMM11] developed an interactive interface to guide procedural  
485 models in generating decorative elements. Guerrero et al. [GBLM16] developed PATEX, a  
486 system that preserves high-level geometric relationships like symmetry and repetition while  
487 ornamental designs are being edited. Gieseke et al. [GALF17] developed a user interface  
488 where an artist can generate a decorative pattern by specifying design principles such as  
489 flow, balance, and symmetry. Li et al. [LBMH19] developed a method that can produce  
490 a 2D quilting pattern by generating connected stitch paths, each ~~is then~~ decorated by  
491 ornamental elements.

492 **Decorative Strokes:** FLOWPAK is related to decorative stroke methods since we aim  
493 to deform long thin elements. The goal of these methods is to create stylized decorative  
494 strokes, and not to fill a container area. Hsu et al. developed Skeletal Strokes [HLW93], a  
495 method to warp decorative patterns along a stroke. Asente [Ase10] improved upon Skeletal  
496 Strokes by correcting severe deformation on high-curvature strokes. Lu et al. ~~developed~~  
497 DecoBrush [LBW<sup>+</sup>14] ~~that~~ can join multiple decorative patterns seamlessly.

498 **Vector Field-Guided Compositions:** FLOWPAK requires elements to be deformed  
499 and oriented according to a vector field. Xu and Mould [XM09] ~~developed a tracing method~~  
500 ~~of particles in a magnetic field to generate decorative curves.~~ Li et al. [LBZ<sup>+</sup>11] developed

501 a method to decorate a surface using a shape grammar that is guided by a tensor field.  
502 Maharik et al. explored Digital Micrography [MBS<sup>+</sup>11], in which lines of small-scale text  
503 are deformed to fit along dense streamlines in a container. These streamlines are traced  
504 from a user defined vector field. We take inspiration from Digital Micrography, but seek  
505 to place fewer, larger elements taken from a small library of elements and to use shorter,  
506 sparser, and less regular streamlines. Lastly, Xu and Mould [XM15] proposed a graph-  
507 based tree synthesis method that is guided by vector fields.

508 **Decorative Ornaments on Surfaces:** Related methods in fabrication has sought  
509 to cover surfaces with arrangements of decorative elements that satisfy manufacturing  
510 constraints such as connectivity [CZX<sup>+</sup>16, ZCT16, BWL18, MSS<sup>+</sup>19]. Elements must be  
511 ~~tightly connected~~ to produce a connected result that will hold together when 3D printed.  
512 In the case of FLOWPAK, we seek to distribute disconnected elements inside a container.

513

### 514 3.3 Problem Formulation

515 We formally define the problem to be solved as follows. The user provides several pieces  
516 of input to our system:

- 517 1. A set of target containers. Each container is a closed curve to be filled with orna-  
518 mental elements.
- 519 2. A set of direction guides that guide the placement algorithm, defining the flow of the  
520 results. Every target container must have at least one guide, and some or all of the  
521 guides typically follow the container boundaries.
- 522 3. An optional set of fixed elements that we transfer directly to the result.
- 523 4. A set of ornamental elements, each with a spine that will control its deformation.

524 The first three inputs are combined into a single diagram, where they are distinguished  
525 by their colors; see the left and right drawings in Figure 3.1. The goal of our algorithm is  
526 to fill each target container with elements, trying to satisfy several guiding principles:

- 527 1. Follow the flow defined by the direction guides.
- 528 2. Have as little empty space as possible.

- 529        3. Make the spacing between elements be as even as possible.  
530        4. Conform to container boundaries.  
531        5. Vary element width and length to avoid an excessively uniform arrangement.

532        The next section describes how we achieve these results.

## 533        **3.4 Approach**

534        Figure 3.3 gives an overall view of our system. The numbers in the following steps corre-  
535        spond to parts of the figure.

- 536        1. Read the input target containers and copy any fixed elements to the output art  
537              (Section 3.4.1).
- 538        2. Analyze the ornamental elements, creating a shape descriptor for each (Section 3.4.2).
- 539        3. Use the direction guides to fill each target container with a vector field then trace  
540              streamlines (Section 3.4.3).
- 541        4. Divide the target containers into blobs around the streamlines (Section 3.4.4).
- 542        5. Use the element shape descriptors to determine the best element for each blob. Place  
543              the element in the blob, treating it as a skeletal stroke and mapping its spine to the  
544              streamline (Section 3.4.5).
- 545        6. Iteratively refine the placement to eliminate empty areas and make the spacing more  
546              even (Section 3.4.6).

### 547        **3.4.1 Target Containers**

548        The input diagram contains a set of target containers. Each is a single closed curve defining  
549        an area to be filled. Most non-trivial examples include more than one target container. For  
550        the most part, our algorithm fills each container separately, and so the following explanation  
551        is given in terms of a single container. Containers will later be merged in the iterative  
552        refinement step.

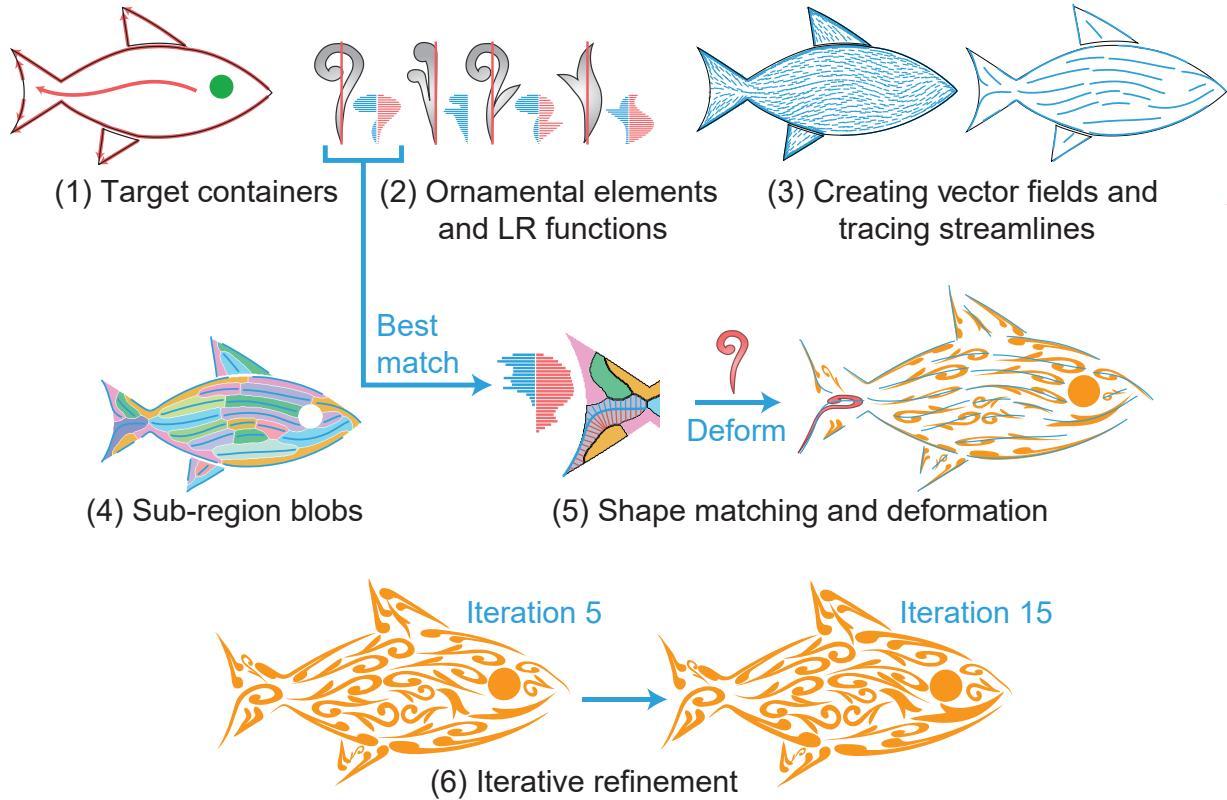


Figure 3.3: A visualization of the steps in our ornamental packing algorithm. The input containers are shown as three black outlines in (a): the body and two fins. They are annotated with directional guides in red and fixed elements (in this case, the eye) in green. The steps in the algorithm follow the description at the beginning of Section 3.4.

553      The artist has the option of including a set of fixed elements that we copy directly into  
 554      the final result. The following sections include descriptions of how the fixed elements affect  
 555      the filling algorithm 

556      We define *input\_size* to be the maximum of the combined width or height of all the  
 557      target containers and fixed elements as laid out by the artist. This value will be used to  
 558      set various parameters in the synthesis process.

### 559    3.4.2 Ornamental Elements and LR functions

560    An ornamental element is defined as one or more closed curves. Our placement method  
 561    will eventually deform copies of the element (Section 3.4.5) using a simple skeletal stroke

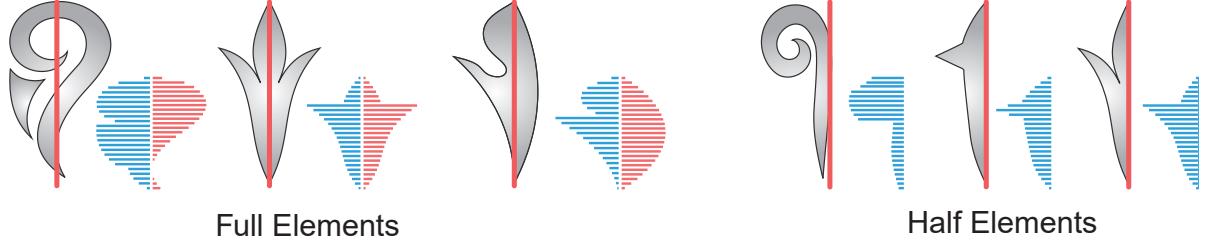


Figure 3.4: Ornamental elements and their LR functions. Full elements have non-empty left and right sides, while half elements have only one non-empty side. We normalize the LR functions to a unit square.

562 algorithm [HLW93], so each element must be annotated with a straight spine to guide the  
 563 deformation. The spine does not need to go through the center of the element—it can be  
 564 anywhere.

565 We define two classes of elements: a *full element* extends across both sides of its spine,  
 566 and a *half element* lies entirely on one side of its spine. Figure 3.4 shows examples of full  
 567 and half elements. If the input to our algorithm includes direction guides that coincide  
 568 with target container boundaries, the placement method will align half elements along  
 569 these boundaries. If half elements have edges that closely follow their spines, they will  
 570 visually reinforce container boundaries, as shown in our examples.

571 We define a simple shape descriptor called an *LR function* that will be used in Section  
 572 3.4.5 to choose which element to place in a particular location. Inspired by the work  
 573 of Gal et al. [GSCO07], we sample the element’s spine at  $n$  locations and at each location  
 574 determine how far the ornament extends to the left and right of the spine. The LR function  
 575 is the set  $\{L, R\}$  where  $L = \{\ell_1, \dots, \ell_{n_f}\}$  is the left function and  $R = \{r_1, \dots, r_{n_f}\}$  is the  
 576 right function. The number of samples is denoted by  $n_f$ .

577 The LR function is made scale-invariant by normalizing its domain and range to  $[0, 1]$ .  
 578 Note that swapping the  $L$  and  $R$  functions corresponds to reflecting the element across its  
 579 spine, and reversing each of  $L$  and  $R$  corresponds to reflecting the element along its spine.  
 580 We will consider all four combinations of these two reflections when placing an element in  
 581 a blob (Section 3.4.4), in order to achieve the best possible fit.

582 Intuitively, LR functions give an approximate area an ornamental element can claim.  
 583 Figure 3.4 shows elements with their left values in blue and their right values in red. We  
 584 have found that  $n_f = 100$  gives sufficient granularity for our algorithm.

585    **3.4.3    Creating Vector Fields and Tracing Streamlines**

586    To implement the flow principle described in the introduction, we fill each target container  
587    with a vector field, constrained by the direction guides in that container.

588    We sample the directional guides  $D = \{d_1, d_2, \dots, d_{n_d}\}$  and use the tangent at every  
589    sampled point as a directional constraint. We then construct a vector field using the  $N$ -  
590    RoSy algorithm of Palacios and Zhang[PZ07]. Note that, as shown in Step 3 of Figure 3.3,  
591    fixed elements do not affect the vector field. The artist can include directional guides to  
592    guide the vector field around fixed elements if desired.

593    The next step is to trace streamlines in the vector field, guided by three input param-  
594    eters:

- 595    *s\_gap* is the desired space between streamlines;  
596    *s\_max* is the maximum desired streamline length; and  
597    *s\_min* is the minimum desired streamline length.

598    Because we will ultimately place elements along streamlines without overlap, *s\_gap* deter-  
599    mines the approximate width of the placed elements, and *s\_max* the maximum length. We  
600    also derive a value *s\_stop* that prevents streamlines from coming too close to each other;  
601    in our implementation we compute  $s\_stop = 0.8 s\_gap$ .

602    We adapt the streamline tracing algorithm of Jobard and Lefer [JL97]. First we generate  
603    a set of potential seed points  $P = \{\mathbf{p}_1, \mathbf{p}_2, \dots, \mathbf{p}_{n_p}\}$  by densely resampling the target  
604    container boundary  $T$  and the directional guides in  $D$ . We use a sampling distance of  
605     $0.005 \text{input\_size}$ . The first streamline  $s$  is generated by randomly removing a seed point  
606    from  $P$  and following the vector field until one of the following conditions holds:

- 607    1. the length of  $s$  would exceed *s\_max*.  
608    2.  $s$  would come within *s\_stop* of another streamline.  
609    3.  $s$  would cross  $T$ , leaving the container.  
610    4.  $s$  would cross the boundary of a fixed element.

611    If the length of  $s$  is less than *s\_min*, we discard it. Otherwise we sample  $s$ , again using  
612     $0.005 \text{input\_size}$ , and at each point generate two more potential seeds that are *s\_gap* away  
613    from  $s$  on either side. If a seed is inside the container, we add it to  $P$ . The process is  
614    repeated until  $P$  is empty. Note that the *s\_stop* distance test combined with the *s\_min*

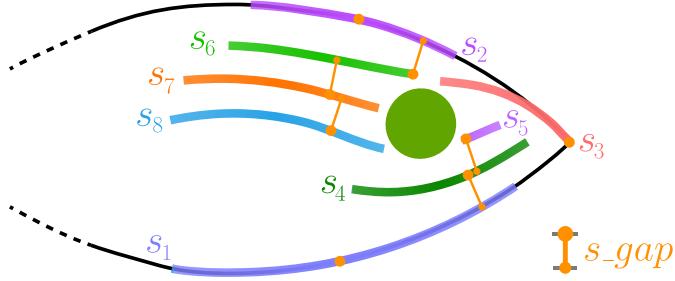


Figure 3.5: The streamline tracing process. The first streamline  $s_1$  always begins on a directional guide or the container boundary. Subsequent streamlines begin on the container boundary, a directional guide, or at a point that is  $s\_gap$  away from a previous streamline.

length test imply that many attempts to form streamlines will stop immediately, especially as the container fills with streamlines.

Figure 3.5 shows the creation process, and Algorithm 1 shows the pseudocode. The sort function  $SORT(P)$  orders the points in  $P$  according to their distance from the boundary  $T$  and the directional guides in  $D$ , with closer points first and equally distant points ordered randomly. Because the initial points are all on  $T$  or on a path in  $D$ , their sort value is zero, and they will be processed before any derived points.

---

#### Algorithm 1 Tracing Streamlines

---

```

Create a seed list  $P = \{p_1, p_2, \dots, p_{n,p}\}$  by uniformly resampling
 $T$  and the guides in  $D$ .
Create an empty set  $S$  of streamlines.
Randomly order the elements of  $P$ .
while  $P$  is not empty do
    Generate a new streamline  $s$  from  $p_1$ .
    Remove  $p_1$  from  $P$ .
    if  $s$  is longer than  $s\_min$  then
        Add  $s$  to  $S$ .
        Create seed points that are  $s\_gap$  away from  $s$  and
        add them to  $P$ .
         $SORT(P)$ .
    end if
end while

```

---

622 **3.4.4 Sub-Region Blobs**

623 To assist in choosing which element to place along each streamline, we first subtract the  
624 areas of any fixed elements from the target container. We then construct an approximate  
625 generalized Voronoi diagram of the interior using the method of Osher and Sethian [OS88].  
626 The streamlines are then extended at each end, following the vector field, until they en-  
627 counter the boundaries of their Voronoi regions. We call the area around each streamline  
628 a *sub-region blob*. Step 4 of Figure 3.3 shows the blobs of the fish.

629 We then compute an LR function for each blob as described in Section 3.4.2, using the  
630 streamline as the spine. Because the streamline is not usually straight, we compute the left  
631 and right distances along the normals to the streamline. The LR function approximates  
632 the blob's shape if the streamline were to be straightened.

633 **3.4.5 Shape Matching and Deformation**

634 The next step is to place an ornamental element in each blob. We choose which element  
635 to place in the blob by finding the element that minimizes a sum of least squares distance,  
636 defined as

$$\sum_{i=1}^N (\alpha_{li} - \beta_{li})^2 + \sum_{i=1}^N (\alpha_{ri} - \beta_{ri})^2 \quad (3.1)$$

637 where

- $\alpha_l$  is the element left function;  
 $\alpha_r$  is the element right function;  
638  $\beta_l$  is the blob left function; and   
 $\beta_r$  is the blob right function.

639 Every element can be placed in one of four orientations, by optionally incorporating  
640 reflections across and along its spine. These reflections correspond, respectively, to swap-  
641 ping the *L* and *R* functions and reversing them. When comparing the LR functions for an  
642 element and a blob, we compute the least squares distances for all four orientations and  
643 choose the orientation with the smallest distance. Note that this matching method ~~automati-~~  
644 ~~cally places~~ half elements along streamlines that follow container boundaries, visually  
645 reinforcing the overall shape.

646 We investigated alternatives for shape matching, using an approach discussed by Gal et  
647 al. [GSP<sup>+</sup>07] that tries to fill a sub-region blob as much as possible, with heavy penalties

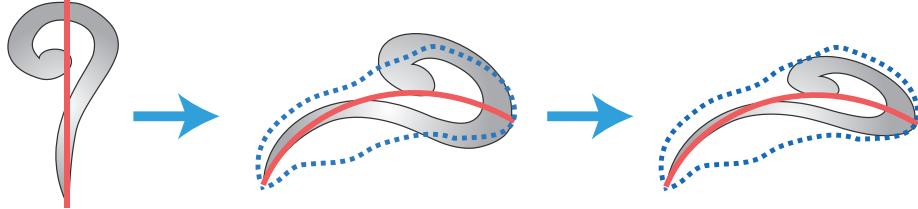


Figure 3.6: The deformation process bends the element along the streamline and scales it to fit inside the blob.

648 if a part of an element protrudes outside the boundary of the blob. However, we found  
 649 this computation to be more expensive without providing significant advantages over our  
 650 LR functions.

651 Once we have chosen an element, we place it along the streamline using a simple skeletal  
 652 stroke algorithm [HLW93]. We uniformly scale the element’s width to make it as wide as  
 653 possible while still staying inside the blob (Figure 3.6).

### 654 3.4.6 Iterative Refinement

655 We now refine the overall composition in an iterative process. We perform this part of  
 656 the algorithm globally, by merging all containers and allowing the elements within them  
 657 to interact. 

658 The refinement process aims to reduce the amount of negative space and make it more  
 659 even by growing and shifting the placed ornamental elements. It would be possible to use  
 660 a greedy approach, improving the placement of each element as much as possible before  
 661 moving on to the next. However, we have found that gradually improving the placement  
 662 of all elements leads to a more even result.

663 Each refinement iteration has two phases. First, we shift the streamlines to more  
 664 accurately follow the space that is available, as shown in Figure 3.7. After shifting, we  
 665 recalculate the LR function for the blob to reflect the new position, and repeat a variant  
 666 of the element placement process that allows the elements to rotate slightly in their space.  
 667 Second, we expand each blob to allow it to use adjacent space that is not filled with another  
 668 element, as shown in Figure 3.9.

669 Each refinement iteration considers the blobs in increasing order of placed element area,  
 670 allowing smaller elements to grow more. While each step usually results in a larger placed  
 671 element, some configurations can result in a smaller one. We only accept the new element

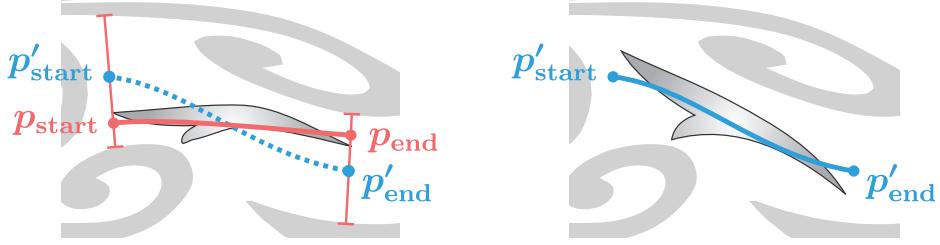


Figure 3.7: Streamline shifting. We move the streamline’s start and end points along perpendiculars, stopping before intersecting neighboring elements.

672 if its area is no smaller than  $\alpha$  times its old area, where  $\alpha$  is a growth tolerance that we  
 673 set to 0.9. Elements therefore have some freedom to grow or shrink, in the search for more  
 674 globally even spacing.

675 Algorithm 2 gives the overall method, and the following sections give details. We have  
 676 found that 15 iterations suffice for most designs. Note that in Algorithm 2, the variable  $E$   
 677 is the list of placed, distorted ornamental elements, and not the set of prototype elements  
 678 discussed earlier.

679 **Shifting Streamlines.** There are two issues that keep the initial placement of elements  
 680 from being evenly distributed. Our streamline placement method keeps streamlines apart,  
 681 but they may not be spaced completely evenly. More significantly, the ornamental elements  
 682 often have unbalanced left and right sides and concavities, leading to extra space on one  
 683 side or the other. Our refinement process shifts streamlines to address these problems.

684 The shifting process allows the endpoints of the streamline to move to the left or the  
 685 right relative to the streamline, depending on which side has more empty space. This  
 686 allows the streamline’s element to become wider and fill more of the space (Figure 3.7). It  
 687 also gives the streamline room to extend if its endpoints were too close to boundaries of  
 688 other placed elements.

689 Given the endpoints  $p_{\text{start}}$  and  $p_{\text{end}}$ , we calculate new endpoints  $p'_{\text{start}}$  and  $p'_{\text{end}}$ . We  
 690 generate perpendicular vectors to the left side and to the right side at each endpoint  
 691 and construct a line segment joining the points where the vectors intersect other placed  
 692 elements. We then move the endpoint of the streamline towards the midpoint of this  
 693 segment. To enforce the principle of gradual refinement, we do not allow the endpoint to  
 694 move more than  $g_{\text{limit}}$  units, where  $g_{\text{limit}} = 0.005 \text{ input\_size}$  (Recall that  $\text{input\_size}$  is the  
 695 maximum dimension of the design as described in Section 3.4.1).

696 We replace the streamline with a path joining  $p'_{\text{start}}$  and  $p'_{\text{end}}$ . Our goal is to create a  
 697 path that is smooth and does not deviate too much from the vector field. We calculate the

---

**Algorithm 2** Iterative Refinement

---

**Input:**  $E = \{e_1, e_2, \dots, e_{n\_e}\}$  as the ornamental element list.

**Input:**  $S = \{s_1, s_2, \dots, s_{n\_s}\}$  as the streamline list.

**Input:**  $B = \{b_1, b_2, \dots, b_{n\_b}\}$  as the blob list.

**Input:**  $\alpha$  as the growth tolerance

**Input:**  $t$  as the number of iterations

**for**  $t$  times **do**

    Sort  $E$ ,  $S$ , and  $B$  by  $\text{AREA}(e_i)$  (smallest first)

**for** Element  $e_i$  in  $E$  **do**

$s_i$  is the corresponding streamline of  $e_i$

        Calculate  $s'_i$  by **shifting**  $s_i$ .

        Recompute the LR function of  $b_i$  to give  $b'_i$

        Calculate  $e'_i$  by **placing**  $e_i$  inside  $b'_i$

**if**  $\text{AREA}(e'_i) \times \alpha > \text{AREA}(e_i)$  **then**

$s_i \leftarrow s'_i$

$b_i \leftarrow b'_i$

$e_i \leftarrow e'_i$

**end if**

**end for**

    Sort  $E$ ,  $S$ , and  $B$  by  $\text{AREA}(e_i)$  (smallest first)

**for** Element  $e_i$  in  $E$  **do**

$b_i$  is the corresponding blob of  $e_i$

        Calculate  $b'_i$  by **growing**  $b_i$ .

        Calculate  $s'_i$  based on  $b'_i$

        Calculate  $e'_i$  by **placing**  $e_i$  inside  $b'_i$

$b_i \leftarrow b'_i$

**if**  $\text{AREA}(e'_i) \times \alpha > \text{AREA}(e_i)$  **then**

$s_i \leftarrow s'_i$

$e_i \leftarrow e'_i$

**end if**

**end for**

**end for**

---

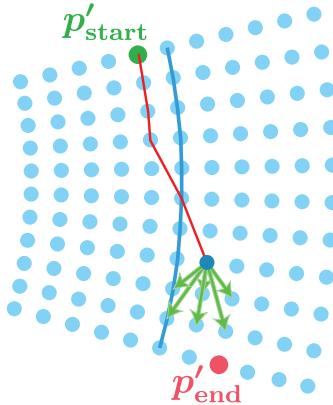


Figure 3.8: Tracing a shortest path using Dijkstra’s algorithm. We generate the blue nodes by resampling and offsetting the blue streamline. The search directions at a node are shown with green arrows.

shifted streamline by performing Dijkstra’s algorithm on a non-rectangular graph that respects the vector field (Figure 3.8), using a method similar to one by Xu and Mould [XM15] for pathfinding in a vector field.

To construct the graph, we begin by densely sampling the original streamline with a distance of  $0.25 g_{\text{limit}}$ . We then duplicate the points, offset to the left and right, again using  $0.25 g_{\text{limit}}$ . The duplication is repeated until the graph extends to the left and right of the streamline by a distance equal to the maximum left and right widths of the blob (i.e., the maximum values in an unnormalized version of the blob’s LR function).

For a node, we check its  $N = 150$  nearest neighbors, considering only neighbors where the angle between the line into the current node and the line to the neighbor form an angle greater than  $90^\circ$ , thereby preventing the streamline from backtracking. The cost of an edge from  $\mathbf{n}_a$  to  $\mathbf{n}_b$  is

$$w = w_f(1 - f^p) + w_d D(s_i, \mathbf{n}_a) \quad (3.2)$$

where

$f$  is  $(\mathbf{n}_b - \mathbf{n}_a) \cdot \mathbf{v}$ ;

$\mathbf{v}$  is a sampled vector of the vector field; 

$s_i$  is the original streamline; and

$D()$  is a distance function between a polyline and a point.

712 In practice, we set  $w_d = 0.1$ ,  $w_f = 1$ , and  $p = 3$ .

713 After finding a set of points, we fit cubic Bézier curves using a method devised by  
714 Schneider [Sch90] and extend the path at both ends by following the vector field until it  
715 intersects the edges of its blob.

716 **Growing Blobs.** The growth process tries to enlarge each sub-region blob to claim  
717 empty space. Given a blob  $b_i$ , we calculate a larger blob  $b'_i$  by offsetting its boundaries  
718 until they intersect other placed elements (Figure 3.9). To enforce gradual growth, the  
719 offset cannot be larger than  $g_{\text{limit}}$ , where  $g_{\text{limit}} = 0.005 \text{ input\_size}$ .

720 The value  $g_{\text{limit}}$  used in growing blobs and shifting streamlines limits the speed of the  
721 refinement. Making it larger would require fewer iterations to fill the available space, but  
722 at a cost of elements growing less evenly.

723 **Element Placement.** In the refinement process, we allow a more flexible element  
724 placement so that the elements can fill more of their blobs. We allow the element to  
725 rotate by a small amount, up to ten degrees, before placing it, as shown in Figure 3.10. We  
726 generate rotated versions of  $e_i$  with varying angles  $r_{\text{angle}} = 1^\circ, 2^\circ, 3^\circ, \dots, 10^\circ$  and precompute  
727 LR functions for each. The shape matching algorithm (Section 3.4.5) automatically chooses  
728 the best rotation. It can also choose to reflect the element across its spine, along its spine,  
729 or both (Figure 3.11).

## 730 3.5 Implementation and Results

731 We design our containers and decorative elements in a vector graphics editor, and then  
732 use them as input to a C++ program that outputs final placed elements in an SVG file.  
733 We use the Clipper library [Joh14] for calculation of LR functions and for testing polygon  
734 intersections during deformation and growth. As a postprocess, we optionally smooth  
735 outlines and replace polygonal paths with Bézier curves. Finally, we apply colors and  
736 other treatments in an editor.

737 Our technique is fast except for the iterative refinement process, which considers a large  
738 number of variations to the composition via brute-force computation. On a computer  
739 with an Intel i7-4790K processor at 4.0 Ghz, 15 iterations of refinement on a packing  
740 of 50 elements takes about an hour. Our software is not intended to run interactively;  
741 still, we believe the performance could be improved significantly through the use of more  
742 sophisticated 2D geometric data structures such as uniform grids or quadtrees.

743 We tested our approach using a variety of container shapes, based mostly on animals,  
744 and many different ornamental elements with varying amounts of geometric complexity.

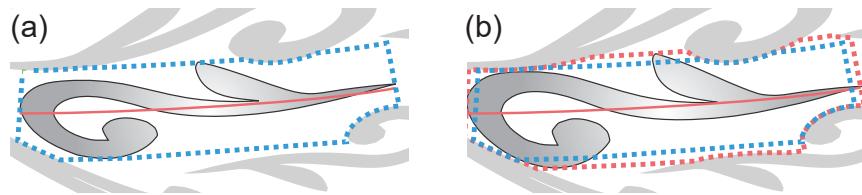


Figure 3.9: (a) An element with its sub-region blob shown in dashed blue line. Note that any blob is constrained by the neighboring elements. (b) The dashed red line is the *grown* blob, which accommodates an enlarged element.



Figure 3.10: **Left:** rotated versions of the original element. The best rotation angle is chosen via least squares matching. **Right:** original, rotated, and enlarged versions of an element.

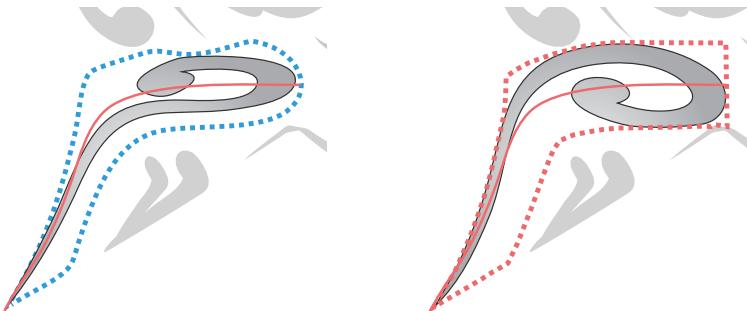


Figure 3.11: An element that reflects across its spine during iterative refinement. LR functions and least squares shape matching allow an element to reflect across its spine, along its spine, or both.

745 In Figure 3.1, we show two packings of a lion and a unicorn. Each packing is generated  
746 with only a set of four elements. The packings demonstrate that FLOWPAK is able to  
747 pack and deform elements inside the containers. In Figure 3.12, we show a packing of  
748 a rhinoceros with simple teardrop elements that demonstrates the variety we achieve in  
749 shape and curvature. We use more complex leaf elements on the bear in Figure 3.13,  
750 and adjust the tracing parameters to obtain shorter placed elements. We also process the  
751 placed elements to create a distressed look.

752 The packing of a cat in Figure 3.14 demonstrates a symmetric packing with a fur  
753 contour. We ~~only compute~~ the left part and reflect the result. The elements around  
754 the cheeks and the chin extend outward, not following the boundary, and creating the  
755 appearance of fur.

756 We experimented with two extensions to our pipeline, which could enhance its aesthetic  
757 value and flexibility. First, in Figure 3.15 we allow the user to draw *fixed spines* in addition  
758 to fixed elements. These fixed spines act like pre-placed streamlines, which will be assigned  
759 blobs and then elements. However, they are not required to follow the surrounding vector  
760 field, and are not shifted during the refinement process. Fixed spines are used in Figure 3.15  
761 for the flower petals in the torso and the paws. Second, in Figure 3.16 we construct explicit  
762 new shapes (drawn in brown) to fill the negative space between placed elements (in black),  
763 by computing offset polygons from the negative space between elements. The result is a  
764 distinct and appealing style.

765 Finally, we asked an artist to draw containers and decorative elements. The result is  
766 the bird design shown in Figure 3.17. The artist requested that different elements and  
767 densities be used in different container regions; the result has sparse “Y” elements in the  
768 breast and head, and denser “O” elements in the wings. The artist was pleased with the  
769 **result**.

## 770 3.6 Conclusions

771 We presented FLOWPAK, a method to create ornamental packings in which the elements  
772 were oriented and deformed to give a sense of visual flow to the final composition. Our  
773 implementation computed a vector field based on user strokes, constructed streamlines  
774 that conform to the vector field, and placed an element over each streamline. An iterative  
775 refinement process then shifted and stretched the elements to improve the composition.

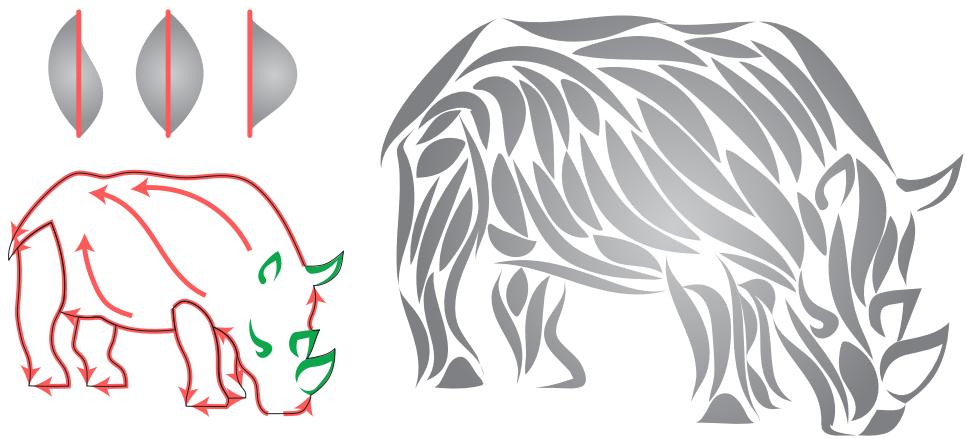


Figure 3.12: A packing of a rhinoceros. Simple teardrop-shaped elements lead to variety in size and curvature.



Figure 3.13: A bear packing with leaf elements. We manually add noise to the elements in the output to create a distressed look.

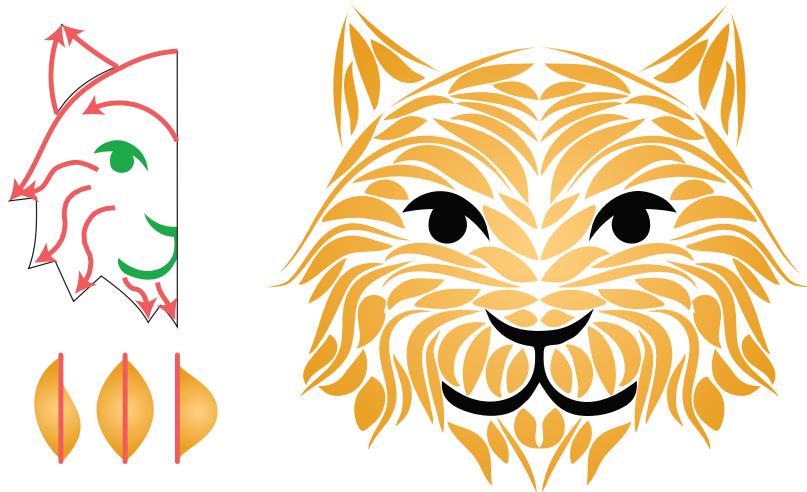


Figure 3.14: A packing with a symmetric layout; we only compute the left half and reflect the result. The elements around the cheeks and the chin are not aligned to the boundary, creating a fur-like effect.

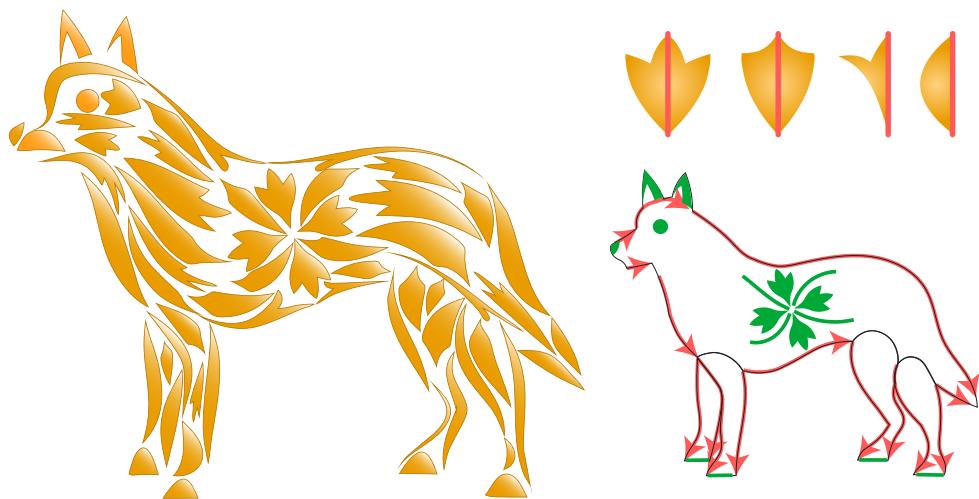


Figure 3.15: A packing of a dog. The fixed elements, shown as green shapes in the diagram, are copied as-is to the output; fixed spines, shown as green paths, force the placement of new elements at the given locations.



Figure 3.16: A packing of the same container as in Figure 3.13. We place longer and sparser elements and synthesize additional forms to fill negative space.

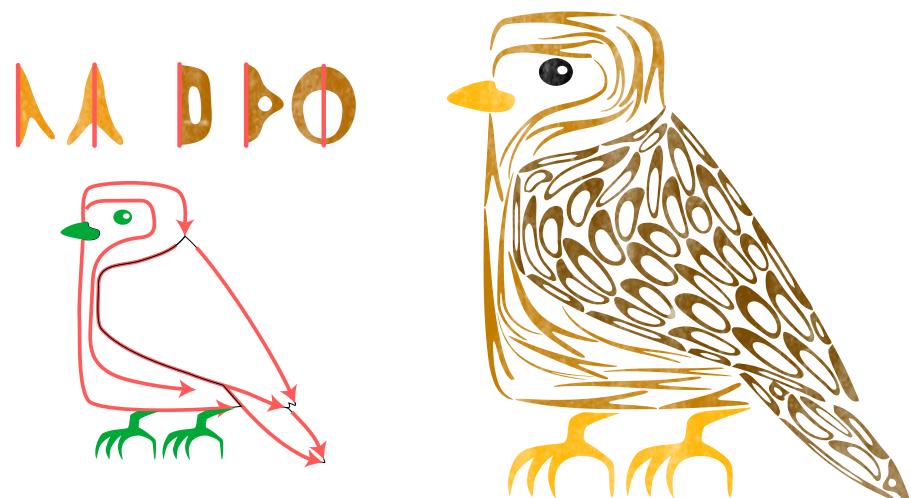


Figure 3.17: A packing of a bird, based on input provided by an artist.

776 Chapter 4

777 **RepulsionPak: Deformation-Driven**  
778 **Element Packing with Repulsion**  
779 **Forces**



Figure 4.1: (a) A packing of six cat elements inside a fish-shaped target container. (b) Input elements. (c) lined up deformed elements in the packing. Controllable deformation and repulsion forces allow the elements to deform, efficiently filling the container and creating a uniform distribution of negative space. We then reduce the remaining negative space by placing smaller cat heads. The gradient fill was added as a post-process.

## 780 4.1 Introduction

781 RepulsionPak is a technique to pack elements using a physical simulation, in which each  
782 element is represented by a mass-spring system called an *element mesh*. *Repulsion forces*  
783 between neighboring meshes work to even out the negative space, inducing displacements  
784 in mesh springs. These displacements translate, rotate, and deform the elements as they  
785 gradually adapt to the shapes of their neighbors and the container boundary. To control  
786 the amount of deformation, we use *spring forces* within a mesh to preserve element shapes. 

787 Most of the elements in a packing are large shapes of real-world objects like animals,  
788 plants, or man-made objects. We refer to these as *primary elements*. An artist distributes  
789 primary elements so that they communicate the shape of the container, while attempting  
790 as much as possible to ensure an even distribution of negative space. When the primary  
791 elements leave behind large pockets of negative space, the artist typically fills those pockets  
792 with small *secondary elements*, often simple abstract shapes like circles or triangles. A  
793 packing example with primary and secondary elements is shown in Figure 4.2.

794 Unlike FLOWPAK that is optimized to deform long thin elements, we design RepulsionPak  
795 to support arbitrary shaped elements. RepulsionPak is also intended to fulfill three of five  
796 design principles: balance, uniformity amidst variety, and boundaries (Section 3.  We  
797 achieve a more even distribution of negative space by building an algorithm with a con-  
798 trollable deformation model at its core. We are able to use repeated copies taken from a  
799 small library of elements but their final shapes are varied thanks to element deformation.  
800 We also show that by carefully placing elements in the initial distribution using a shape  
801 matching algorithm, the shape of the container boundary can be ~~more~~ emphasized.



## 802 4.2 Related Work

803 **Lloyd’s Method:** Variants of Lloyd’s method have been proposed by Hausner [Hau01],  
804 Hiller et al. [HHD03], Smith et al. [SLK05], and Dalal et al. [DKLS06]. Although these  
805 approaches can generate appealing results, they can only accept convex or near convex  
806 elements. RepulsionPak can be categorized as a variant of point-based Lloyd’s method  
807 without explicitly computing a Voronoi diagram. Forces do not act directly onto elements,  
808 but their mesh vertices. During the simulation, these mesh vertices are spread apart but  
809 still connected using elastic springs, effectively creating an element representation that can  
810 be packed and deformed.

811 **Data-driven Methods:** Data-driven methods include Image Mosaics (JIM) [KP02],  
812 the Pyramid of Arclength Descriptor (PAD) [KSH<sup>+</sup>16]. JIM is able to pack nearly-convex



Figure 4.2: A packing created by Balabolka. (a) The packing with primary elements only.  
(b) Secondary elements are added.

elements tightly by placing one element at a time and backtracking as needed. As a shape descriptor, PAD can find new elements that partially match existing element boundaries as a container was being filled. Both methods permit some elements to overlap. While they can both correct gaps and overlaps using deformation, the deformation is applied locally near edges in a post-processing step after elements were frozen in place. 2D Arcimboldo Collage [HZZ11] and 3D Arcimboldo Collage [GSP<sup>+</sup>07] are also data-driven, but they permit large overlaps to resemble the aesthetic style of Arcimboldo's artwork.

**Deformation-Driven Methods:** Xu and Kaplan [XK07] and Zou et al. [ZCR<sup>+</sup>16] deformed letters of one or two words to fit into a container. Because the order of the letterforms was defined by the text, their solutions require significant distortion of the individual letters. Peng et al. [PYW14] computed layouts by packing and deforming simple polygons and polyominoes. Their method cannot handle more complicated shapes, making it unsuitable for our style of packings.

Zehnder et al. [ZCT16] proposed a method to cover 3D surfaces with deformed ornamental elastic curves. Our method has some similarities to theirs in that both start with scaled-down copies of elements and grow them, but the growth process is quite different. Unlike their approach, our elements exert forces on each other throughout the growth process, allowing them a greater opportunity to translate, rotate, and deform in search of more even negative space. Furthermore, the goal of their work (3D fabrication) is quite different from ours (2D ornamentation) and our results appear qualitatively different.



**Physically-Based Methods for Deformable Objects:** Contemporary research has yielded many more sophisticated physical simulation methods. the Finite Element Method (FEM) is accurate and realistic but notoriously slow. Position Based Dynamics (PBD) [MHHR07] is simple, robust, and efficient to be used in real-time applications, but

837 suffers from inaccuracy. More recently, Projective Dynamics [BML<sup>+</sup>14] bridges a gap be-  
838 tween FEM and PBD. Projective dynamics uses a variational implicit time integration and  
839 optimizes by alternating a local constraint step and a global solve step.

840 **Sample-Based Representation:** A few off-texture synthesis methods represent an  
841 element as a collection of sparse points [MWT11, MWLT13, HWYZ20]. This representation  
842 allows them to synthesize textures with elements that have large concave shapes. Our  
843 RepulsionPak has a similar approach, but we represent an element with denser vertices  
844 that are connected with elastic springs.

845 **Curve Simulation:** Pedersen and Singh [PS06] grew curves to create organic labyrinthine  
846 paths. Their algorithm is related to ours by the use of repulsion forces to maintain even  
847 spacing and parallel segments. More recently, Yu et al. [YSC20] used tangent-point energy  
848 to pack curves inside a container while avoiding self-intersections.

### 849 4.3 System Overview

850 Our system requires three main pieces of input:

- 851 • A library of primary and optional secondary elements, such as those shown in Fig-  
852 ure 4.3a. Each element is a collection of open or closed polygonal paths—any curves  
853 must be flattened ahead of time.
- 854 • One or more closed polygonal target containers, such as the heart in Figure 4.3b.  
855 Target containers can optionally have internal holes.
- 856 • The desired element spacing distance  $d_{\text{gap}} > 0$ .

857 RepulsionPak starts by preprocessing the elements, creating additional space around  
858 each to enforce the spacing distance, and fitting a triangle mesh over each element. The  
859 containers are then seeded with randomly placed copies of scaled-down elements (Sec-  
860 tion 4.4).

861 RepulsionPak then performs a physics simulation on the meshes, making them simu-  
862 taneously grow and repel each other. As a proof of concept, we implement a spring-based  
863 simulation, which yields satisfactory results despite its simplicity; many alternatives are  
864 possible (see Section 4.9). Forces in the simulation push mesh vertices away from vertices in  
865 other meshes, attempt to keep the meshes from undergoing excessive deformation, and re-  
866 solve places where meshes overlap or vertices move outside target containers (Section 4.5).

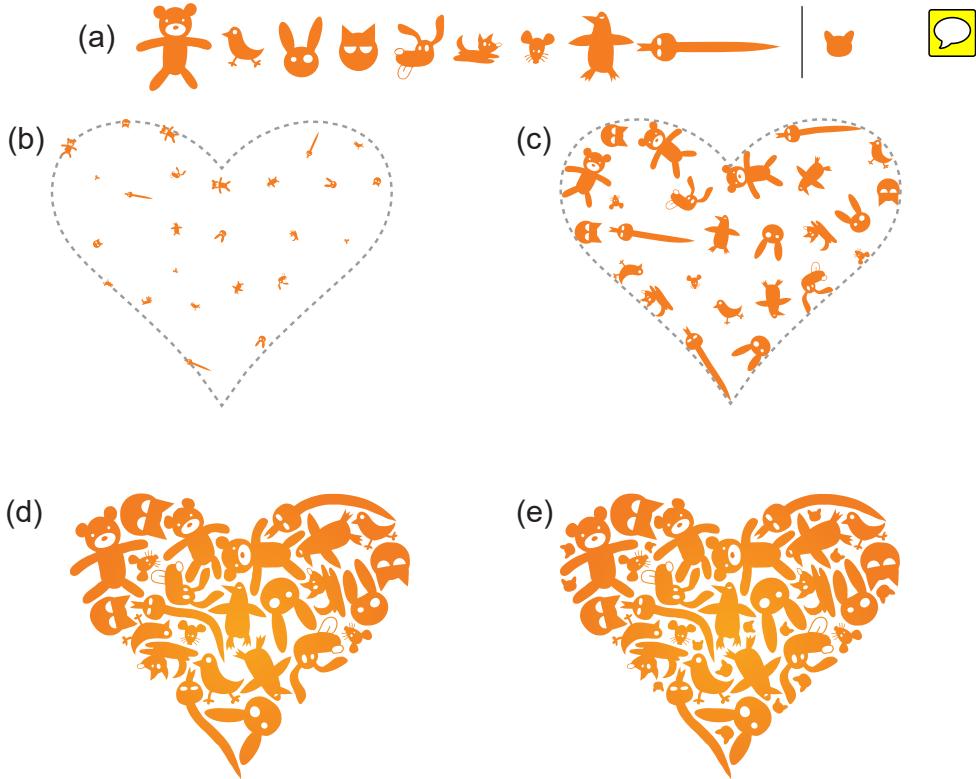


Figure 4.3: The creation of a packing using RepulsionPak. (a) A library of elements, comprising nine primary elements and a single secondary element. (b) A target container with the initial distribution of scaled-down elements. (c) The simulation in progress, showing the elements growing, translating, rotating, and deforming. (d) The resulting packing of primary elements. (e) The final result, after adding secondary elements and allowing them to grow. Figure 4.4 shows the deformations of some of the elements.

867 After each iteration of the simulation, meshes grow into adjacent space if possible, so that  
868 they gradually consume the negative space in the container.

869 The simulation concludes either when the elements occupy a sufficient proportion of  
870 the container area, or when some number of simulation steps fail to significantly reduce  
871 the negative space (Section 4.5.2).

872 An optional second simulation further reduces and evens out the negative space. It  
873 begins by placing small secondary elements in large pockets of negative space. This sim-  
874 ulation is the same as the first, except that vertices of primary element meshes are not  
875 allowed to move (Section 4.6).

876 Final SVG output is created by using barycentric coordinates to map each element's

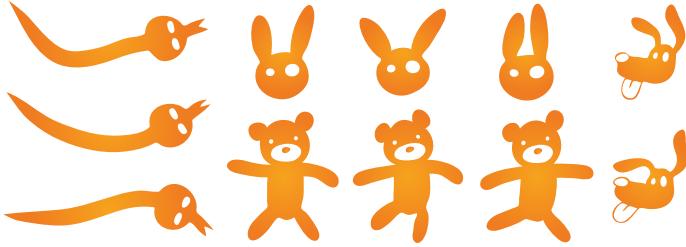


Figure 4.4: Some of the elements in the final packing in Figure 4.3, showing the effect of deformation in our simulation.

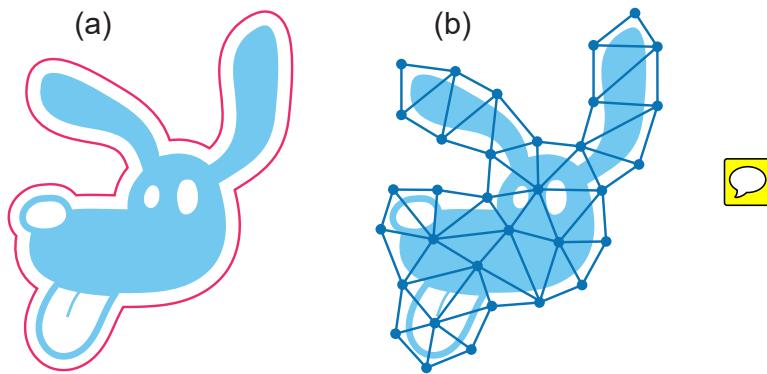


Figure 4.5: An illustration of element discretization for preprocessing. (a) An element with its boundary displaced to create a skin, drawn in red. (b) A triangle mesh with boundary vertices on the skin.

877 paths from the element's initial mesh into the deformed mesh produced through simulation. 

## 878 4.4 Preprocessing

879 The *skin* of an element is a simple closed polygon that fully encloses it, as in the red shape  
880 in Figure 4.5(a). We generate the skin by offsetting an element's boundary outward by  
881  $d_{\text{gap}}/2$ . The simulation aims to produce an approximate tessellation of the target container  
882 by *deformed* skins, thereby achieving the desired element spacing and suppressing overlaps.

883 We triangulate the element skin to obtain a triangle mesh. To create the mesh we uni-  
884 formly sample the skin polygon  $s$ , with samples spaced apart by distance  $d_{\text{gap}}$ , to obtain  
885 a simpler polygon  $s'$  (the outer boundary of the mesh). We then construct a Delaunay

886 triangulation of  $s'$ . The vertices and edges of this mesh become unit masses and longi-  
 887 tudinal springs in a physical simulation, allowing elements to deform in response to their  
 888 neighbors. We also add extra edges to prevent folding and self-intersections during sim-  
 889 ulation (Figure 4.6b). First, if two triangles  $ABC$  and  $BCD$  share an edge  $BC$ , then  
 890 we add a *shear edge* connecting  $A$  and  $D$ . Second, we triangulate the negative space in-  
 891 side the convex hull of the original Delaunay triangulation, and create new *negative space*  
 892 *edges* corresponding to the newly created triangulation edges. These negative space edges  
 893 are used exclusively for internal bracing. The element’s concavities can still be occupied  
 894 by its neighbors. The construction of negative space edges is a simpler variant of Air  
 895 Meshes [MCKM15], a technique to detect and resolve collisions.

896 Due to discretization, a low mesh resolution does not guarantee a separation of  $d_{\text{gap}}$ .  
 897 Increasing the mesh resolution produces a more precise result at the expense of greater  
 898 running time. 

899 **Barycentric Coordinates:** The simulation operates on meshes, not element geome-  
 900 try. In the final rendering phase, we will redraw an element relative to a deformed copy  
 901 of its mesh. To do so, we first re-express every vertex of an element path in terms of the  
 902 mesh triangles. Every element vertex lies either inside a mesh triangle or just beyond a  
 903 border edge. We encode each vertex in barycentric coordinates relative to its enclosing or  
 904 nearest triangle.

905 **Initial Element Placement:** We prepare our simulation by randomly placing non-  
 906 overlapping elements. 

- 907 1. Generate random points  $P = \{\mathbf{p}_1, \mathbf{p}_2, \dots, \mathbf{p}_n\}$  inside the target container via blue  
 908 noise sampling [Bri07]. The user controls the number of points; using more points  
 909 gives results with smaller elements. We can automatically estimate  $n$  by dividing the  
 910 container area by the desired average area of the element skins.
- 911 2. Cycle through the primary elements, assigning each element to a random unused  $p_i$   
 912 with a random orientation, repeating until every point has an element.
- 913 3. Shrink all the elements so that they do not overlap and occupy only a small fraction  
 914 of the container’s area; in our implementation we have found that 5 – 10% of the  
 915 area gives good results. Making them larger would speed up the simulation but does  
 916 not allow enough freedom of movement to generate successful packings. Figure 4.3b  
 917 shows an initial placement.

## 918 4.5 Simulation

919 We design a simulation in which we generate forces that pack and deform elements by  
 920 transforming their meshes. Let  $\mathbf{x} = (x, y)$  be a vertex of an element mesh. The total force  
 921  $\mathbf{F}_{\text{total}}$  applied to  $\mathbf{x}$  is

$$\mathbf{F}_{\text{total}} = \mathbf{F}_{\text{rpl}} + \mathbf{F}_{\text{edg}} + \mathbf{F}_{\text{bdr}} + \mathbf{F}_{\text{ovr}} + \mathbf{F}_{\text{tor}} \quad (4.1)$$

922 where

- 923  $\mathbf{F}_{\text{rpl}}$  is the repulsion force;
- 924  $\mathbf{F}_{\text{edg}}$  is the edge force;
- 925  $\mathbf{F}_{\text{bdr}}$  is the boundary force;
- 926  $\mathbf{F}_{\text{ovr}}$  is the overlap force; and
- 927  $\mathbf{F}_{\text{tor}}$  is the torsional force.

928 These forces combine with the growth process, described in Section 4.5.1, to completely  
 929 fill the target container.

930 **Repulsion Force:** The repulsion force tries to push element meshes apart when they  
 931 approach each other, with the goal of making them transform to align their boundaries  
 932 (Figure 4.6a).

933 The vertex  $\mathbf{x}$  will experience an inverse square repulsive force, inspired by Coulomb's  
 934 law, from all nearby meshes. We use the following formula:

$$\mathbf{F}_{\text{rpl}} = k_{\text{rpl}} \sum_{i=1}^n \frac{\mathbf{u}}{\|\mathbf{u}\|} \frac{1}{\varsigma + \|\mathbf{u}\|^2} \quad (4.2)$$

935 where

- 936  $k_{\text{rpl}}$  is the strength of the repulsion force relative to other forces in the simulation;
- 937  $n$  is the number of nearest neighboring meshes to  $\mathbf{x}$ ;
- 938  $\mathbf{x}_i$  is the closest point on the skin of the  $i$ th neighbor;
- 939  $\mathbf{u} = \mathbf{x} - \mathbf{x}_i$ ; and
- 940  $\varsigma$  is a *soft parameter*; it places an upper bound on the magnitude of  $\mathbf{F}_{\text{rpl}}$ , avoiding  
 941 explosive instability when  $\|\mathbf{u}\|$  is very small.

942 An imbalance in the repulsion forces across a mesh's vertices will naturally induce  
 943 translation and rotation in meshes, helping their boundaries discover compatible segments  
 944 and consuming more of the remaining negative space.

945 If  $\mathbf{x}$  lies inside of another element's mesh, then the aggregate repulsion force from other  
 946 neighbors can push  $\mathbf{x}$  further inside and worsen the overlap. If we discover such an overlap,  
 947 we set  $\mathbf{F}_{\text{rpl}}$  to 0 and use the overlap force  $\mathbf{F}_{\text{ovr}}$ , discussed below, to correct it.

948 **Edge Force:** A mesh's edges are treated as longitudinal  springs; displacements of these  
 949 springs allow a mesh to deform in response to repulsion forces by neighboring meshes. An  
 950 undeformed element mesh provides the rest lengths for all of its springs; as mesh vertices  
 951 move relative to each other, the springs attempt to restore these rest lengths. Let  $\mathbf{x}_a$  and  
 952  $\mathbf{x}_b$  be mesh vertices connected by a spring. We compute the spring force as follows:

$$\mathbf{F}_{\text{edg}} = k_{\text{edg}} \frac{\mathbf{u}}{\|\mathbf{u}\|} s (\|\mathbf{u}\| - \ell)^2 \quad (4.3)$$

953 where

954  $k_{\text{edg}}$  is the strength of the edge force relative to other forces;  
 955  $\mathbf{u} = \mathbf{x}_b - \mathbf{x}_a$ ;  
 956  $\ell$  is the rest length of the spring; and  
 957  $s$  is +1 or -1, according to whether  $(\|\mathbf{u}\| - \ell)$  is positive or negative.

958 We apply  $\mathbf{F}_{\text{edg}}$  to  $\mathbf{x}_a$  and  $-\mathbf{F}_{\text{edg}}$  to  $\mathbf{x}_b$ . The equation is a modification of Hooke's law  
 959 in which the strength of the force increases quadratically with displacement. This change  
 960 allows the meshes to resist severe deformations when subjected to powerful forces. 

961 **Overlap Force:** Occasionally, a vertex  $\mathbf{x}$  from one mesh can be pushed inside the skin  
 962 of a neighboring mesh. In such cases, we temporarily disable the repulsion force on this  
 963 vertex by setting it to 0, and instead apply an overlap force that attempts to eject the  
 964 intruding vertex. In particular, every mesh triangle having  $\mathbf{x}$  as a vertex will pull  $\mathbf{x}$  in the  
 965 direction of its centroid. The overlap force is thus given by:

$$\mathbf{F}_{\text{ovr}} = k_{\text{ovr}} \sum_{i=1}^n (\mathbf{p}_i - \mathbf{x}) \quad (4.4)$$

966 where

967  $k_{\text{ovr}}$  is the relative strength of the overlap force;  
 968  $n$  is the number of mesh triangles that have  $\mathbf{x}$  as a vertex; and  
 969  $\mathbf{p}_i$  is the centroid of the  $i$ th triangle incident on  $\mathbf{x}$ .

970 The overlap force is zero for vertices that are not within another mesh.

971 **Boundary Force:** The boundary force causes element meshes to stay inside the target  
 972 container and conform to its boundary. It applies to any vertex  $\mathbf{x}$  that is outside the

973 container, and moves the vertex towards the closest point on the container's boundary, by  
 974 an amount proportional to the distance to the boundary:

$$\mathbf{F}_{\text{bdr}} = k_{\text{bdr}}(\mathbf{p}_b - \mathbf{x}) \quad (4.5)$$

975 where

976  $k_{\text{bdr}}$  is the relative strength of the boundary force; and  
 977  $\mathbf{p}_b$  is the closest point on the target container to  $\mathbf{x}$ .

978 The boundary force is zero for any point inside the container.

979 **Torsional Force:** As forces propagate through an element mesh, the aggregate velocity  
 980 vectors of the vertices can induce a rotation of the entire element. However, some elements  
 981 may have a preferred orientation, either for aesthetic reasons or because the shape is  
 982 comprehensible only at certain orientations. We introduce a torsional force that penalizes  
 983 individual vertices for which the orientation, relative to their element's center of mass,  
 984 drifts too far from its initial orientation.

985 Consider a vertex  $\mathbf{x}$  belonging to an element, and let  $\mathbf{c}_r$  be the element's center of mass  
 986 in its undeformed state. We may define the "rest orientation" of  $\mathbf{x}$  as the orientation of  
 987 the vector  $\mathbf{u}_r = \mathbf{x} - \mathbf{c}_r$ . During simulation we compute the current center of mass  $\mathbf{c}$  of the  
 988 element and let  $\mathbf{u} = \mathbf{x} - \mathbf{c}$ . Then the torsional force is

$$\mathbf{F}_{\text{tor}} = \begin{cases} k_{\text{tor}} \mathbf{u}^\perp, & \theta > 0 \\ -k_{\text{tor}} \mathbf{u}^\perp, & \theta < 0 \end{cases} \quad (4.6)$$

989 where

990  $k_{\text{tor}}$  is the relative strength of the torsional force;  
 991  $\theta$  is the signed angle between  $\mathbf{u}_r$  and  $\mathbf{u}$ ; and  
 992  $\mathbf{u}^\perp$  is a unit vector rotated 90° counterclockwise relative to  $\mathbf{u}$ ;

993 Using the equation above,  $\mathbf{F}_{\text{tor}}$  is always perpendicular to  $\mathbf{u}$  and the direction of  $\mathbf{F}_{\text{tor}}$   
 994 points to the left or to the right depending on  $\theta$ . Unlike the first four force types, the  
 995 torsional force is optional.

996 **Simulation Details:** We use explicit Euler integration to simulate the motions of the  
 997 mesh vertices under the forces described above. Every vertex has a position and a velocity  
 998 vector; in every iteration, we update velocities using forces, and update positions using  
 999 velocities. These updates are scaled by a simulation time step  $\Delta t_{\text{sim}}$ , typically chosen from  
 1000 the range [0.01, 0.1]. A smaller time step results in a more stable simulation at the cost

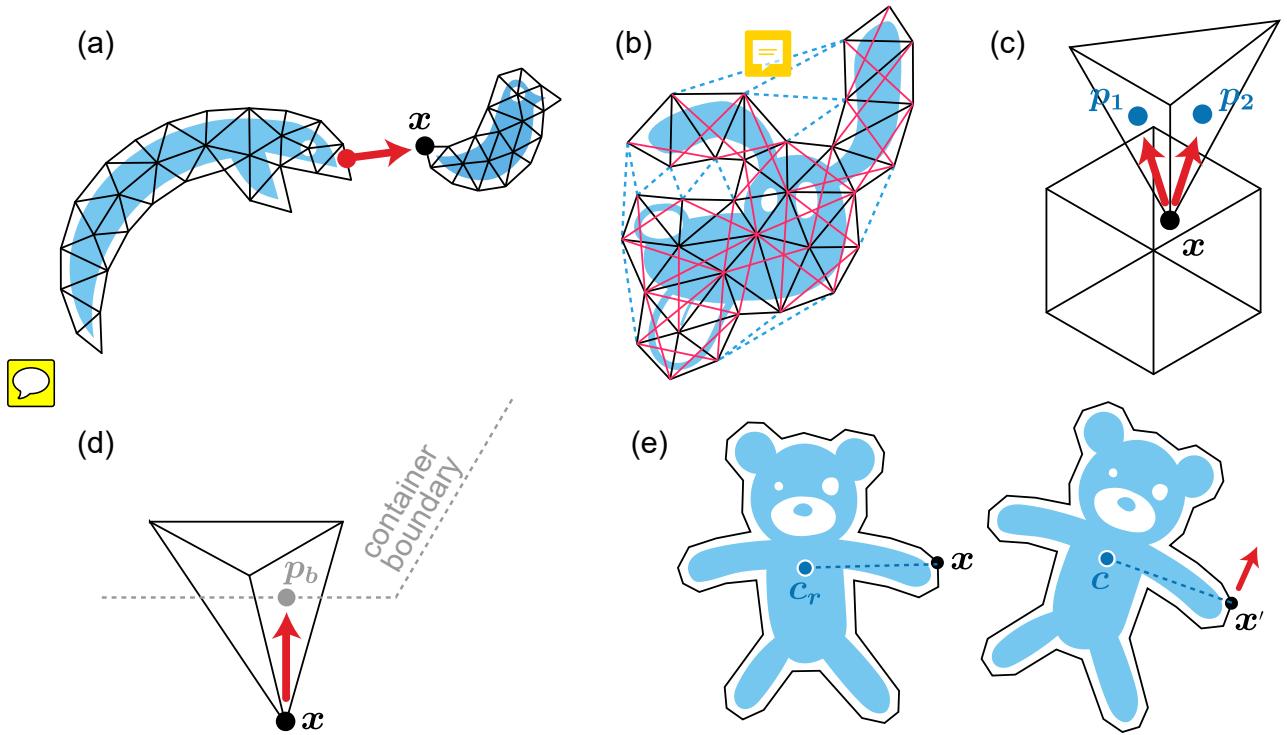


Figure 4.6: Illustrations of the forces in our simulation. **(a) Repulsion force:** The closest point on the snake’s mesh repels a vertex  $x$  in the bird mesh. **(b) Edge force:** We generate edge forces using **edge springs (black)**, **shear springs (red)**, and **negative-space springs (dashed blue)**. **(c) Overlap force:** Centers of triangles  $p_1$  and  $p_2$  attract a vertex  $x$  that lies in the interior of another mesh. **(d) Boundary force:** A vertex  $x$  moves toward  $p_b$ , the closest point on a target container, when it is outside the container. **(e) Torsional force:** We restrict the orientation of the element by generating a torsional force at every boundary vertex  $x$  to restore its angular position relative to  $c$ , the center of mass of the element.

1001 of additional running time. We cap velocities at  $5\Delta t_{\text{sim}}$  to dissipate extra energy from the  
1002 system.

1003 The repulsion and overlap forces rely on nearest-neighbor queries on the set of all  
1004 vertices. We accelerate these queries by storing vertices in a uniform spatial subdivision  
1005 grid that covers the container. In our implementation, cell width and height are 6 – 10%  
1006 of the larger dimension of the grid. We define the neighbors of a vertex  $\mathbf{x}$  as all vertices in  
1007 a  $3 \times 3$  window of cells centered on the cell containing  $\mathbf{x}$ . This approximation ignores the  
1008 negligible interactions between distant vertices.

1009 The constants  $k_{\text{rpl}}$ ,  $k_{\text{ovr}}$ ,  $k_{\text{bdr}}$ ,  $k_{\text{edg}}$ , and  $k_{\text{tor}}$  control the relative strengths of the five  
1010 forces in the simulation. They must also be chosen relative to the time step  $\Delta t_{\text{sim}}$  and the  
1011 overall width and height of the container. We find that our simulation produces satisfactory  
1012 results when  $k_{\text{rpl}} \approx k_{\text{ovr}} \approx k_{\text{bdr}} \geq k_{\text{edg}} > k_{\text{tor}}$ . For example, if the container’s bounding  
1013 box is approximately  $1000 \times 1000$ , we set  $k_{\text{rpl}} = k_{\text{ovr}} = k_{\text{bdr}} = 80$ , and  $k_{\text{tor}} = 1$ . We set  
1014  $k_{\text{edg}} = 40$  for edge springs and shear springs, and  $k_{\text{edg}} = 10$  for negative-space springs, since  
1015 weaker negative-space springs are sufficient to avoid self-intersections. We also set  $\varsigma = 1$   
1016 to avoid explosive repulsion forces. Increasing  $k_{\text{edg}}$  relative to the other forces suppresses  
1017 deformation, yielding a close approximation of packing with rigid elements.

#### 1018 4.5.1 Element Growth

1019 RepulsionPak starts with small initial elements to avoid intersections, and gradually en-  
1020 larges them until they tightly fill the target container. Figure 4.3c-d shows elements growing  
1021 and gradually consuming negative space. Elements have different intrinsic sizes, which are  
1022 respected in the initial placement. Because they all grow at roughly the same rate, their  
1023 relative sizes tend to be maintained.

1024 After each iteration of the physics simulation, the element meshes undergo a growth  
1025 step. If an element mesh has no vertices that lie inside of neighboring meshes, it is permitted  
1026 to grow in this iteration. A mesh with overlaps may still grow in subsequent iterations, if  
1027 local changes to the packing open up more negative space. This approach produces slight  
1028 variations in skin offsets in the output packing but the effect is negligible.

1029 We implement growth in the context of the physics simulation by scaling the rest lengths  
1030 of a mesh’s springs, allowing it to expand as the simulation progresses. Every mesh  $M$  has  
1031 a counter  $n_M$  that records the number of times it has grown. When a mesh is permitted  
1032 to grow, we add 1 to the counter. Then, if  $L_i$  is the rest length of the  $i$ th spring in the  
1033 original undeformed mesh, we increase its rest length to  $(1 + n_M k_g \Delta t_{\text{sim}}) L_i$ . The constant  
1034  $k_g$ , usually 0.01 in our system, controls the growth rate.

1035 **4.5.2 Stopping Criteria**

1036 We choose one of two criteria to stop the simulation. First, the artist specifies the desired  
1037 positive space ratio at which the simulation immediately terminates. The ratio should  
1038 not be set too high, and we find that most packings have positive space ratios between  
1039 45% and 65% depending on the element shapes. For example, concave elements with long  
1040 extensions are more difficult to pack, so a lower positive space ratio is recommended to  
1041 generate a satisfactory packing. Second, we stop the simulation when the element meshes  
1042 are no longer able to maneuver enough to consume the remaining negative space. After  
1043 each iteration, we compute an *area fraction*  $A$ , defined to be the fraction of the container  
1044 area taken up by element meshes. We then compute a measurement of the recent change  
1045 in area fraction in a sliding window that covers the  $w$  most recent iterations of the system;  
1046 we use  $w = 100$ . If  $A_0, \dots, A_w$  are the area fractions in the  $w + 1$  iterations up to the  
1047 current one, then we define

$$\text{RMS} = \sqrt{\frac{1}{w} \sum_{i=1}^w (A_i - \bar{A})^2} \quad (4.7)$$

1048 We stop iterating when  $\text{RMS} < \epsilon$ , where  $\epsilon$  is 0.01 in our system.

1049 **4.6 Secondary Elements**

1050 The iteration process described above can leave behind isolated pockets of empty space,  
1051 which will be visible in the final composition. We imitate the approach taken by human  
1052 artists by filling these pockets with small, usually simple secondary elements.

1053 We seed the container with secondary elements by finding points that are far from any  
1054 existing element mesh. Specifically, we compute a discrete approximation of the distance  
1055 transform of the negative space. We then create an initial candidate list of all points  
1056 for which the distance value is above a threshold  $d_{\min}$ , sorted by decreasing distance.  
1057 We consider each of these candidates in turn, adding it to a final list of seed locations  
1058 provided that no previously chosen seed is within distance  $d_{\text{sep}}$  of the candidate. In our  
1059 implementation, if the distance transform is computed on a  $1000 \times 1000$  grid fit to the  
1060 container's bounds, then we typically set  $5 \leq d_{\min} \leq 10$  and  $d_{\text{sep}} = 10$ .

1061 Next, we assign random secondary elements to these chosen seed points, scaled down as  
1062 before to avoid overlaps. We then run the simulation and growth process again, but freeze

1063 the primary elements: they exert repulsion forces on secondary elements and can induce  
 1064 overlaps, but primary mesh vertices cannot move. The secondary elements gradually grow  
 1065 to consume some of the remaining negative space until the packing satisfies the same  
 1066 stopping criteria described above. 

## 1067 4.7 Shape Matching

1068 The motions and deformations of elements, as described in the previous sections, give them  
 1069 an opportunity to conform to each other and to the target container. However, in some  
 1070 cases the random seeding may position some elements in such a way that the simulation  
 1071 process will still leave undesirable artifacts. In particular, when a round element is placed  
 1072 near a sharp convex corner of the container, it cannot deform enough to extend into the  
 1073 corner, but cannot yield its position to a pointy element that offers a better fit. In the  
 1074 final packing, the sharp corners of the containers will appear “rounded off”. Another  
 1075 problem occurs when a long narrow element is initially placed diagonally across a corner,  
 1076 in which case the simulation pushes the element’s middle into the corner, causing extreme  
 1077 deformation. 

1078 To overcome these deficiencies, we optionally perform an initial fit-guided placement  
 1079 pass before seeding the rest of the container at random. Here we take inspiration from  
 1080 existing rigid packing algorithms [KSH<sup>+</sup>16], which are driven primarily by shape compat-  
 1081 ibility. We use a simplified shape descriptor; we can tolerate a less perfect initial fit, with  
 1082 the expectation that deformation will improve the quality later.

1083 We begin by building shape descriptors for the elements. Each element is first scaled to  
 1084 have area  $0.6A_c/n_e$ , where  $A_c$  is the container area and  $n_e$  is the number of elements. This  
 1085 step resizes the element to a rough estimate of its final size in the packing, an approximation  
 1086 that is adequate for our fit-based placement. We then sample the target container and  
 1087 the boundaries of the elements, with adjacent samples separated by distance  $\delta$ . We set  
 1088  $\delta = 0.002L_s$ , where  $L_s$  is the side length of the container’s bounding square.

1089 We define a local descriptor based on integral of absolute curvature [CFH<sup>+</sup>09, KSH<sup>+</sup>16].  
 1090 Let  $P(t)$  be an arclength-parameterized 2D curve, and let  $\kappa(t)$  represent the curvature of  
 1091 the curve at  $P(t)$ . For a given interval  $[s, t]$  within the curve’s domain, we may define **the**  
 1092 **integral of absolute curvature**:

$$\tau(s, t) = \int_s^t |\kappa(x)| dx. \quad (4.8)$$

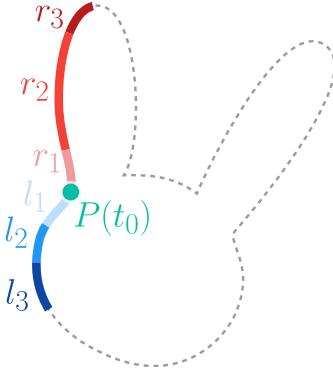


Figure 4.7: An illustration of a local shape descriptor with  $n = 3$ . These segments have varying arclengths but they all have the same value of  $\tau$ , the integral of absolute curvature along their lengths.

1093 In practice, we estimate curvature from the discrete sample points using second-order  
 1094 forward finite differences [BL15] and compute  $\tau$  by summing curvature estimates using the  
 1095 trapezoid rule.

1096 Let an objective  $\tau_{\text{obj}}$  be a positive real number and  $P(t_0)$  be a given point on a curve.  
 1097 If  $\tau_{\text{obj}}$  is sufficiently small, then as we traverse the curve on either side of  $P(t_0)$  we will  
 1098 eventually reach points  $P(t_{-1})$  and  $P(t_1)$  such that  $\tau(t_{-1}, t_0) = \tau(t_0, t_1) = \tau_{\text{obj}}$ . We let  
 1099  $l_1 = t_0 - t_{-1}$  and  $r_1 = t_1 - t_0$  be the arclengths that produce these integrals. We can  
 1100 continue this process for any number of steps, walking in both directions along the curve  
 1101 away from previous sample points until we reach  $\tau_{\text{obj}}$ , yielding new arclengths  $l_k$  and  $r_k$   
 1102 (see Figure 4.7). Finally, for a given number of steps  $n$  we define the shape descriptor at  
 1103  $P(t_0)$  as

$$(l_n, l_{n-1}, \dots, l_2, l_1, r_1, r_2, \dots, r_{n-1}, r_n) \quad (4.9)$$

1104 Like PAD [KSH<sup>+</sup>16], our shape descriptor is rotation invariant. Effectively it is one level  
 1105 of a PAD, which suffices because we do not require scale invariance. Descriptors can be  
 1106 compared using simple Euclidean distance, accelerated by storing them in a k-D tree. In  
 1107 our implementation we set  $\tau_{\text{obj}} = 0.001L_s$ ; the dependence on  $L_s$  makes the measurement  
 1108 robust against changes in absolute container size, and the factor of 0.001 was determined  
 1109 through experimentation. We further choose  $n = 5$ , yielding a 10-dimensional descriptor.  
 1110 We compute this descriptor for all the container and element samples defined above.

1111 Based on these descriptors, we now use a simple greedy heuristic to identify salient  
 1112 container features where shape matching will be used. Here we restrict our attention  
 1113 to convex protrusions with high curvature, which benefit the most from careful element

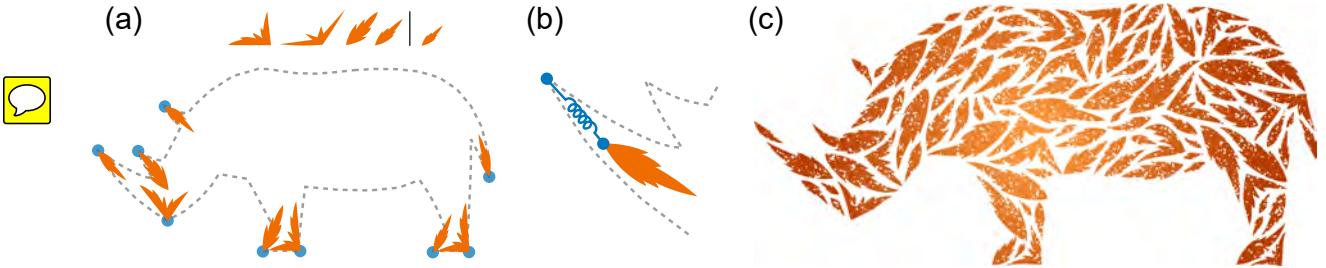


Figure 4.8: A demonstration of shape matching of leaf shapes inside a rhinoceros. (a) We detect nine salient features, namely sharp convex corners, and assign an element to each. (b) A spring holds each element in place. (c) The final result. 

placement (See Figure 4.8a). Given a shape descriptor, we define its total length to be  $\sum_{i=1}^n l_i + r_i$ . We iterate over all sample points on the container boundary in increasing order by the total lengths of their descriptors. For each sample point  $P(t_0)$ , we add it to a list of salient features under two conditions. First, we require that the angle formed by  $P(t_0)$  with two samples to either side of it be sufficiently acute; we use a threshold of 0.3 radians to ensure convex-to-convex matching. Second, we ensure that salient features are not too close to each other by requiring that every new sample point be separated by a distance of at least  $0.2L_s$  from all previously identified salient features.

Even when an element's descriptor is a close match to a segment of target container, it may still not be safe to place that element at a given location. For example, one part of an element may extend into a corner of the container, while a different part, outside the purview of the local shape descriptor, could protrude outside the container entirely. We augment our descriptor-based fit calculation with an additional score adapted from Gal et al. [GSP<sup>+</sup>07] in order to ensure a more global element fit. Let  $d(\mathbf{x})$  be a signed distance function for the container, with negative values inside the container and positive outside. In practice we superimpose a  $1000 \times 1000$  grid over the container's bounding square and compute a discrete approximation of the distance transform. For an element sample point  $\mathbf{x}$ , we then compute a score

$$q(\mathbf{x}) = \begin{cases} e^{-d(\mathbf{x})^2/\mu^2\beta}, & \text{if } d(\mathbf{x}) < 0 \\ 1, & \text{if } d(\mathbf{x}) = 0 \\ -\alpha d(\mathbf{x}), & \text{if } d(\mathbf{x}) > 0 \end{cases} \quad (4.10)$$

With this scoring function, a sample point that lies on the container boundary will have a score of 1, the maximum for  $q(\mathbf{x})$ . Scores decay exponentially towards 0 for points

1134 farther inside the container. Sample points outside the container are penalized by assigning  
1135 a negative score proportional to distance.

1136 We then define a quality measurement for an entire element by summing over all of its  
1137 sample points:

$$Q = \sum_{i=1}^{n_p} q(\mathbf{x}_i) \quad (4.11)$$

1138 The parameter  $\alpha$  penalizes elements with parts that protrude outside the container;  $\beta$  favors  
1139 elements with parts close to the target container;  $\mu = 0.5\sqrt{2}L_s$  is half of the diagonal of  
1140 the container's bounding square; and  $n_p$  is the number of sample points on the element. In  
1141 our implementation we choose  $\alpha = 1$  and  $\beta = 0.001$ . For every salient container feature,  
1142 we obtain the 10 closest descriptors from the k-D tree, and choose one that has the highest  
1143  $Q$ . Finally, we place the selected element in the container by aligning the endpoints of the  
1144 element's and container's matching shape descriptors.

1145 The shape matching process yields a set of elements that should be attached to par-  
1146 ticular locations on the target container. However, during simulation they could wander  
1147 away from these initial positions, under the influence of the many other forces at play. As  
1148 shown in Figure 4.8b, we encourage these elements to remain in place by attaching them  
1149 to the salient container points with additional springs.

## 1150 4.8 Implementation and Results

1151 Our software was written in C++, and reads in text files describing elements and containers;  
1152 we prepared these files using Adobe Illustrator. We ran our software on a computer with a  
1153 2.4 GHz Intel i7-4700HQ processor and 16 GB of RAM. As a post-process, we optionally  
1154 read packings back into Illustrator, fit smooth curves to polygonal paths, and apply colors  
1155 and other visual effects. Table 4.1 shows statistics for the results [in this chapter](#). All results  
1156 in this [chapter](#) use  $\Delta t_{\text{sim}} = 0.1$ . Torsional forces are used only in Figure 4.10, and shape  
1157 matching is used only in Figure 4.8.

1158 The supplemental materials include movies that visualize the simulation process<sup>1</sup>.   
1159 These movies make it clear that elements can jostle each other around, inducing trans-  
1160 lation, rotation, and deformation throughout the simulation.

1161 The packing in Figure 4.1 uses six cat-shaped primary elements and one secondary cat  
1162 head. RepulsionPak naturally bends legs and tails to fill the container more evenly. 

---

<sup>1</sup>[https://cs.uwaterloo.ca/~radhitya/repulsionpak/supplemental\\_material.zip](https://cs.uwaterloo.ca/~radhitya/repulsionpak/supplemental_material.zip)

Table 4.1: Data and statistics for the RepulsionPak results. The table shows the numbers of primary and secondary elements ( $n_p$ ,  $n_s$ ), the number of vertices ( $v$ ), the running times of the primary and the secondary simulations ( $t_p$ ,  $t_s$ ) in seconds, and the number of iterations ( $i$ ). In these results, torsional forces are used only in Figure 4.10 and shape matching is used only in Figure 4.8.

Packing	$n_p$	$n_s$	$v$	$t_p$	$t_s$	$i$	💡
Cats (Figure 4.1)	41	69	3598	185	62	8531	
Animals (Figure 4.3)	25	14	2412	133	65	16670	
Rhino (Figure 4.8)	107	0	4833	237	0	15521	
Birds (Figure 4.9a)	43	43	2309	102	54	11571	
Bats (Figure 4.9b)	47	22	3048	165	56	13120	
Butterflies (Figure 4.9c)	123	135	11916	696	616	14379	
Giraffes & Penguins (Figure 4.10)	60	0	2250	163	0	9347	
Paisley (Figure 4.11)	162	0	4860	128	0	23040	
Circles in Paisleys (Figure 4.11)	144	0	2544	403	0	29584	

1163     The animal packing in Figure 4.3 has several elements with limbs (the bear, fox, chick,  
 1164     and penguin), extensions (the dog and bunny ears), and long shapes (the snake). Figure 4.4  
 1165     highlights the deformations for some of these elements; they are noticeably more deformed  
 1166     than nearly convex elements like the cat and mouse.

1167     The butterfly packing in Figure 4.9 is an attempt to reproduce the visual style of a  
 1168     dense packing (or tessellation), similar to Jigsaw Image Mosaics [KP02] or the “Butterflies  
 1169     in Butterfly” example from the Pyramid of Arclength Descriptor paper [KSH<sup>+</sup>16]. The  
 1170     target container is made from two regions, one with internal holes. The resulting packing  
 1171     is tight but overlap-free.

1172     The bird packing in Figure 4.9 exhibits significant deformation in the wings and the  
 1173     tails of the birds. In particular, the thin tails of the swallows have some unaesthetic sharp  
 1174     bends. We would like to investigate ways to ensure these bends are smoother.

1175     The packing in Figure 4.10 demonstrates torsional forces that gradually turn the ele-  
 1176     ments upside down. The rhinoceros packing in Figure 4.8 demonstrates initial placement  
 1177     via shape matching. The entire shape matching process took 922 milliseconds with a  
 1178     4-element library.

1179     We have also experimented with creating tileable packings, as shown in Figure 4.11.  
 1180     We seed a central square with elements, but also place clones of those elements in all the  
 1181     squares of the 8-neighborhood around the center. These clones track the transformations

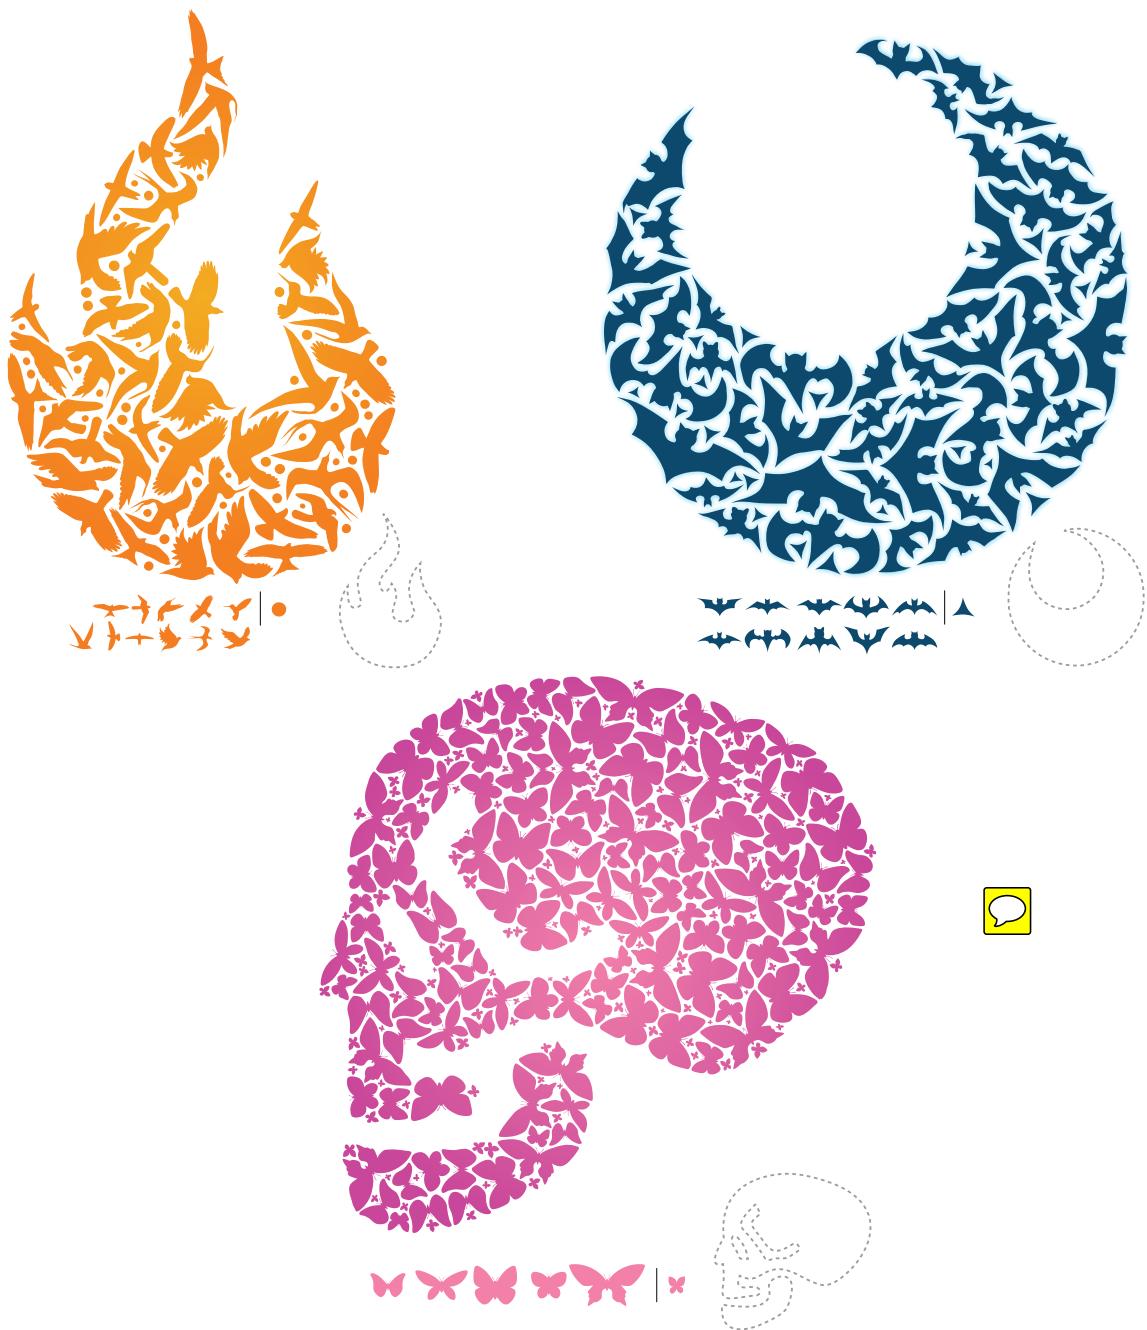


Figure 4.9: Three packings created using RepulsionPak: Birds, Bats, and Butterflies. The results are visually appealing overall, though some birds' tails suffer from excessive deformation in the packing on the left.

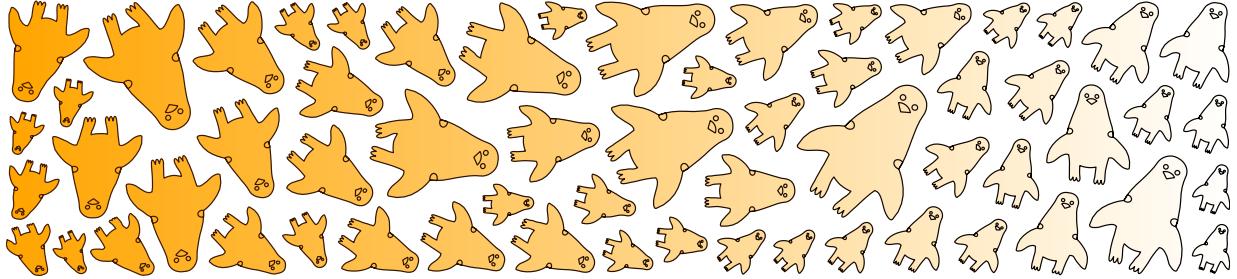


Figure 4.10: A packing that demonstrates torsional forces. The packing uses copies of a single element shape, but every copy is given a rest orientation between  $0^\circ$  and  $180^\circ$ , based on its horizontal position in the container. In the final packing the elements transition from upright to upside-down, recreating an illusion in which giraffe heads become penguins.



1182 and deformations of their originals, but also exert forces on them during simulation, leading  
1183 to an even, seamless packing in a toroidal domain.

1184 We have found that RepulsionPak is robust to parameter variation, and produces pre-  
1185 dictable, high quality results without the need for fine-grained adjustments. However, as  
1186 shown in Figure 4.12, extreme parameter settings can still produce degenerate results. In  
1187 Figure 4.12a, the repulsion force is made much stronger than the edge force, leading to  
1188 excessive element deformation and self-intersections in the pursuit of even negative space.  
1189 In Figure 4.12b, edge and torsional forces dominate repulsion, producing a packing with  
1190 stiff, upright elements that do not fill their container effectively.

## 1191 4.9 Conclusions

1192 We presented RepulsionPak, a method to create packings with deformable elements. Each  
1193 element is represented as a mass-spring system, allowing it to deform to achieve a better fit  
1194 with its neighbors and the container. The combination of repulsion forces and controlled  
1195 deformation allows RepulsionPak to create shape compatibilities that eliminate the need  
1196 for a large element library and fill the target container effectively.

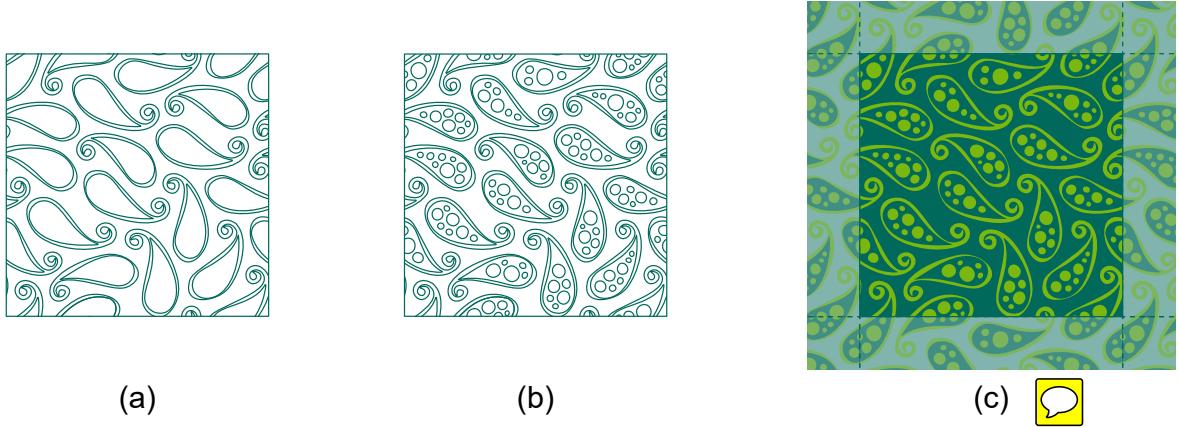


Figure 4.11: A paisley-inspired toroidal packing that can tile the plane. (a) An initial paisley packing. (b) In a separate simulation, we fill each paisley with circles to demonstrate a packing inside a packing. (c) The final result.

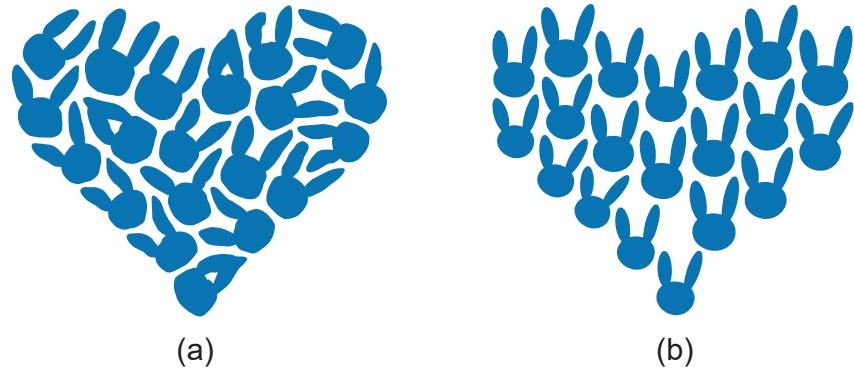


Figure 4.12: Two **examples** of how extreme parameter values can lead to low-quality results. In (a), we allow repulsion to overwhelm element shape by setting  $k_{\text{rpl}}$  to 200 and  $k_{\text{edg}}$  to 20; the resulting packing has even negative space, but elements suffer from high deformation and **self-intersections**. In (b), we minimize repulsion and prioritize orientation by setting  $k_{\text{edg}}$ ,  $k_{\text{tor}}$ , and  $k_{\text{rpl}}$  to 200, 100, and 50, respectively. The elements deviate minimally from their original shapes and orientations, but cannot fill the container effectively.

<sub>1197</sub> Chapter 5

<sub>1198</sub> **AnimationPak: Packing Elements  
with Scripted Animations**

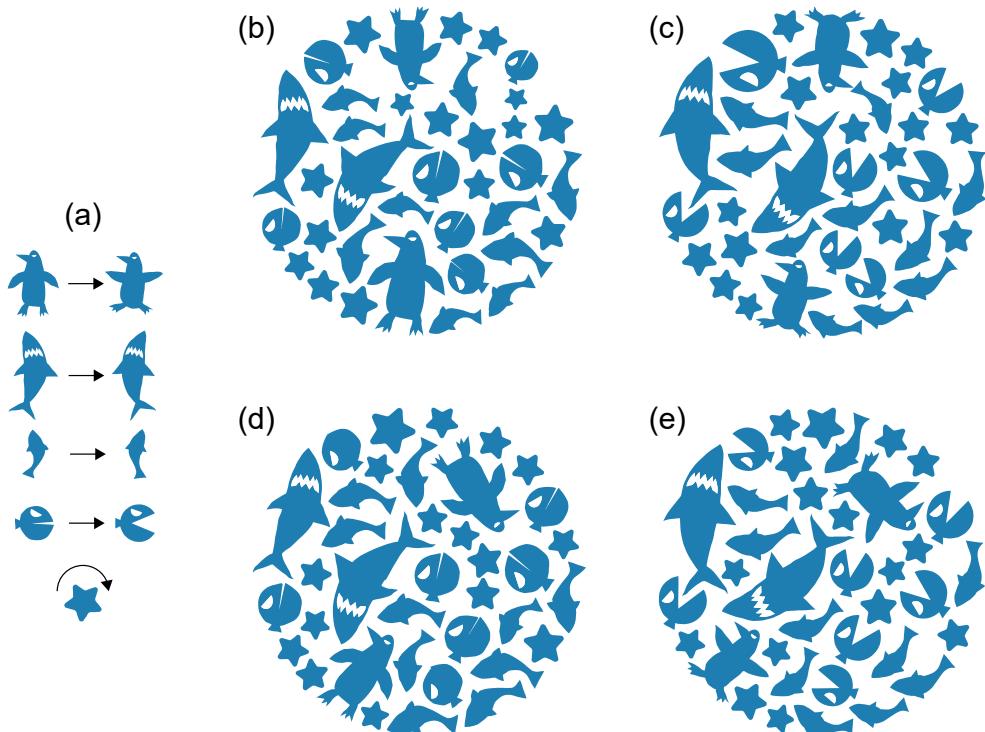


Figure 5.1: (a) Input animated elements, each with its own animation: swimming penguins, swimming sharks and fish, Pac-Man fish that open ~~or~~ close their mouths, and rotating stars. (b-e) four selected frames from an animated packing.

1200 **5.1 Introduction**



1201 AnimationPak is a method to pack elements with scripted animations. An element can  
1202 have an animated deformation, such as a slithering snake or a dancing bear. It can also  
1203 have an animated transformation, giving a changing position, size, and orientation within  
1204 the container. Our goal is producing an *animated packing*, with elements playing out their  
1205 animations while simultaneously filling the container shape evenly. A successful animated  
1206 packing should balance among the evenness of the negative space, the preservation of  
1207 element shapes, and the comprehensibility of their scripted animations.

1208 We consider an animated element to be a geometric extrusion along a time axis, a  
1209 three-dimensional object that we call a “spacetime element”. We use a three-dimensional  
1210 physical simulation similar to RepulsionPak to pack spacetime elements into a volume  
1211 created by extruding a static container shape. The animated packing emerges from this  
1212 three-dimensional volume by rendering cross sections perpendicular to the time axis.

1213 Animated packings are a largely unexplored style of motion graphics, presumably be-  
1214 cause of the difficulty of creating an animated packing by hand. Finding motivating exam-  
1215 ples created by artists is difficult, and we only found a single animated packing shown in  
1216 Figure 5.2, which is animated gear-like meshing of copies of Lisa Simpson’s spiky head. The  
1217 discussion in Section 5.2 shows there is also limited past research on animated packings.

1218 **5.2 Related Work**



1219 **Packings and Mosaics:** Researchers have explored many approaches to creating 2D pack-  
1220 ings and simulated mosaics, including using Centroidal Area Voronoi Diagrams (CAVDs)  
1221 to position elements [Hau01, HHD03, SLK05], spectral approaches to create even negative  
1222 space [DKLS06], energy minimization [KP02], and shape descriptors [KSH<sup>+</sup>16].

1223 **Animated Packings and Tilings:** Animosaics by Smith et al. [SLK05] constructs  
1224 animations in which static elements without scripted animations follow the motion of an  
1225 animated container. Elements are placed using CAVDs, and advected frame-to-frame  
1226 using a choice of methods motivated by Gestalt grouping principles. As the container’s  
1227 area changes, elements are added and removed as needed, while attempting to maximize  
1228 overall temporal coherence. Dalal et al. [DKLS06] showed how the spectral approach  
1229 they introduced for 2D packings could be extended to pack animated elements in a static  
1230 container. Like us, they recast the problem in terms of three-dimensional spacetime; they  
1231 computed optimal element placement using discrete samples over time and orientation.

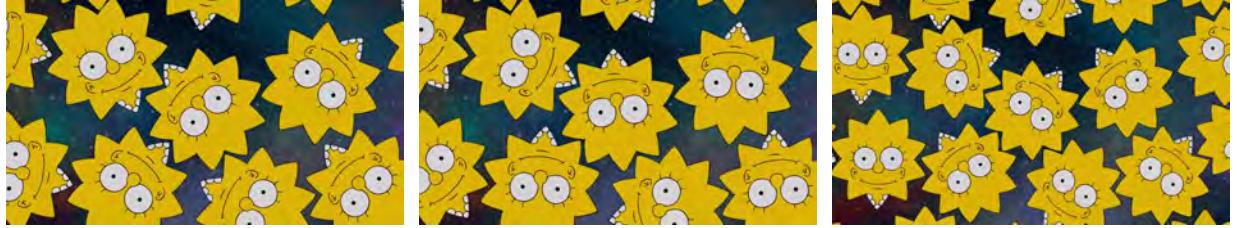


Figure 5.2: Three frames of an animated packing from The Simpsons Season 31, Episode 19 (© FOX), showing gear-like meshing of oppositely-rotating copies of Lisa Simpson.

1232 However, their spacetime elements have fixed shapes and are made to fit together using  
 1233 only translation and rotation, limiting their ability to consume the container's negative  
 1234 space.

1235 Liu and Veksler created animated decorative mosaics from video input [LV09]. Their  
 1236 technique combines vision-based motion segmentation with a packing step similar to Ani-  
 1237 mosaics. Kang et al. [KOHY11] extracted edges from video and then oriented rectangular  
 1238 tesserae relative to edge directions.

1239 Kaplan [Kap19] explored animations of simple tilings of the plane from copies of a single  
 1240 shape. Elements in a tiling fit together by construction, and therefore always consume all  
 1241 the negative space in the animation.

1242 **3D Packings:** AnimationPak fills a 3D container with 3D elements, and is therefore  
 1243 related to approaches that construct 3D packings, although the time axis in the spacetime  
 1244 domain behaves differently than the third spatial axis. Related work includes 3D collages  
 1245 methods [GSP<sup>+</sup>07, TRdAS07, HWFL14], cutting and packing methods [CSR10, VGB<sup>+</sup>14,  
 1246 CZL<sup>+</sup>15], non-overlapping 3D element packing methods [BTW09, MCHW18], and packing  
 1247 methods for engineering purposes [CSY02].

1248 **Animated 3D Collages:** Theobalt et al. [TRdAS07] developed a data-driven method  
 1249 to generate animated 3D collages that resemble an Arcimboldo's paintings by arranging  
 1250 static 3D elements into an animated 3D container. They partitioned the input animated  
 1251 container by identifying rigid parts, each is replaced by a matching element taken from a  
 1252 library. The result is a 3D collage with overlapping elements, and the animation is created  
 1253 from rigid transformations of the elements.

1254 **Animated Textures:** Ma et al. [MWLT13] developed Dynamic Element Textures  
 1255 (DET), a method to synthesize an animated texture of a group of identical elements, such  
 1256 as stars or particles. For each animated element in an exemplar, DET generates sample  
 1257 points both in spatial and temporal domains. DET synthesizes a larger texture by copying



1258 and distributing the sample points, which are then repositioned to resemble the original  
1259 distribution in the exemplar. Similar to DET, AnimationPak represents an animated  
1260 element as a collection of discrete vertices that reside in a spacetime domain. However,  
1261 as a texture synthesis method, the goal of DET is to create an element arrangement that  
1262 resembles an exemplar, and not fill a container tightly.

1263 **Derived Animations:** AnimationPak falls into the category of systems that create a  
1264 derived animation based on some input animation. This problem, which requires preserving  
1265 the visual character of the input, is a longstanding one in computer graphics research.  
1266 Spacetime constraints [WK88, Coh92] allow an animator to specify an object’s constraints  
1267 and goals, and then calculates the object’s trajectory via spacetime optimization. Motion  
1268 warping [WP95] is a method that deforms an existing motion curve to meet user-specified  
1269 constraints. Bruderlin and Williams [BW95] used signal processing techniques to modify  
1270 motion curves. Gleicher [Gle01] developed a motion path editing method that allows user  
1271 to modify the traveling path of a walking character.

1272 Previous work has also investigated geometric deformation of animations. Ho et al.  
1273 [HKT10] encoded spatial joint relationships using tetrahedral meshes, and applied as-  
1274 rigid-as-possible shape deformation to the mesh to retarget animation to new characters.  
1275 Choi et al. [CKHL11] developed a  method to deform character motion to allow characters  
1276 to navigate tight passages. Masaki [Osh17] developed a motion editing tool that deformed  
1277 3D lattice proxies of a character’s joints. Dalstein et al. [DRvdP15] presented a data struc-  
1278 ture to animate vector graphics with complex topological changes. Kim et al. [KHHL12]  
1279 explored a packing algorithm to avoid collisions in a crowd of moving characters. They  
1280 defined a motion patch containing temporal trajectories of interacting characters, and ar-  
1281 ranged deformed patches to prevent collisions between characters.

## 1282 5.3 Animated Elements

1283 The input to AnimationPak is a library of animated elements and a fixed container shape.  
1284 AnimationPak currently supports two kinds of animation: the user can animate the shape  
1285 of each individual element and can also give elements trajectories that animate their pos-  
1286 ition within the container. This section explains how we animate the element shapes using  
1287 as-rigid-as-possible deformation, and then construct spacetime-extruded objects that form  
1288 the basis of our packing algorithm. These elements animate “in place”: they change shape  
1289 without translating. The next section describes how these elements can be given ~~trans-~~  
1290 ~~formation~~ trajectories within the container. Size and orientation of an element can be

1291 animated either way; they can be specified as an animation of the element’s shape, or they  
1292 can be part of the transformation trajectory.

### 1293 5.3.1 Spacetime Extrusion

1294 Each element begins life as a static shape defined using vector paths. Similar to RepulsionPak,  
1295 we construct a discrete geometric proxy of the element that will interact with other proxies  
1296 in a physical simulation. The construction of this proxy for a single shape is shown in  
1297 Figure 5.3, and the individual steps are explained in greater detail below.

1298 In order to produce a packing with  even distribution of negative space, we first  
1299 offset the shape’s paths by a distance  $\Delta s$ , leaving the shape surrounded by a channel of  
1300 negative space (Figure 5.3a). In our system we scale the shape to fit a unit square and set  
1301  $\Delta s = 0.04$ .

1302 Next, we place evenly-spaced samples around the outer boundary of the offset path and  
1303 construct a Delaunay triangulation of the samples (Figure 5.3b). We will later treat the  
1304 edges of the triangulation as springs, allowing the element to deform in response to forces  
1305 in the simulation. We also follow RepulsionPak by adding shear edges to prevent folding  
1306 and negative space edges to avoid self-intersections during simulation (Figure 5.3c).

1307 We refer to the augmented triangulation shown in Figure 5.3c as a *slice*. The entire  
1308 spacetime packing process operates on slices. However, we will eventually need to compute  
1309 deformed copies of the element’s original vector paths when rendering a final animation  
1310 (Section 5.6). To that end, we re-express all path information relative to the slice trian-  
1311 gulation: every path control point is represented using barycentric coordinates within one  
1312 triangle.

1313 To extend the element into the time dimension, we now position evenly-spaced copies  
1314 of the slice along the time axis. Assuming that the animation will run over the time  
1315 interval  $[0, 1]$ , we choose a number of slices  $n_s$  and place slices  $\{s_1, \dots, s_{n_s}\}$ , with slice  $s_i$   
1316 being placed at time  $(i - 1)/(n_s - 1)$ . Higher temporal resolution will produce a smoother  
1317 final animation at the expense of more computation. In our examples, we set  $n_s = 100$ .  
1318 Figure 5.3d shows a set of time slices, with  $n_s = 5$  for visualization purposes.

1319 To complete the construction of a spacetime element without animation, we stitch the  
1320 slices together into a single 3D object. Let  $s_j$  and  $s_{j+1}$  be consecutive slices constructed  
1321 above. The outer boundaries of the element triangulations are congruent polygons offset in  
1322 the time axis. We stitch the two polygons together using a new set of *time edges*: if  $AB$  is  
1323 an edge on the boundary of  $s_j$  and  $CD$  is the corresponding edge on the boundary of  $s_{j+1}$ ,

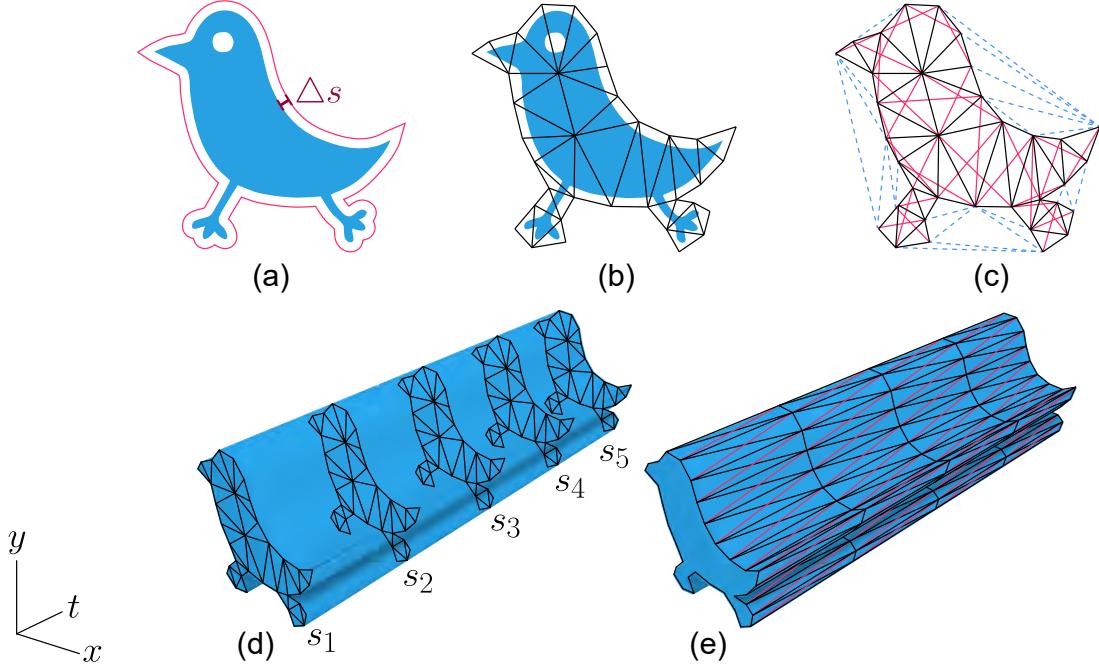


Figure 5.3: The creation of a discretized spacetime element. (a) A 2D element shape offset by  $\Delta s$ . (b) A single triangle mesh slice. (c) Shear edges (red) and negative space edges (blue). (d) A set of five slices placed along the time axis. (e) The vertices on the boundaries of the slices are joined by time edges. The black edges in (e) define a triangle mesh called the envelope of the element. In practice we use a larger number of slices in (d) and (e).

1324 then we add time edges  $AC$ ,  $AD$ , and  $BC$ . During simulation, time edges will transmit  
 1325 forces backwards and forwards in time, maintaining temporal coherence by smoothing out  
 1326 deformation and transformations. Figure 5.3e shows time edges for  $n_s = 5$ .

### 1327 5.3.2 Animation

1328 The 3D information constructed above is a parallel extrusion of a slice along the time axis,  
 1329 representing a shape with no scripted animation. We created a simple interactive applica-  
 1330 tion for adding animation to spacetime elements, inspired by as-rigid-as-possible shape  
 1331 manipulation [IMH05]. The artist first designates a subset of the slices as keyframes.  
 1332 They can then interactively manipulate any triangulation vertex of a keyframe slice. Any  
 1333 vertex that has been positioned manually has its entire trajectory through the animation

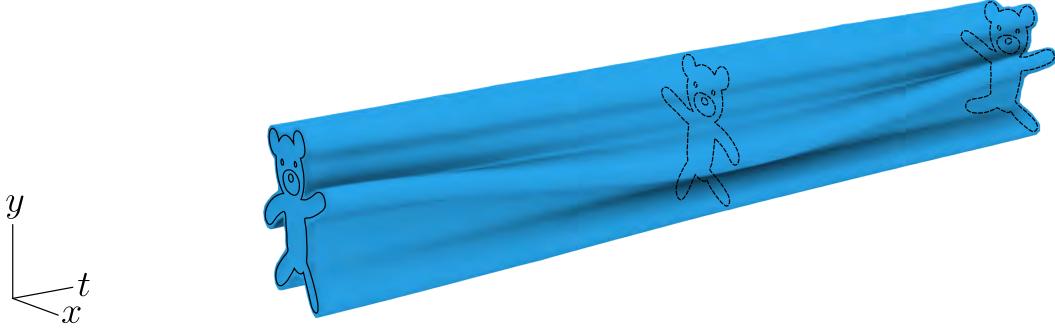


Figure 5.4: A spacetime element with a scripted animation.

1334 computed using spline interpolation. Then, at any other slice, the positions of all other vertices can be interpolated using the as-rigid-as-possible technique. The result is a smoothly  
 1335 animated spacetime volume like the one visualized in Figure 5.4. 

1337 Unlike data-driven packing methods, AnimationPak allows distortions so it does not  
 1338 require a large library of distinct elements to generate successful packings. The results in  
 1339 this [chapter](#) all use fewer than ten input elements, and some use only one. The physical  
 1340 simulation induces deformation to enhance the compatibility of nearby shapes in the final  
 1341 animation.

## 1342 5.4 Initial Configuration

1343 We begin the packing process by constructing a 3D spacetime volume for the container  
 1344 by extruding its static shape in the time direction. The container is permitted to have  
 1345 internal holes, which are also extruded. The resulting volume is scaled to fit a unit cube.  
 1346 We also shrink each of the spacetime elements, in the spatial dimensions only, to 5–10% of  
 1347 its original size. These shrunken elements are thin enough that we can place them in the  
 1348 container without overlaps. 

1349 The artist can optionally specify trajectories for a subset of the elements, which we call  
 1350 *guided elements*. A guided element attempts to pass through a sequence of fixed target  
 1351 points in the container, imbuing the animation with a degree of intention and narrative  
 1352 structure. To define a guided element, we designate the triangulation vertex closest to its  
 1353 centroid to be the anchor point for the element. The artist then chooses a set of spacetime  
 1354 target points  $\mathbf{p}_1, \dots, \mathbf{p}_n$ , with  $\mathbf{p}_i = (x_i, y_i, t_i)$ , that the anchor should pass through during  
 1355 the animation. In our interface, the artist uses a slider to choose the time  $t_i$  for a target

1356 point, and clicks in the container to specify the spatial position  $(x_i, y_i)$ . The artist can also  
1357 optionally specify scale and orientation at the target points. We require  $t_1 = 0$  and  $t_n = 1$ ,  
1358 fixing the initial and final positions of the guided element. We then linearly interpolate the  
1359 anchor position for each slice based on the target points, and translate the slice so that its  
1360 anchor lies at the desired position. The red extrusions in Figure 5.6a are guided elements.

1361 If the artist wishes to create a looping animation, the  $(x_i, y_i)$  position for target points  $\mathbf{p}_1$   
1362 and  $\mathbf{p}_n$  must match up, either for a single guided element or across elements. In Figure 5.6  
1363 the two guided elements form a connected loop;  $(x_1, y_1)$  for each one matches  $(x_n, y_n)$  for  
1364 the other.

1365 In this initial configuration, the guided elements abruptly change direction at target  
1366 points. However, because the slices are connected by springs, the trajectories will smooth  
1367 out as the simulation runs. Also, the simulation is not constrained to reach each target  
1368 position exactly. Instead, we attach the anchor to the target using a *target-point spring*  
1369 that attempts to draw the element towards it while balancing against the other physical  
1370 forces in play (Figure 5.6b). The strength of these springs determines how closely the  
1371 element will follow the trajectory.

1372 We then seed the container with an initial packing of non-guided spacetime elements.  
1373 We generate points within the container at random, using blue-noise sampling [Bri07] to  
1374 prevent points from being too close together, and assign a spacetime element to each seed  
1375 point, selecting elements randomly from the input library. Depending upon the desired  
1376 effect, we either randomize their orientations or give them preferred orientations. We reject  
1377 any candidate seed point that would cause an unguided element's volume to intersect a  
1378 guided element's volume.

1379 Finally we shrink each element, guided and unguided, uniformly in the spatial dimen-  
1380 sion towards its centroid. These shrunken elements are guaranteed not to intersect one  
1381 another; as the simulation runs, they will grow and consume the container's negative space,  
1382 while avoiding collisions. The blue extrusions in Figure 5.6a show an initial placement of  
1383 spacetime elements.

## 1384 5.5 Simulation

1385 We perform a physics simulation on the spacetime elements and the container. Elements  
1386 are subjected to a number of forces that cause them to simultaneously grow, deform, and  
1387 repel each other (Figure 5.6). In Section 5.5.2 we introduce some new hard constraints  
1388 that must be applied after every time step. Note that we must distinguish two notions of

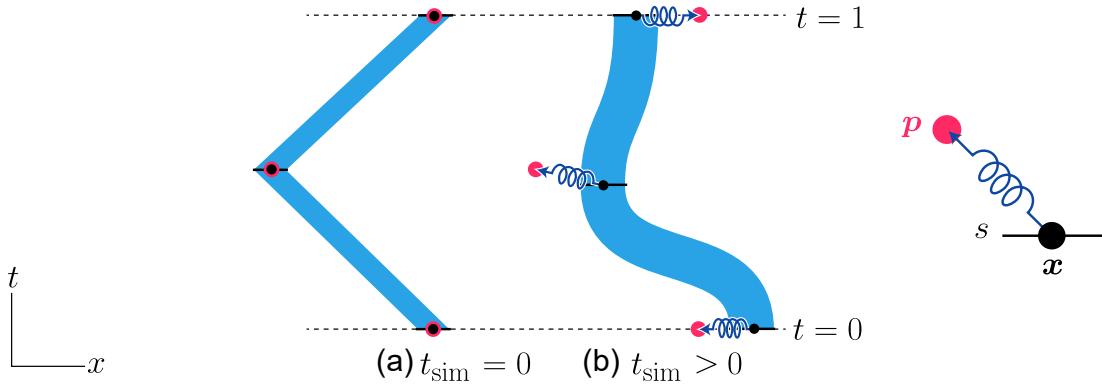


Figure 5.5: A 2D illustration of a guided element. Slices are depicted as black lines and slice vertices as black dots. A spring connects the centermost vertex  $x$  of a slice  $s$  to a target point  $p$ . (a) The initial shape of a guided element is a polygonal extrusion. (b) The spacetime element deforms but the springs pull it back towards the target points.

time in this simulation. We use  $t$  to refer to the time axis of our spacetime volume, which will become the time dimension of the final animation, and  $t_{\text{sim}}$  to refer to the time domain in which the simulation is taking place.

Let  $\mathbf{x} = (x, y, t)$  be a vertex of a slice. The total force  $\mathbf{F}_{\text{total}}$  applied to  $\mathbf{x}$  is

$$\mathbf{F}_{\text{total}} = \mathbf{F}_{\text{rpl}} + \mathbf{F}_{\text{edg}} + \mathbf{F}_{\text{bdr}} + \mathbf{F}_{\text{oov}} + \mathbf{F}_{\text{tor}} + \mathbf{F}_{\text{tmp}} \quad (5.1)$$

where

$\mathbf{F}_{\text{rpl}}$  is the repulsion force;

$\mathbf{F}_{\text{edg}}$  is the edge force;

$\mathbf{F}_{\text{bdr}}$  is the boundary force;

$\mathbf{F}_{\text{oov}}$  is the overlap force;

$\mathbf{F}_{\text{tor}}$  is the torsional force; and

$\mathbf{F}_{\text{tmp}}$  is the temporal force.

These forces are the spacetime analogues of the ones used in RepulsionPak, except the new temporal force. In this section, we explain the first five forces briefly and readers can refer back to Section 4.5 for equations used to generate them.

**Repulsion forces** allow elements to push away vertices of neighboring elements, inducing deformations and transformations that lead to an even distribution of elements within the container (Figure 5.7). Since the simulation operates in the spacetime domain, vertex  $\mathbf{x}$  accumulates repulsion forces from points at various time positions. To locate

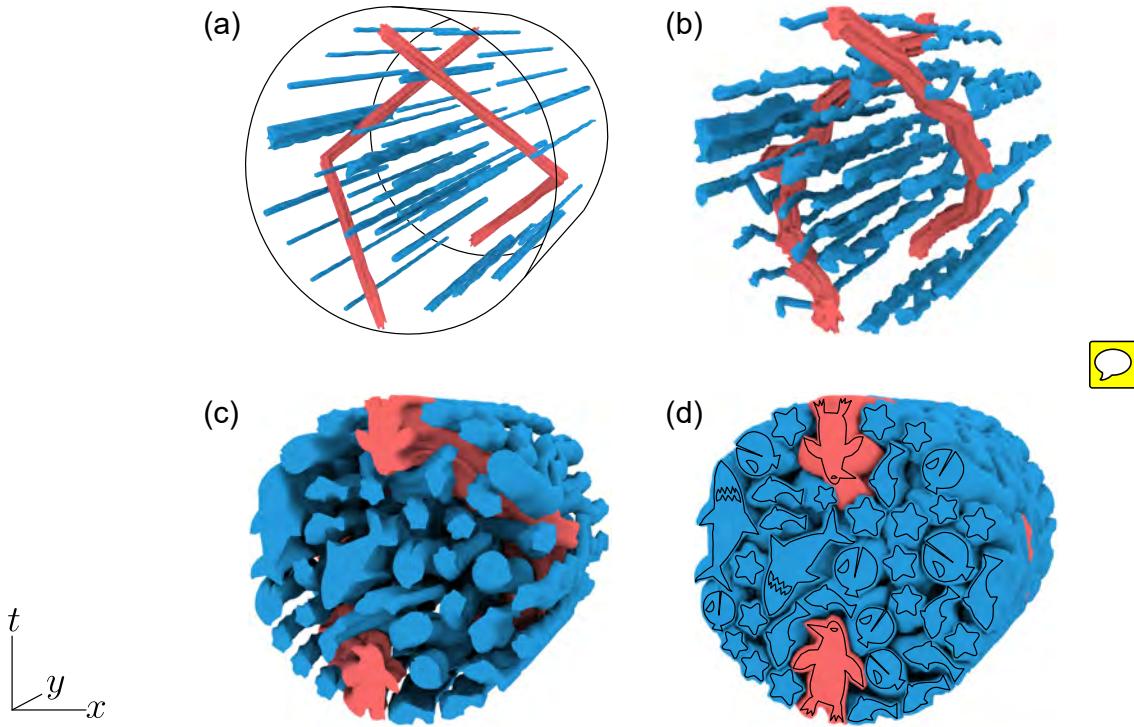


Figure 5.6: The simulation process. (a) Initial placement of shrunken spacetime elements inside a static 2D disc, extruded into a cylindrical spacetime domain. Guided elements are shown in red and unguided elements in blue. (b) A physics simulation causes the spacetime elements to bend. They also grow gradually. (c) The spacetime elements occupy the container space. (d) The simulation stops when elements do not have sufficient negative space in which to grow, or have reached their target sizes.

1407 these points on neighboring elements that are considered nearest, we use a collision grid  
1408 data structure, described in greater detail in Section 5.5.1.

1409 **Edge forces** allow elements to deform in response to repulsion forces. As discussed in  
1410 Section 5.3.1 and Section 5.4, we have five types of springs: edge springs, shear springs,  
1411 negative-space springs, time springs, and target-point springs. Each of the spring type has  
1412 a different relative strength.

1413 **Overlap forces** resolve a vertex penetrating a neighboring spacetime element. Over-  
1414 laps can occur later in the simulation when negative space is limited. Once we detect a  
1415 penetration, we temporarily disable the repulsion force on vertex  $\mathbf{x}$ , and apply an overlap  
1416 force  $\mathbf{F}_{ovr}$  to push it out.

1417 **Boundary forces** keep vertices inside the container. If an element vertex  $\mathbf{x}$  is outside

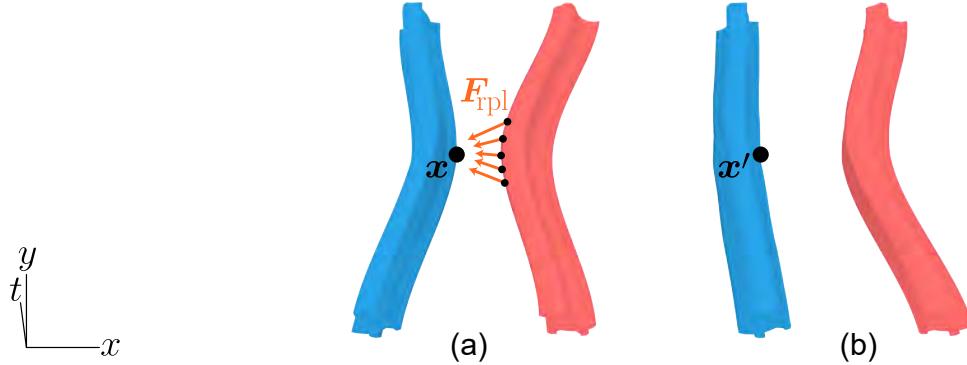


Figure 5.7: Repulsion forces applied to a vertex  $x$ , allowing the element to deform and move away from a neighboring element.

1418 the container, the boundary force  $\mathbf{F}_{\text{bdr}}$  moves it towards the closest point on the container's  
 1419 boundary by an amount proportional to the distance to the boundary.

1420 **Torsional forces** allow an element's slices to be given preferred orientations, to which  
 1421 they attempt to return.

1422 **Temporal forces** prevent slices from drifting too far from their original positions  
 1423 along the time axis **positions** (Figure 5.8), which could cause unexpected accelerations and  
 1424 decelerations in the final animation. For every vertex, we compute the temporal force  $\mathbf{F}_{\text{tmp}}$   
 1425 as

$$\mathbf{F}_{\text{tmp}} = k_{\text{tmp}} \mathbf{u}^t (t - t') \quad (5.2)$$

1426 where

1427  $k_{\text{tmp}}$  is the relative strength of the temporal force;  
 1428  $t$  is the initial time of the slice to which the vertex belongs;  
 1429  $t'$  is the current time value of the vertex; and  
 1430  $\mathbf{u}^t = (0, 0, 1)$ .

1431 **Simulation Details:** We use explicit Euler integration to simulate the motions of the  
 1432 mesh vertices under the forces described above. Every vertex has a position and a velocity  
 1433 vector; in every iteration, we update velocities using forces, and update positions using  
 1434 velocities. These updates are scaled by a time step  $\Delta t_{\text{sim}}$  that we set to 0.01. We cap  
 1435 velocities at  $10\Delta t_{\text{sim}}$  to dissipate extra energy from the simulation.

1436 The constants  $k_{\text{rpl}}$ ,  $k_{\text{ovr}}$ ,  $k_{\text{bdr}}$ ,  $k_{\text{edg}}$ ,  $k_{\text{tor}}$  and  $k_{\text{tmp}}$  control the relative strengths of repul-  
 1437 sion, overlap, boundary, edge, torsional, and temporal forces, respectively. They must be  
 1438 adjusted relative to the time step  $\Delta t_{\text{sim}}$  and the size of the container. Since container's

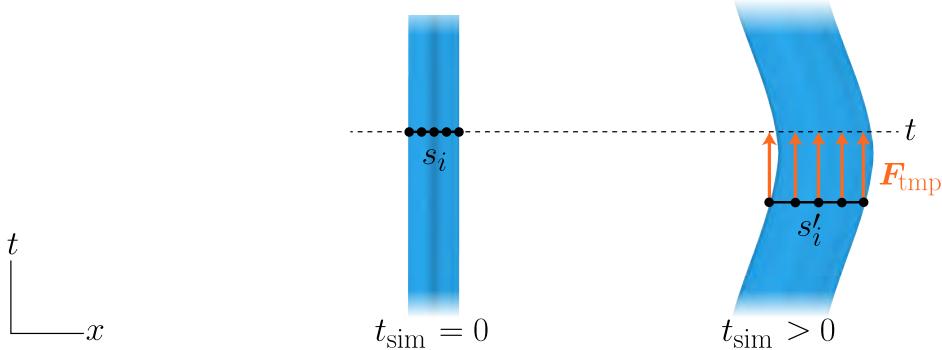


Figure 5.8: An illustration of the temporal force. The vertices in slice  $s_i$  are drawn back towards time  $t$ .

1439 bounding box is resized to a unit cube, we set  $k_{\text{rpl}} = 10$ ,  $k_{\text{ovr}} = k_{\text{bdr}} = 5$ ,  $k_{\text{tor}} = k_{\text{tmp}} = 1$ .  
1440 We set  $k_{\text{edg}} = 0.01$  for time springs,  $k_{\text{edg}} = 0.1$  for negative-space springs, and  $k_{\text{edg}} = 10$   
1441 for edge springs, shear springs, and target-point springs.

### 1442 5.5.1 Spatial Queries

1443 Repulsion and overlap forces rely on being able to find points on neighboring elements that  
1444 are close to a given query vertex. To find these points, we use each element's *envelope*, a  
1445 triangle mesh implied by the construction in Section 5.3.1. Each triangle of the envelope  
1446 is made from two time edges and one edge of a slice boundary, as shown in Figure 5.9a.  
1447 Given a query vertex  $\mathbf{x}$ , we need to find nearby envelope triangles that belong to other  
1448 elements.

1449 To accelerate this computation, we first compute and store the centroids of every element's envelope triangles in a uniformly subdivided 3D grid that surrounds the spacetime  
1450 volume of the animation. In using this data structure, we make two simplifying assumptions;  
1451 first, that because envelope triangles are small, their centroids are adequate for  
1452 finding triangles near a given query point; and second, that the repulsion force from a  
1453 more distant triangle is well approximated by a force from its centroid.

1455 Given a query vertex  $\mathbf{x}$ , we first find all envelope triangle centroids in nearby grid  
1456 cells that belong to other elements. For each centroid, we use a method described by  
1457 Ericson [Eri05] to find the point on its triangle closest to  $\mathbf{x}$  and include that point in the  
1458 list of points in Eq. (1). These nearby triangles will also be used to test for interpenetration  
1459 of elements. We then find centroids in more distant grid cells, and add those centroids

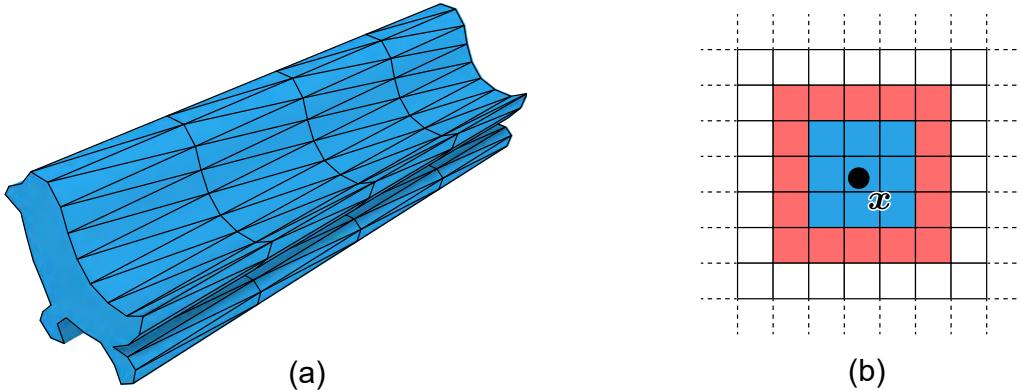


Figure 5.9: (a) The triangles that connect consecutive slices define the envelope of the element. The midpoints of these triangles are stored in a collision grid. (b) A 2D visualization of the region of collision grid cells around a query point  $x$  in which repulsion and overlap forces will be computed. In the central blue region, we check overlaps and compute exact repulsion forces relative to closest points on triangles of neighboring elements; in the peripheral red region we do not compute overlaps, and repulsion forces are approximated using triangle midpoints only.

directly to the Eq. (1) list, skipping the closest point computation. In our system we set the cell size to 0.04, giving a  $25 \times 25 \times 25$  grid around the simulation volume. A query point's nearby grid cells are the 27 cells making up a  $3 \times 3 \times 3$  block around the cell containing the point; the more distant cells are the 98 that make up the outer shell of the  $5 \times 5 \times 5$  block around that (Figure 5.9).

### 5.5.2 Slice Constraints

There are three hard geometric constraints on the configuration of slices, which must be enforced throughout the simulation. Each of the following constraints is reapplied after each physical simulation step described above.

1. **End-To-End Constraint:** A spacetime element must be present for the full length of the animation from  $t = 0$  to  $t = 1$ . After every simulation step, every vertex belonging to an element's first slice has its  $t$  value set to 0, and every vertex of the last slice has its  $t$  value set to 1 (Figure 5.10a). 
2. **Simultaneity Constraint:** During simulation, the vertices of a slice can drift away from each other in time, which could lead to rendering artifacts in the animation.

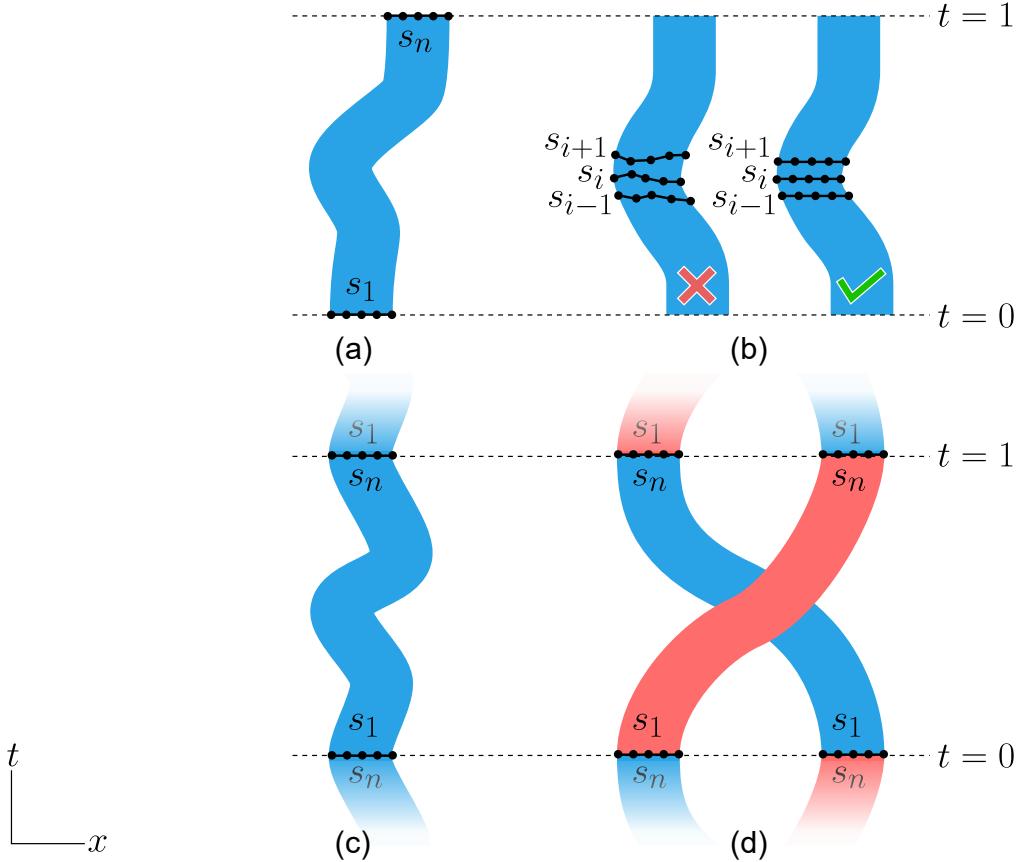


Figure 5.10: a) End-to-end constraint: slice  $s_1$  and  $s_n$ , located at  $t = 0$  and  $t = 1$ , should never change their  $t$  positions but can change their  $x, y$  positions. b) Simultaneity constraint: all vertices on the same slice should have the same  $t$  position. c) Loop constraint with a single element: the  $x, y$  positions for  $s_1$  and  $s_n$  must match. d) Loop constraint with two elements: the  $x, y$  position for  $s_1$  for one element matches the  $x, y$  position for  $s_n$  of the other.

After every simulation step, we compute the average  $t$  value of all vertices belonging to each slice other than the first and last slices, and snap all the slice's vertices to that  $t$  value (Figure 5.10b).

3. **Loop Constraint:** AnimationPak optionally supports looping animations. When looping is enabled, we must ensure that the  $t = 0$  and  $t = 1$  planes of the spacetime container are identical. The  $t = 1$  slice of every element  $e_1$  must then coincide with the  $t = 0$  slice of *some* element  $e_2$ . We can have  $e_1 = e_2$  (Figure 5.10c), but more general loops are possible in which the elements arrive at a permutation of their

1483 original configuration (Figure 5.10d). We require only that there is a one-to-one  
1484 correspondence between the vertices of the  $t = 1$  slice of  $e_1$  and the  $t = 0$  slice of  
1485  $e_2$ . If  $\mathbf{p}_1 = (x_1, y_1, 1) \in e_1$  and  $\mathbf{p}_2 = (x_2, y_2, 0) \in e_2$  are in correspondence, then after  
1486 every simulation step we move  $\mathbf{p}_1$  to  $(\frac{x_1+x_2}{2}, \frac{y_1+y_2}{2}, 1)$  and  $\mathbf{p}_2$  to  $(\frac{x_1+x_2}{2}, \frac{y_1+y_2}{2}, 0)$ .



### 1487 5.5.3 Element Growth and Stopping Criteria

1488 We begin the spacetime packing process with all element slices scaled down in  $x$  and  $y$ ,  
1489 guaranteeing that elements do not overlap. As the simulation progresses we gradually grow  
1490 the slices, consuming the negative space around them (Figure 5.11a,b). A perfect packing  
1491 would fill the spacetime container completely with the elements. Because each element  
1492 wraps the underlying animated shape with a narrow channel of negative space, this would  
1493 yield an even distribution of shapes in the resulting animation. For real-world elements, the  
1494 goal of minimizing deformation of irregular element shapes will lead to imperfect packings  
1495 with additional pockets of negative space.

1496 **Element Growth:** We induce elements to grow spatially by gradually increasing  
1497 the rest lengths of their springs. The initial rest length of each spring is determined  
1498 by the vertex positions in the shrunken version of the spacetime element constructed in  
1499 Section 5.4. We allow an element's slices to grow independently of each other, which  
1500 complicates the calculation of new rest lengths for time springs. Therefore, we create a  
1501 duplicate of every shrunken spacetime element in the container, with a straight extrusion  
1502 for unguided elements, and a polygonal extrusion for guided elements. This duplicate is  
1503 not part of the simulation; it serves as a reference. Every element slice maintains a current  
1504 scaling factor  $g$ . When we wish to grow the slice, we increase its  $g$  value. We can compute  
1505 new rest lengths for all springs by scaling every slice of the reference element by a factor  
1506 of  $g$  relative to the slice's centroid, and measuring distances between the scaled vertex  
1507 positions. These new rest lengths are then used as the  $\ell$  values in Equation 4.3 to calculate  
1508 edge forces.

1509 Every element slice has its  $g$  value initialized to 1. After every simulation step, if none  
1510 of the slice's vertices were found to overlap other elements we increase that slice's  $g$  by  
1511  $0.001\Delta t_{\text{sim}}$ , where  $\Delta t_{\text{sim}}$  is the simulation time step. If any overlaps are found, then that  
1512 slice's growth is instead paused to allow overlap and repulsion forces to give it more room  
1513 to grow in later iterations. This approach can cause elements to fluctuate in size during  
1514 the course of an animation, as slices compete for shifting negative space (Figure 5.11).



1515 **Stopping Criteria:** We halt the simulation when the space between neighboring  
1516 elements drops below a threshold. When calculating repulsion forces, we find the distance

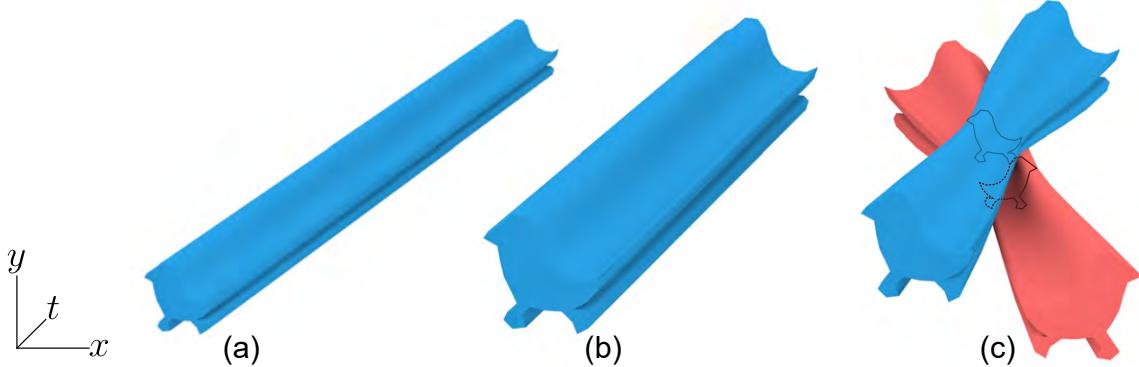


Figure 5.11: A spacetime element shown (a) shrunken at the beginning of the simulation, and (b) grown later in the simulation. (c) When two elements overlap somewhere along their lengths, they are temporarily prohibited from growing there.

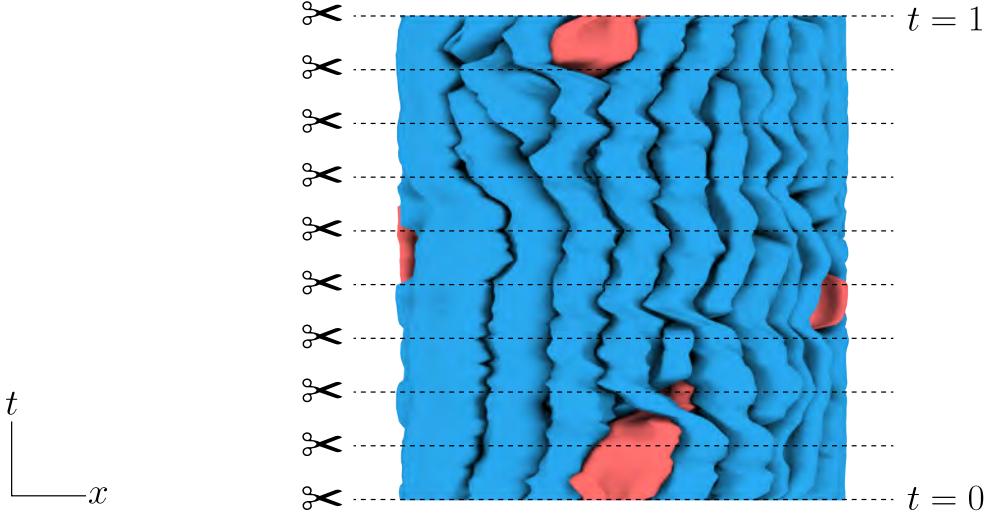
from every slice vertex to the closest point in a neighboring element. The minimum of these distances over all vertices in an element slice determines that slice's closest distance to neighboring elements. We halt the simulation when the maximum per-slice distance falls below 0.006 (relative to a normalized container size of 1). That is, we stop when every slice is touching (or nearly touching) at least one other element.

In some cases it can be useful to stop early based on cumulative element growth. In that case, we set a separate threshold for the slice scaling factors  $g$  described above, and stop when the  $g$  values of all slices exceed that threshold.

## 5.6 Rendering

The result of the simulation described previously is a packing of spacetime elements within a spacetime container. We can render an animation frame-by-frame by cutting through this volume at evenly spaced  $t$  values from  $t = 0$  to  $t = 1$ , as shown in Figure 5.12. For our results, we typically render 500-frame animations.

During simulation, a given spacetime element's slices may drift from their original creation times. However, time springs keep the sequence monotonic, and the simultaneity constraint ensures that every slice is fixed to one  $t$  value. To render this element at an arbitrary frame time  $t_f \in [0, 1]$ , we find the two consecutive slices whose time values bound the interval containing  $t_f$  and linearly interpolate the vertex positions of the triangulations at those two slices to obtain a new triangulation at  $t_f$ . We can then compute a deformed



**Figure 5.12:** An illustration of rendering spacetime elements to generate 2D frames by cutting them at evenly spaced time values from  $t = 0$  to  $t = 1$ .

1536 copy of the original element paths by “replaying” the barycentric coordinates computed  
 1537 in Section 5.3.1 relative to the displaced triangulation vertices. We repeat this process for  
 1538 every spacetime element to obtain a rendering of the frame at  $t_f$ .

1539 This interpolation process can occasionally lead to small artifacts in the animation. A  
 1540 rendered frame can fall between the discretely sampled slices for two elements at an inter-  
 1541 mediate time where physical forces were not computed explicitly. It is therefore possible  
 1542 for neighboring elements to overlap briefly during such intervals.

## 1543 5.7 Implementation and Results

1544 The core AnimationPak algorithm consists of a C++ program that reads in text files de-  
 1545 scribing the spacetime elements and the container, and outputs raster images of animation  
 1546 frames.

1547 Large parts of AnimationPak can benefit from parallelism. In our implementation we  
 1548 update the cells of the collision grid (Section 5.5.1) in parallel by distributing them across  
 1549 a pool of threads. When the updated collision grid is ready, we distribute the spacetime  
 1550 elements over threads. We calculate forces, perform numerical integration, and apply the  
 1551 end-to-end and simultaneity constraints for each element in parallel. We must process any  
 1552 loop constraints afterwards, as they can affect vertices in two separate elements.

Table 5.1: Data and statistics for the AnimationPak results. The table shows the number of elements, the number of vertices, the number of springs, the number of envelope triangles, and the running time of the simulation in hours, minutes, and seconds.

Packing	Elements	Vertices	Springs	Triangles	Time
Aquatic fauna (Figure 5.1)	37	97,800	623,634	106,000	01:06:35
Snake and bird (Figure 5.13)	37	58,700	370,571	58,700	01:01:32
Penguin to giraffe (Figure 5.14)	33	124,300	824,164	143,000	01:19:50
Traveling bird (Figure 5.15)	32	63,700	389,373	70,100	00:19:24
Lion (Figure 5.16b)	16	39,400	236,086	41,800	00:41:56
Animals (Figure 5.17b)	34	69,600	444,337	69,800	01:00:19
Heart stars (Figure 5.18c)	26	85,200	598,218	85,800	00:23:08

1553 We created the results in this chapter using a Windows PC with a 3.60 GHz Intel i7-  
 1554 4790 processor and 16 GB of RAM. We used a pool of eight threads, corresponding to the  
 1555 number of logical CPU cores. Table 5.1 shows statistics for our results. Each packing has  
 1556 tens of thousands of vertices and hundreds of thousands of springs, and requires about an  
 1557 hour to complete. We enable the loop constraint in all results. This chapter shows selected  
 1558 frames from the results; see supplemental videos for full animations<sup>1</sup>. 

1559 Figure 5.1 is an animation of aquatic fauna featuring two penguins as guided elements.  
 1560 During one loop the penguins move clockwise around the container, swapping positions at  
 1561 the top and the bottom. Each ends at the other's starting point, demonstrating a loop  
 1562 constraint between distinct elements. All elements are animated, as shown in Figure 5.1a.  
 1563 Note the coupling between the Pac-Man fish's mouth and the shark's tail on the left side  
 1564 of the second and fourth frames.

1565 A snake chases a bird around an annular container in Figure 5.13, demonstrating a  
 1566 container with a hole and giving a simple example of the narrative potential of animated  
 1567 packings. Figure 5.14 animates the giraffe-to-penguin illusion shown as a static packing in

---

<sup>1</sup><https://cs.uwaterloo.ca/~radhitya/animationpak/videos.zip>

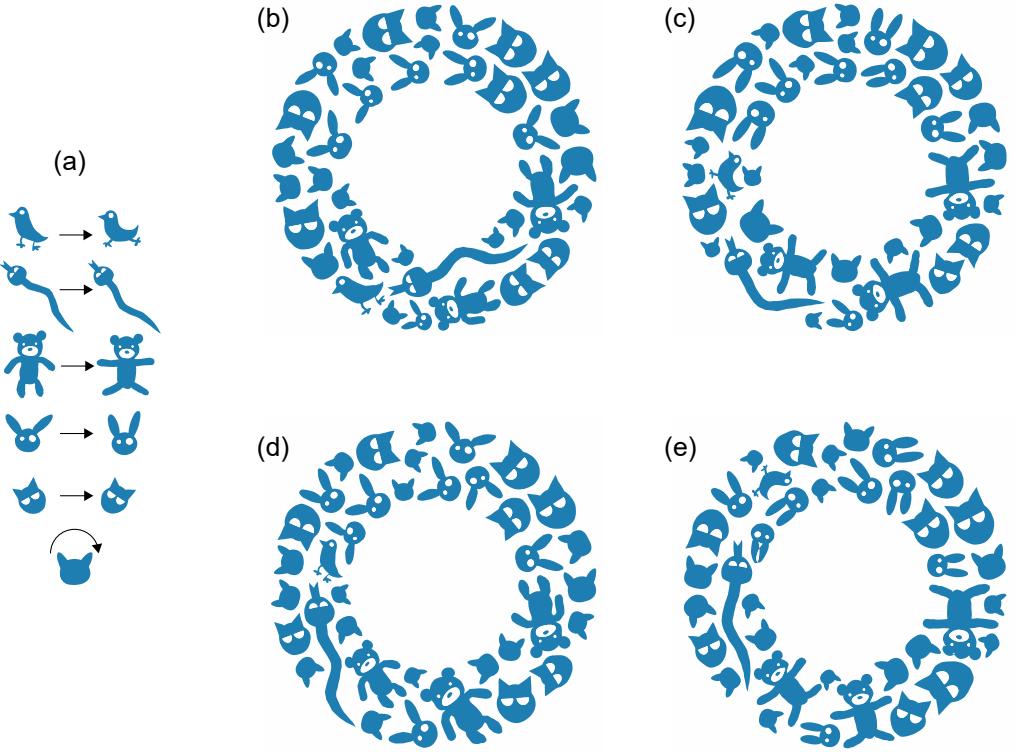


Figure 5.13: A snake chasing a bird through a packing of animals. The snake and bird are both guided elements that move clockwise around the annular container.

1568 RepulsionPak This example uses torsional forces to control slice orientations.

1569 ~~Figure 5.15 shows a guided element as a traveling bird which is not constricted inside~~  
 1570 ~~a container.~~ The traveling bird moves from left to right, passing through a packing of  
 1571 birds. As some slices of the traveling bird taking up container space, the other birds have  
 1572 to stop some of their slice growths prematurely. This is a demonstration of the benefit  
 1573 of the non-uniform slice growths so that elements can adapt to changes in container area.  
 1574 In the rendered 2D animation, the elements have a subtle shrinking-and-expanding effect.  
 1575 The composition can also tile horizontally, which can be seen in the supplemental video.

1576 Figure 5.16a is a static packing of lion's mane created by an artist and used as an  
 1577 example in FLOWPAK (Chapter 3). In Figure 5.16b, we reproduce it with animated ele-  
 1578 ments for the mane. The orientations of elements follow a vector field inside the container,  
 1579 and are maintained during the animation by torsional forces. We simulate only half of the  
 1580 packing and reflect it to create the other half. The facial features were added manually in  
 1581 a post-processing step.

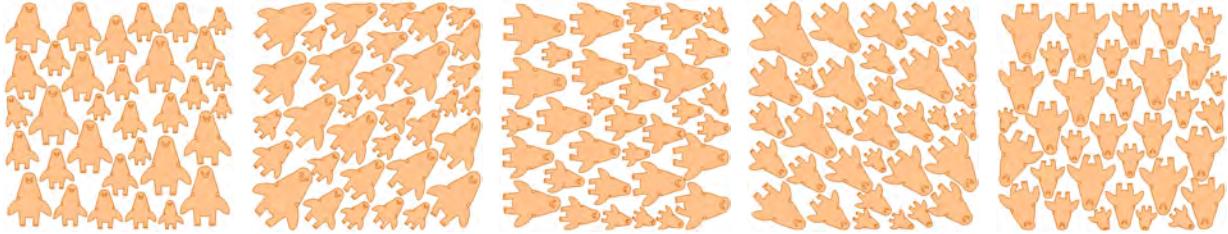


Figure 5.14: Penguins turning into giraffes. The penguins animate by rotating in place. Torsional forces are used to preserve element orientations. Frames are taken at  $t = 0$ ,  $t = 0.125$ ,  $t = 0.25$ ,  $t = 0.375$ , and  $t = 0.5$ .

1582     Figure 5.17 compares a static 2D packing created by RepulsionPak with a frame from  
 1583     an animated packing created by AnimationPak. The extra negative space in AnimationPak  
 1584     comes partly from the trade-off between temporal coherence and tight packing, and partly  
 1585     from the lack of secondary elements, which were used in a second pass in RepulsionPak to  
 1586     fill pockets of negative space.

1587     Figure 5.18 offers a direct comparison between packings computed using Centroidal  
 1588     Area Voronoi Diagrams (CAVD) [SLK05], the spectral approach [DKLS06], and AnimationPak.  
 1589     These packings use stars that rotate and pulsate. For each method we show the initial frame  
 1590     ( $t = 0$ ) and the halfway point ( $t = 0.5$ ). The CAVD approach produces a satisfactory—  
 1591     albeit loosely coupled—packing for the first frame, but because the algorithm was not  
 1592     intended to work on animated elements, the evenness of the packing quickly degrades in  
 1593     later frames. The spectral approach is much better than CAVD, but their animated ele-  
 1594     ments still have fixed spacetime shapes and can only translate and rotate to improve their  
 1595     fit. Repulsion forces and deformation allow AnimationPak to achieve a tighter packing that  
 1596     persists across the animation, including gear-like meshing of ~~oppositely rotating~~ stars.

1597     Figure 5.19 emphasizes the trade-off between temporal coherence and evenness of neg-  
 1598     ative space by creating two animations with different time springs stiffness. In (a), the  
 1599     time springs are 100 times stronger than in (b). The resulting packing has larger pockets  
 1600     of negative space, but the accompanying video shows that the animation is smoother. The  
 1601     packing in (b) is tighter, but the elements must move frantically to maintain that tightness. 

1602     Figure 5.20 is a failed attempt to animate a “blender”. The packing has a beam that  
 1603     rotates clockwise and a number of small unguided circles. In a standard physics simulation  
 1604     we might expect the beam to push the circles around the container, giving each one a  
 1605     helical spacetime trajectory. Instead, as elements grow, repulsion forces cause circles to  
 1606     explore the container boundary, where they discover the lower-energy solution of slipping  
 1607     past the edge of the beam as it sweeps past. If we extend the beam to the full diameter

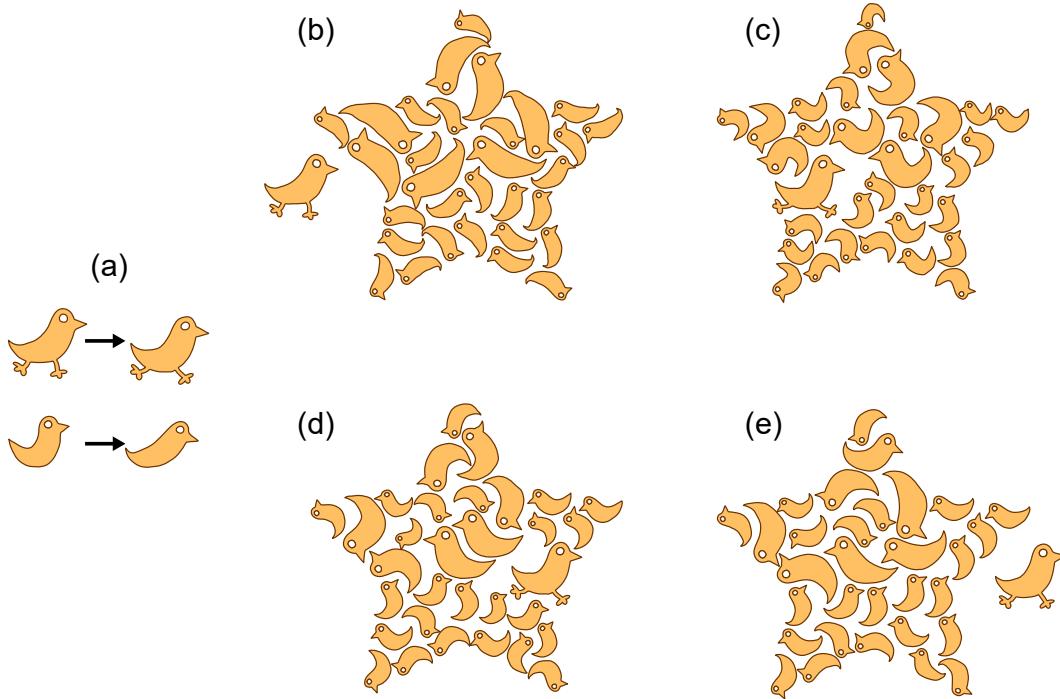


Figure 5.15: An animation of a traveling bird that comes in and out a packing of birds. The other birds adapt to the reduced space by prematurely stopping some of their slices' growths.

of the container, consecutive slices simply teleport across the beam, hiding the moment of overlap in the brief time interval where physical forces were not computed. AnimationPak is not directly comparable to a 3D physics simulation; it is better suited to improving the packing quality of an animation that has already been blocked out at a high level.

## 5.8 Conclusions

We introduced AnimationPak, a system for generating animated packings by filling a static container with animated elements. Every animated 2D element is represented by an extruded spacetime tube. We discretize elements into triangle mesh slices connected by time edges, and deform element shapes and animations using a spacetime physical simulation. The result is a temporally coherent 2D animation of elements that attempt both to perform their scripted motions and consume the negative space of the container. We show a variety of results where 2D elements move around inside the container.

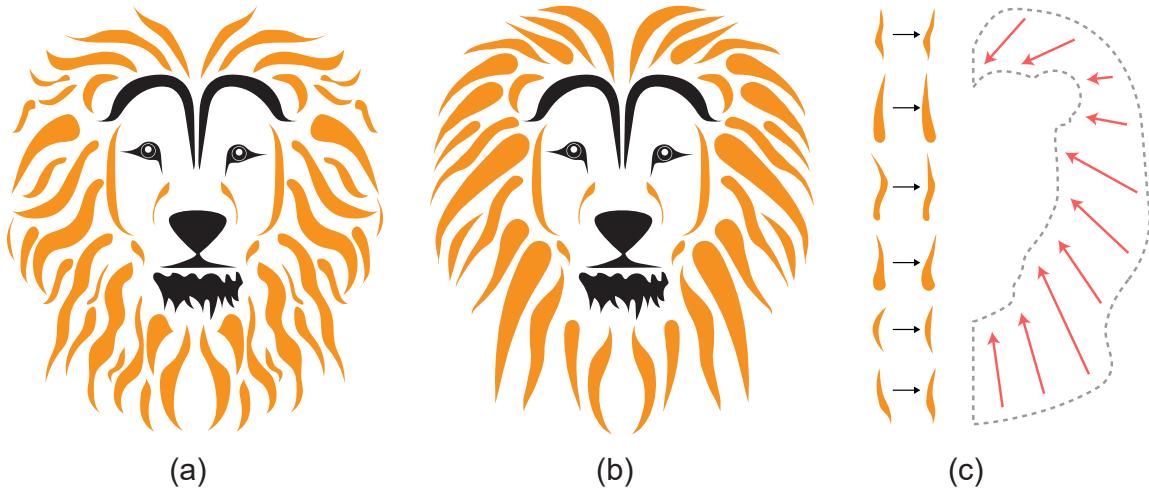


Figure 5.16: (a) A static packing made by an artist, taken from StockUnlimited. (b) The first frame from an AnimationPak packing. (c) The input animated elements and the container shape with a vector field. Torsional forces keep elements oriented in the direction of the vector field. We simulate half of the lion's mane and render the other half using a reflection, and add the facial features by hand.

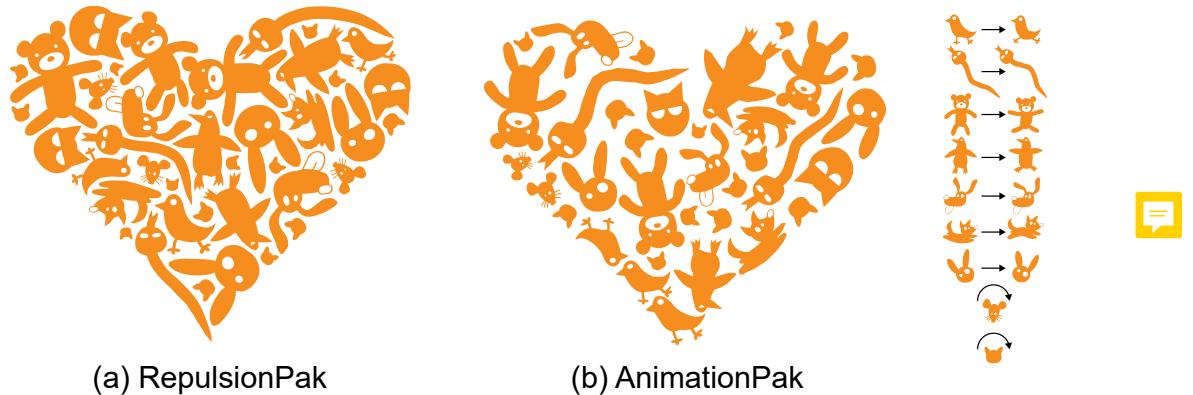


Figure 5.17: (a) A static packing created with RepulsionPak. (b) The first frame of a comparable AnimationPak packing. The input spacetime elements are shown on the right. The AnimationPak packing has more negative space because we must tradeoff between temporal coherence and packing density.

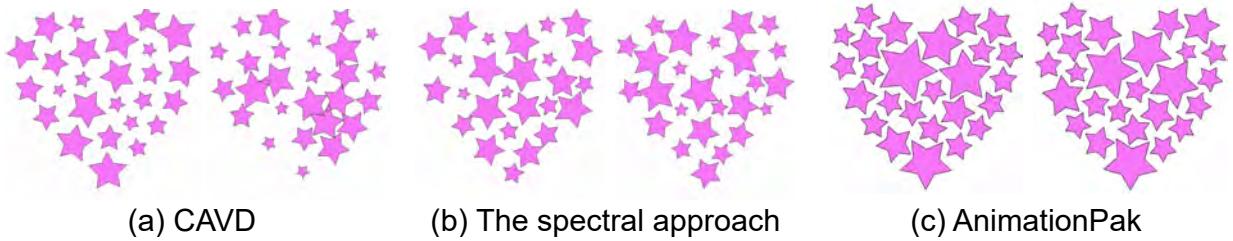


Figure 5.18: A comparison of (a) Centroidal Area Voronoi Diagrams (CAVDs) [SLK05], (b) spectral packing [DKLS06], and (c) AnimationPak. We show two frames for each method, taken at  $t = 0$  and  $t = 0.5$ . The CAVD packing starts with evenly distributed elements but the packing degrades as the animation progresses. The spectral approach improves upon CAVD with better consistency, but still leaves significant pockets of negative space. The AnimationPak packing has less negative space that is more even.

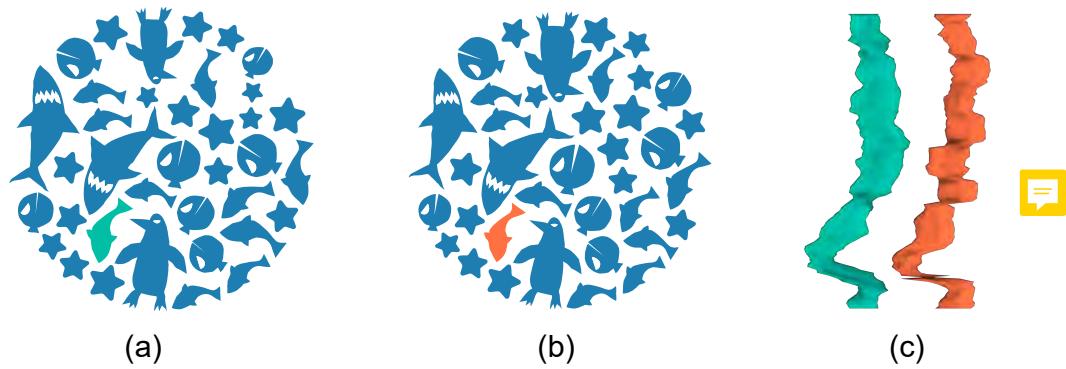


Figure 5.19: (a) One frame from Figure 5.1. (b) The same packing with time springs that are 1% as stiff. Reducing the stiffness of time springs leads to a more even packing with less negative space, but the animated elements must move frantically to preserve packing density. The spacetime trajectories of the highlighted fish in (a) and (b) are shown in (c). The orange fish in (b) exhibits more high frequency fluctuation in its position.

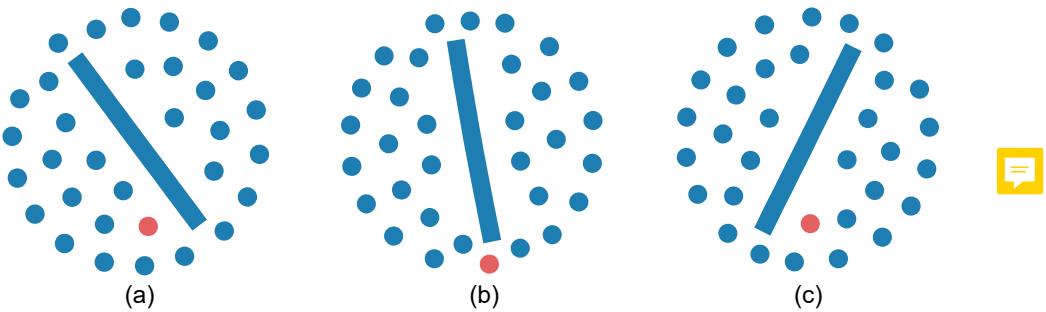


Figure 5.20: A failure case for AnimationPak, consisting of a rotating beam and a number of small circles. Instead of being dragged around by the beam, the circles dodge it entirely by sneaking through the gap between the beam and the container. The red circle demonstrates one such maneuver.

1620

# Chapter 6

1621

## Quantitative Metrics for 2D Packings



1622

### 6.1 Introduction

1623 As with many techniques in artistic applications of computer graphics, evaluating the  
1624 quality of a computer-generated element packing is a challenge. The visual appeal of  
1625 packings follows in part from aligning neighboring elements along compatible segments of  
1626 their boundaries, suggesting that they interlock by design. We believe that the evenness  
1627 of negative space is an indicator of the quality of a packing. The separation between  
1628 neighboring elements should be roughly the same everywhere. In this chapter, we have  
1629 developed several measurements of evenness, which allow us to examine the behavior of  
1630 manually constructed reference packings and to compare RepulsionPak with other packing  
1631 algorithms.



1632

### 6.2 Related Work

1633 In this section, we discuss a few related methods for evaluating discrete texture synthesis,  
1634 stippling, or packings. For more general discussion, Isenberg compiled a survey of  
1635 evaluation methods in NPR research [Ise13].

1636 **Qualitative Evaluation:** AlMeraj et al. [AKA13b] proposed qualitative evaluations  
1637 to find the similarity between exemplars and synthesized textures. They conducted two  
1638 user studies: a pile-sorting study and a pairwise comparison study. These user studies  
1639 concerned about small-sized element distributions, but not element interlocking or the





1640 evenness of negative space, so their work is not suitable for our packing evaluation. Kwan  
1641 et al. performed a simple user study of 13 participants to compare Pyramid of Arclength  
1642 Descriptor (PAD) [KSH<sup>+</sup>16] with previous packing methods. A participant was randomly  
1643 shown a computer-generated packing and they had to give ratings from 1 (worst) to 6  
1644 (best) to each of three categories: their preference, stylishness, interlocking. The study  
1645 revealed that PAD received the highest scores on all three categories. However, we argue  
1646 that the first two categories are ambiguous and do not ~~give better understanding to the~~  
1647 ~~quality of a packing.~~

1648 **Quantitative Evaluation:** In Jigsaw Image Mosaics (JIM) [KP02], Kim et al. used  
1649 an energy minimization approach that penalizes a packing if it has too much negative space,  
1650 too many overlaps, or severely deformed elements. They did not further discuss whether  
1651 it can be used for a quantitative evaluation, but we are inspired by their work to develop  
1652 a metric based on shape overlaps (Section 6.3.3). Maciejewski et al. [MIA<sup>+</sup>08] compared  
1653 computer-generated stippling artwork with artist-made stippling artwork using the Gray-  
1654 Level Co-Occurrence Matrix (GLCM), which measures spatial relationships between pixel  
1655 intensities. GLCM works best for analyzing dense point distributions, but it unfortunately  
1656 cannot be used to analyze arrangements with larger shapes. In medical research, Aliy et  
1657 al. [ASF<sup>+</sup>13] proposed a method to analyze the arrangement of small biological objects  
1658 in a histology image. They treated these objects as a point distribution by computing  
1659 their centroids and encoded their spatial relationships, so their work is unsuitable for our  
1660 research. In fabrication, the packing quality is dictated by manufacturing constraints,  
1661 such as connectivity, so that printed objects can resist from breaking [CZX<sup>+</sup>16, ZCT16,  
1662 MSS<sup>+</sup>19]. However, our objective is more about the evenness of negative space and less  
1663 about connectivity strengths.

### 1664 6.3 Quantitative Metrics



1665 **Calibration:** Two packings must be calibrated to each other before our measurements  
1666 can be compared meaningfully. We can compare packings of different sizes by normalizing  
1667 their containers to have unit area. We must also arrange for the packings to have the  
1668 same *negative space ratio* (the overall amount of negative space as a fraction of container  
1669 area), so that our measurements can focus on the distribution of negative space and not  
1670 just the amount. Given two calibrated packings, we examine three main evenness metrics:  
1671 spherical contact probabilities, distance histograms, and overlap of offset elements.

### 6.3.1 Spherical Contact Probability

The spherical contact probability (SCP) is the probability that a disc of radius  $r$ , chosen uniformly at random within the container region, lies entirely within the packing's negative space [CSKM13]. The SCP can be summarized via a function  $Q_s(r)$  that gives this probability for each radius  $r$ . In order to interpret the SCP, it is helpful first to examine a “packing” with perfectly even negative space (Figure 6.1a). Consider a pattern of infinite horizontal stripes of width  $d_s$ , separated from each other by negative space of width  $d_{\text{gap}}$ . For this pattern,  $Q_s(0) = d_{\text{gap}}/(d_{\text{gap}} + d_s)$ ; it is also clear that  $Q_s(d_{\text{gap}}/2) = 0$ , because no disc of diameter greater than  $d_{\text{gap}}$  can fit in the negative space (our  $d_{\text{gap}}$  is twice the radius of the ball). Furthermore,  $Q_s(r)$  will decrease linearly between these two points, and remain at zero thereafter; its graph will consist of a tilted line segment connected to a horizontal ray.

No real-world packing exhibits this SCP. Even in a perfect arrangement of squares (Figure 6.1b), the intersections of horizontal and vertical channels produce pockets of negative space that can accommodate balls of radius  $d_{\text{gap}}\sqrt{2}/2$ . These pockets tend to raise the SCP slightly everywhere, and cause it to bend into a small tail that approaches zero gradually. For a given set of elements in a container, the best packings will have a steeply-decreasing SCP that stays close to the idealized stripe function most of the way down, has a low value at  $r = d_{\text{gap}}/2$ , and then bends towards horizontal near that value. Note that there is always a largest disc that can fit in a packing's negative space, and hence a largest  $r$  for which  $Q_s(r) > 0$ . In our graphs, we plot  $Q_s(r)$  only until this point, allowing us to compare the largest empty gaps of two packings. In less effective packings (Figs. 6.1c,d), the negative space will be narrower in some places and wider in others, recognizable as a shallower SCP with a longer tail.

We geometrically compute SCP by offsetting the negative space inward. Let  $N$  be the shape of the negative space (essentially the container region with holes where elements are). For a given radius  $r$ , we compute  $N(r)$ , the Minkowski difference of  $N$  with a disc of radius  $r$ . Then  $Q_s(r)$  is simply the ratio of the area of  $N(r)$  to the area of the container.

### 6.3.2 Histograms of the Distance Transform

The distance transform of the negative space provides insights into how negative space varies. In the context of 2D packings, a distance transform is a 2D image where each pixel contains the radius  $r$  value to the nearest element. Figure 6.2 shows an example for a distance transform, lighter a color of a pixel, higher the  $r$  value. We investigated

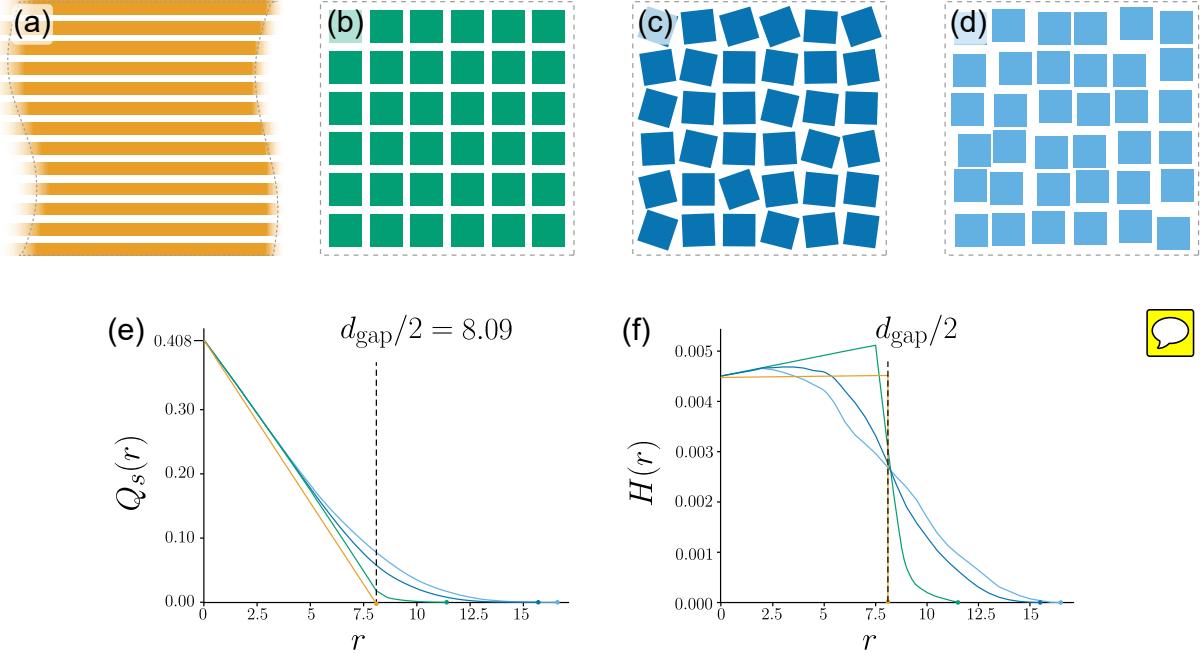


Figure 6.1: Spherical contact probabilities and distance histograms for reference packings. A “perfect packing” of infinite stripes is shown in (a), followed by a square packing with the same area fraction and **negative space width**  $d_{\text{gap}}$  in (b). The square packing is then perturbed with random rotations in (c) and translations in (d). The corresponding SCP functions and histograms are plotted in (e) and (f).

the idea of analyzing  $r$  values only on a medial axis since it bisects the gaps of negative space. However, the challenge is to prune a significant number of problematic medial axis’s branches that are either too long or too short, as shown in Figure 6.3.

As creating a high quality medial axis is difficult, we decide to analyze the entire negative space by computing the histogram of the distance transform. However, this would require quantizing  $r$  values into bins. Instead, note that the SCP  $Q_s(r)$  is precisely the normalized area of negative space for which the distance transform is at least  $r$ . Looked at another way, if we offset each element by a radius  $r$ , then  $1 - Q_s(r)$  must be the normalized area of the union of these offset elements. This area can then be interpreted as a cumulative distribution function of distance. From this observation we can compute a continuous variant of the distance histogram as a probability density function via the derivative of the SCP:  $H(r) = -Q'_s(r)$ .

Given two calibrated packings, the areas under their distance histograms are the same.

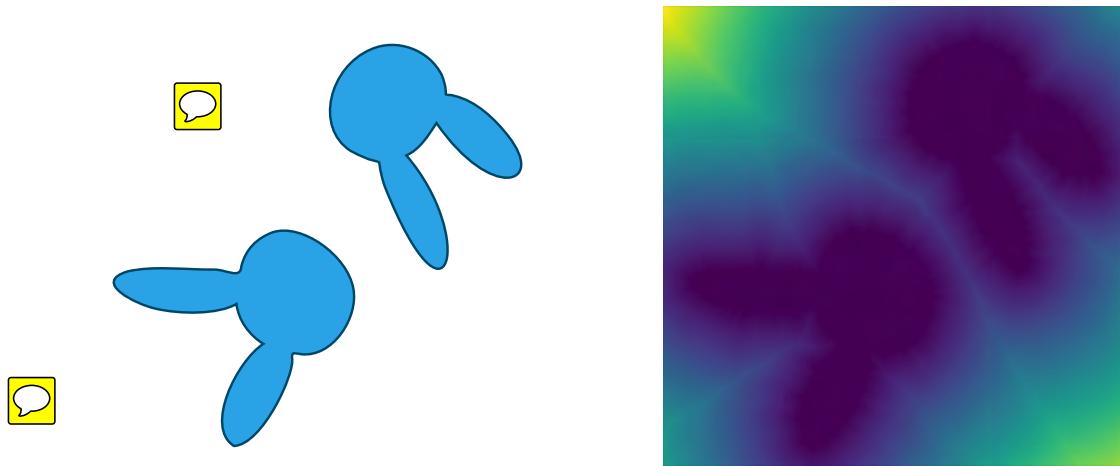


Figure 6.2: The image on the right is a distance transform generated from negative space around two rabbits. The brighter the color of a pixel, the farther it is from the nearest boundary.

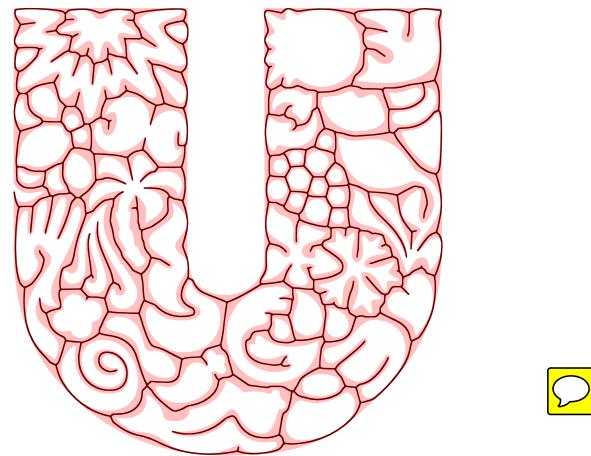


Figure 6.3: A medial axis generated from the negative space of the Unilever packing using Zhang-Suen thinning algorithm [ZS84]. The medial axis contains unnecessarily long branches that touch the element boundaries and inconsequential short branches that should be deleted.

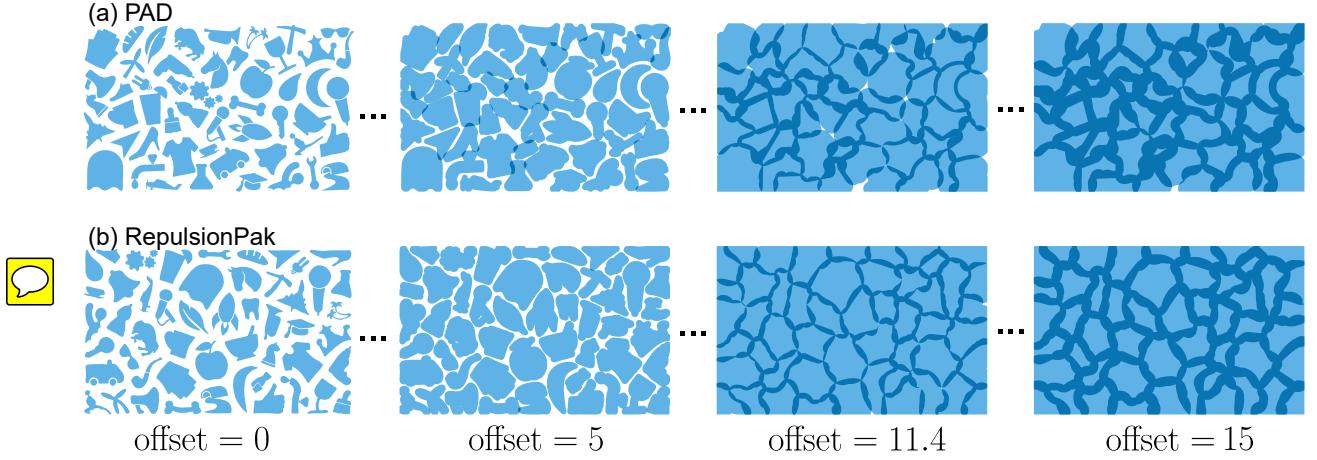


Figure 6.4: An illustration of offsetting elements outward. The packings have  $d_{\text{gap}}/2 = 5.7$ . At an offset of 5, which is slightly less than  $d_{\text{gap}}/2$ , the overlap for the PAD packing is 1.462% of the total area, while the overlap for our packing is only 0.039%. At an offset of 11.4, which equals  $d_{\text{gap}}$ , the PAD packing shows more empty space (1.05%) than RepulsionPak (0.07%). As the offset is increased, overlaps in the PAD packing create channels with uneven widths, whereas ours are more uniform.

1718 But a more even packing will have a shorter tail, indicating a tighter upper bound on gap  
 1719 size, and it will have a larger concentration of density around  $d_{\text{gap}}/2$ . In Figure 6.1, the  
 1720 histogram for the perfect stripe pattern is a step function that drops to zero at  $d_{\text{gap}}/2$ .  
 1721 The ideal square packing has a histogram that climbs gently until around  $d_{\text{gap}}/2$  before  
 1722 dropping steeply. The other two packings have shallower, smoother histograms. Note  
 1723 also that high values of the distance histogram near  $d_{\text{gap}}/2$  correspond to a rapid negative  
 1724 change in the SCP, suggesting a more even packing.

### 1725 6.3.3 The Overlap Function

1726 The overlap function is a function of a non-negative offset amount  $r$ . For any given  $r$ ,  
 1727 we offset every element by computing its Minkowski sum with a disc of radius  $r$ . As  $r$   
 1728 grows, elements will start overlapping; the overlap function measures the total area of  
 1729 these overlaps, normalized by container area, as a function of  $r$ . We can also visualize  
 1730 these overlapping areas directly as in Figure 6.4. In a perfect packing, we would expect  
 1731 no overlaps until  $r = d_{\text{gap}}/2$ , our desired gap distance, at which point overlapping areas  
 1732 would start to grow into channels of roughly even width.

1733 **6.4 Comparisons**



1734 For all comparisons below, we estimate  $d_{\text{gap}}/2$  as the average of skin widths of all elements in  
1735 a RepulsionPak packing after the simulation stops. We assume the other non-RepulsionPak  
1736 packing has approximately the same  $d_{\text{gap}}/2$  value since both packings have the identical  
1737 negative space ratio due to the calibration.

1738 **Comparison to PAD:** Figure 6.5 compares RepulsionPak and PAD [KSH<sup>+</sup>16]. Pack-  
1739 ing (a) is a result from the PAD paper; Packing (b) was created with RepulsionPak using  
1740 the same elements, and calibrated to have the same negative space as (a). Note that the  
1741 PAD packing actually has several overlapping elements (for example, the tooth and the  
1742 horse), but white haloes around elements artfully conceal overlaps with little degradation  
1743 in visual quality. Our packing avoids overlaps by design. The SCP plot in (c) shows that  
1744 our packing has a lower value at  $d_{\text{gap}}/2$ , indicating more even negative space, and has a  
1745 shorter tail, indicating fewer large empty areas. Our result also has a histogram bump  
1746 around  $d_{\text{gap}}/2$ , and a lower overlap function.

1747 **Comparison to an Artist-Made Packing:** In Figure 6.6 we show a RepulsionPak  
1748 result created using the elements from the artist-made packing. Our packing was calibrated  
1749 to match the artist's. Looking closely, the artist's packing has a few elements separated  
1750 by narrow gaps, such as the cherries and the corn on the top left. Our result has fewer  
1751 large empty gaps, as indicated by a short tail in its SCP. Our result also has a histogram  
1752 bump around  $d_{\text{gap}}/2$ , and a lower overlap function. The result shows the effectiveness of  
1753 the repulsion forces in successfully discovering compatibilities in the element boundaries  
1754 and filling the space effectively.

1755 **Comparison to Rigid Packings:** To evaluate the effect of deformation on negative  
1756 space, we compute all three metrics under increasing values of  $k_{\text{edg}}$ , the edge force relative  
1757 strength. Note that we only modify the strengths of edge springs and shear springs, and  
1758 the strength of negative space springs are unchanged. Increasing  $k_{\text{edg}}$  allows the element  
1759 meshes to resist deformation, ultimately approximating a rigid packing algorithm. We  
1760 created 15 packings, five for each of three values of  $k_{\text{edg}}$  (5, 250, and 1000). Each packing  
1761 used 25 elements chosen at random from a library of 60, with no secondary elements. All  
1762 packings are partly calibrated: they all have the same negative space ratio, but elements  
1763 are chosen at random and have some variation in their sizes. As shown in Figure 6.7, a  
1764 low value of  $k_{\text{edg}}$  leads to greater deformation and steeper SCPs. The more pronounced  
1765 histogram bump at  $r = d_{\text{gap}}/2$  suggests more even negative space. The lower overlap  
1766 functions indicate fewer narrow gaps between elements.



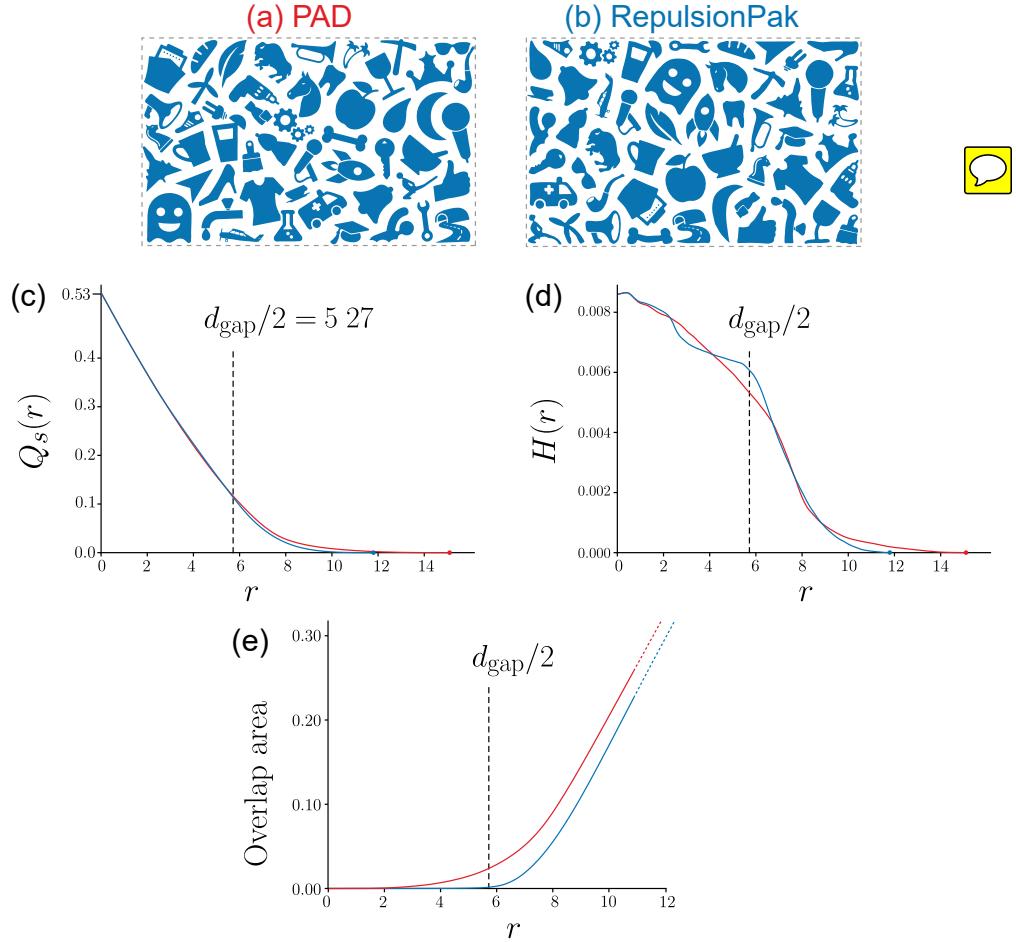


Figure 6.5: A comparison between a PAD packing shown in (a) and a RepulsionPak packing in (b) with their corresponding SCPs (c), distance histograms (d), and overlap functions (e). The PAD and RepulsionPak packings are calibrated to have the same negative space ratio. Our SCP is lower and shorter than the PAD's result, our histogram shows higher concentration around  $d_{\text{gap}}/2$ , indicating more even negative space, and our packing has a lower overlap function.

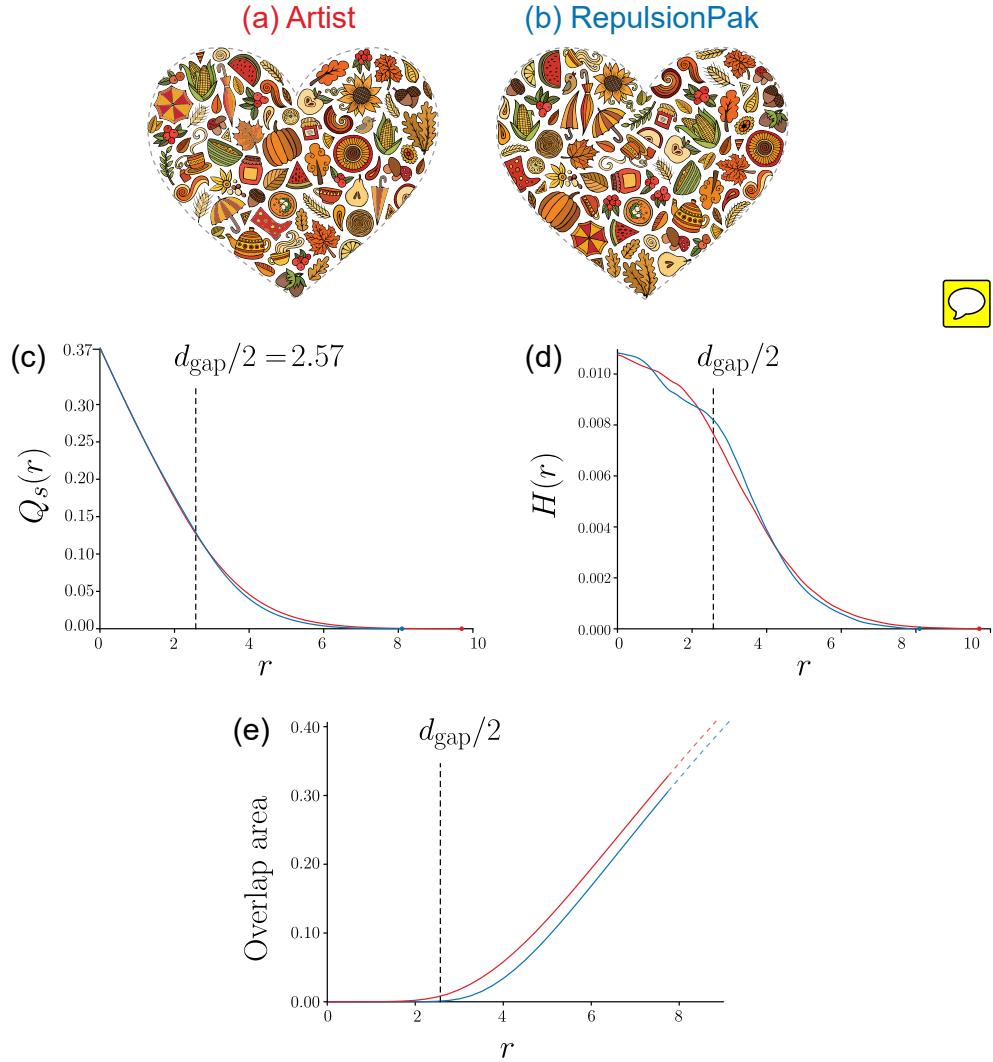


Figure 6.6: A comparison between the artist-made packing (Artist: Balabolka on Shutterstock) shown in (a), and a RepulsionPak packing with the same elements in (b). We plot the corresponding SCPs (c), distance histograms (d), and overlap functions (e). For comparison purposes we remove secondary elements from the artist's packing. Our SCP is lower and shorter than the artist's result, our histogram shows more concentration around  $d_{\text{gap}}/2$ , and RepulsionPak also has a lower overlap function.

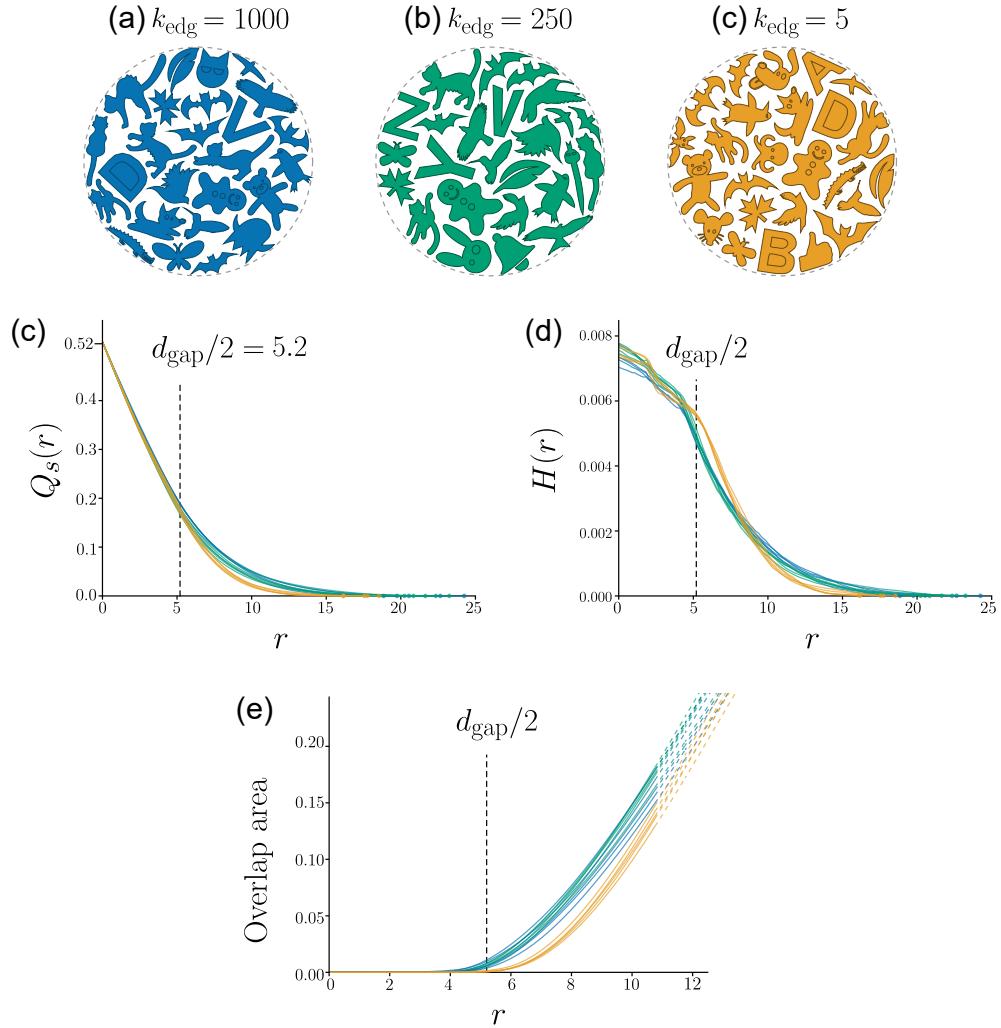


Figure 6.7: A demonstration of the effect of deformation on the evenness of negative space. The packings in (a), (b) and (c) are representative results using three values of the edge force strength  $k_{\text{edg}}$ , from rigid ( $k_{\text{edg}} = 1000$ ) to moderate ( $k_{\text{edg}} = 250$ ) to deformable ( $k_{\text{edg}} = 5$ ). We construct five random packings for each value of  $k_{\text{edg}}$ , and plot their SCPs (d), histograms (e), and overlap functions (f). Packings with the most deformation have steeper and shorter SCPs, more histogram concentrations around  $d_{\text{gap}}/2$ , and lower overlap functions.



1767 **6.5 Conclusions**

1768 The evenness of negative space is an indicator of the quality of a packing. The more  
1769 the elements interlock, the more even the negative space. We evaluated the evenness of  
1770 negative space using spherical contact probabilities, histograms of distance transform, and  
1771 overlap functions. A packing with a more even negative space has a steeper and shorter  
1772 SCP, more histogram concentrations around  $d_{\text{gap}}/2$ , and a lower overlap function.

<sub>1773</sub> **Chapter 7**

<sub>1774</sub> **Conclusions and Future Work**

<sub>1775</sub> **7.1 Conclusions**

<sub>1776</sub> We presented three deformation-driven packing methods. FLOWPAK is a method to  
<sub>1777</sub> deform and pack long thin elements to follow a vector field. RepulsionPak is a method  
<sub>1778</sub> that uses repulsion forces to pack and deform elements that are represented as mass-spring  
<sub>1779</sub> systems. AnimationPak is an extension of RepulsionPak that packs animated 2D elements,  
<sub>1780</sub> each **is** an extruded 3D shape in a spacetime domain.

<sub>1781</sub> Given a small size element library, we demonstrated that deformation-driven methods  
<sub>1782</sub> can create element compatibilities and fill the container effectively. Repeated elements give  
<sub>1783</sub> a sense of uniformity but deformation creates a sense of variety.

<sub>1784</sub> We discussed **that** the evenness of negative space **is** an indicator of the quality of  
<sub>1785</sub> a packing. We measured the evenness using three statistical metrics: spherical contact  
<sub>1786</sub> probabilities, histograms of distance transforms, and overlap functions.

<sub>1787</sub> **7.2 Future Work**

<sub>1788</sub> We see many possibilities for further improvements to our methods and packing research.



### 1789 7.2.1 FLOWPAK

- 1790 • **Element Threading:** In FLOWPAK, an element must completely use up its streamline, causing a few elements to stretch too long. We would like to combine multiple shorter elements by threading them along streamlines, creating a longer *compound element* (Figure 7.1a). First, we need a way to connect the overlapping segments of the elements together to create smooth and plausible transitions. This possibly leads us to a deeper research topic: how to develop a content-aware blending method on vector graphics patterns. Second, we should carefully decide on the selection of elements we have to thread and the number of elements we need. We would like to investigate a greedy approach with backtracking [KP02].
- 1799 • **Element Branching:** We also would like to support branching structures to ~~resemble~~ floral ornaments (Figure 7.1b). We start ~~with~~ tracing a long streamline that follows the vector field, and then generate shorter streamline branches, which produce blobs that partly overlap the main blob. Similarly, we also need to connect and blend element segments at every intersection. We also should consider the semantics of the element shapes. Not every part of an element can support branching; for example, a leaf cannot be connected with twigs. In another example, a stem or a trunk is an ideal segment where we can glue other elements.
- 1807 • **High Curvature Streamlines:** FLOWPAK results do not have significant high-curvature streamlines like u-turns, since they could unpleasantly fold the decorative elements. This could be solved with a folding avoidance algorithm [Ase10].
- 1810 • **Better Iterative Refinement:** The iterative refinement process uses a greedy approach, in which we iterate over all placed elements in a fixed order from smallest to largest. We would like to investigate global optimizations that could be applied to improve the overall composition in an order-independent way. A natural first choice would be an approach based on simulated annealing, although performance could become a more serious issue in that case. Another idea is to use RepulsionPak to refine the packing ~~but we need to enforce a constraint on elements' spines to ensure they adhere to the vector field.~~
- 1818 • **Automatic Fixed Element Creation:** We would like to explore the automatic creation and placement of the fixed elements, perhaps by discovering them as salient regions in source photographs, and extracting and vectorizing them. This extraction must be carried out carefully, yielding enough fixed elements to communicate a container clearly without disrupting the uniformity of the design.

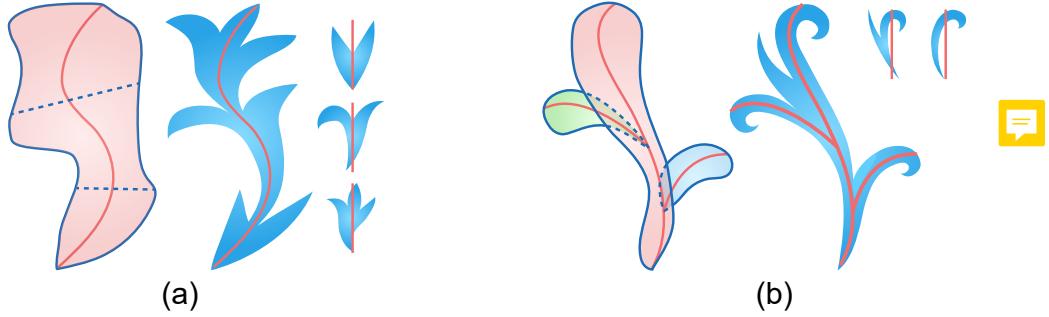


Figure 7.1: (a) Thread multiple elements on a single streamline. (b) Combine multiple streamlines to create branching.

### 1823 7.2.2 RepulsionPak

- 1824 • **Improving the Numerical Integrator:** We deliberately chose a simple simulation  
1825 model based on springs and forward Euler integration, because our main goal  
1826 was to demonstrate the validity of a deformation-driven approach, and not to con-  
1827 tribute a new physical simulation method. Contemporary research has yielded many  
1828 more sophisticated physical simulation methods, such as Position Based Dyna-  
1829 mics [MHHR07], Projective Dynamics [BML<sup>+</sup>14], and the Finite Element Method. No  
1830 one method is obviously best suited to this problem, and we intend to experiment  
1831 with several to investigate if any offers a suitable trade-off between performance and  
1832 quality.
- 1833 • **RepulsionPak for Fabrication:** We would like to explore the use of RepulsionPak  
1834 in a fabrication context. For example, our boundary compatibilities might be used  
1835 to create a connected object. Alternatively, it would be interesting to 3D print the  
1836 negative space, which is already connected, leaving holes that surround the element  
1837 shapes.
- 1838 • **Vector Graphics Warping:** Our barycentric warping method can introduce unde-  
1839 sirable artifacts in highly deformed elements, as in the swallow tails in the left result  
1840 of Figure 4.9. We would like to explore other methods for warping an element’s geo-  
1841 metry based on the correspondences between the triangles of its original mesh and  
1842 the deformed meshes in the final packing, based for example on the work of Jacobson  
1843 et al. [JBPS11] and Liu et al. [LJG14].
- 1844 • **Interactive packings:** We would like to develop human-in-the-loop interactions,  
1845 in which the user and the computer work together to create a desired composition.

1846 Examples of recent work in this style include research by Zehnder et al. [ZCT16],  
1847 Gieseke et al. [GALF17], and Hsu et al. [HWYZ20], which let the user directly place  
1848 and manipulate elements while a composition is being created. It would also be  
1849 interesting to display an interactive packing on a large touch screen, and let people  
1850 ~~to move elements around by touching and dragging gestures.~~

### 1851 7.2.3 AnimationPak

- 1852 • **2D Vector Graphics Interpolation:** Because we use linear interpolation to ~~ren-~~  
1853 ~~der~~ an element’s shape between slices, we require elements not to undergo changes  
1854 in topology. More sophisticated representations of vector shapes, such as that of  
1855 Dalstein et al. [DRvdP15], could support interpolations between slices with complex  
1856 topological changes. We would also need to synthesize a watertight envelope around  
1857 the animating element in order to compute overlap and repulsion forces.
- 1858 • **Dynamic Mesh Resolution:** We would like to improve the performance of the  
1859 physical simulation. One option may be to increase the resolution of element meshes  
1860 progressively during simulation. Early in the process, elements are small and dis-  
1861 tant from each other, so lower-resolution meshes may suffice for computing repulsion  
1862 forces.
- 1863 • **Continuous Collision Detection:** Our discrete simulation can miss element over-  
1864 laps that occur between slices. A more robust continuous collision detection (CCD)  
1865 algorithm such as that of Brochu et al. [BEB12] could help us find all collisions  
1866 between the envelopes of spacetime elements.
- 1867 • **Secondary Elements:** In RepulsionPak ~~[SKA18]~~, an additional pass with small  
1868 secondary elements had a significant positive effect on the distribution of negative  
1869 space in the final packing. It may be possible to identify stretches of unused spacetime  
1870 that can be filled opportunistically with additional elements. The challenge would  
1871 be to locate tubes of empty space that run the full duration of the animation ~~and~~  
1872 ~~have sufficient diameters to accommodate added elements.~~
- 1873 • **Animated Containers:** Like the spectral method [DKLS06], and unlike Animo-  
1874 saics [SLK05], AnimationPak can pack animated elements into a static container.  
1875 We would like to extend our work to also handle animated containers. ~~As demon-~~  
1876 ~~strated by the packing with a loose bird in Figure 5.15, elements can vary their slice~~  
1877 ~~sizes to adapt to changes in container area. A more challenging solution is adding and~~

removing elements, as this would certainly affect the initial element placement, which would need to ensure that elements are placed fully inside the spacetime volume of the container.

- **Deformation-Driven 3D Packings:** AnimationPak implements forces and constraints geared towards spacetime animation, but many of the same ideas could be adapted to develop a deformation-driven method for packing purely spatial 3D objects into a 3D container. We would like to evaluate the expressivity and visual quality of deformation-driven 3D packings in comparison to other 3D packing techniques.

 **User Study with Artists:** There are many examples of static two-dimensional packings created by artists, which can serve as inspiration for an algorithm like RepulsionPak. We were only able to find a single example of animated packing from The Simpsons (Figure 5.2). We think its unpopularity is because it is difficult and time-consuming to create by hand. Therefore, we would like to engage with artists to understand the aesthetic value and limitations of AnimationPak.

#### 7.2.4 Packing Evaluation

- **Human Perception of the Evenness:** We would like to conduct experiments that investigate the extent to which quantitative measurements of the evenness of negative space in a packing correlate with the human perception of a packing’s quality. In informal evaluations, some viewers found that the packing in Figure 6.6b, created with RepulsionPak, was packed more tightly than the artist’s packing in Figure 6.6a, even though both have the same total amount of negative space. 
- **Measuring Other Design Principles:** We would like to develop additional metrics to evaluate how well an ornamental design fulfills other design principles. A measure of element deformation in a composition would permit a comparison against future deformation-driven techniques. In Chapter 3, we argue that visual flow and “uniformity amidst variety” are important to attractive packings. In another study, Wong et al. [WZS98] describe basic design principles for decorative arts: repetition, balance, and conformation to geometric constraints. The rigorous expression of aesthetic principles is a fascinating area for future research.
- **Different Offsetting for SCP:** Our validation metrics are all based on Minkowski sums or differences with discs, corresponding to a form of shape offsetting where corners become round. The result of this rounding is visible in our graphs, for

example in the gradual flattening of the SCP for the squares in Figure 6.1e. It would be worthwhile to repeat these measurements using mitered offsetting, and to evaluate whether rounded or mitered offsetting is a closer match to human perceptual judgment of evenness.



- **Better Visualization for SCP:** When comparing calibrated packings, SCPs communicate differences in the evenness of negative space, but the differences between SCPs can be subtle. In addition to distance histograms, we would like to investigate other visualizations of this information that might amplify these differences to make evaluation easier.
- **Evaluating 2D Animated Packings:** All three metrics are only for evaluating 2D packings. While they extend naturally to three purely spatial dimensions, it is not clear whether they can be adapted to the spacetime context. We would like to investigate spatial statistics for the quality of animated packings created by AnimationPak that correlate with human perceptual judgments.



<sup>1924</sup>

# References

- <sup>1925</sup> [AGYS14] Rinat Abdrashitov, Emilie Guy, JiaXian Yao, and Karan Singh. Mosaic: Sketch-Based Interface for Creating Digital Decorative Mosaics. In *Proceedings of the 4th Joint Symposium on Computational Aesthetics, Non-Photorealistic Animation and Rendering, and Sketch-Based Interfaces and Modeling*, SBIM '14, pages 5–10, New York, NY, USA, 2014. ACM.
- <sup>1926</sup>
- <sup>1927</sup>
- <sup>1928</sup>
- <sup>1929</sup>
- <sup>1930</sup> [AKA13a] Zainab AlMeraj, Craig S. Kaplan, and Paul Asente. Patch-Based Geometric Texture Synthesis. In *Proceedings of the Symposium on Computational Aesthetics*, CAE '13, pages 15–19, New York, NY, USA, 2013. ACM.
- <sup>1931</sup>
- <sup>1932</sup>
- <sup>1933</sup> [AKA13b] Zainab AlMeraj, Craig S. Kaplan, and Paul Asente. Towards Effective Evaluation of Geometric Texture Synthesis Algorithms. In *Proceedings of the Symposium on Non-Photorealistic Animation and Rendering*, NPAR 13, page 514, New York, NY, USA, 2013. Association for Computing Machinery.
- <sup>1934</sup>
- <sup>1935</sup>
- <sup>1936</sup>
- <sup>1937</sup> [AKMS20] Paul Asente, Craig Kaplan, Radomír Měch, and Reza Adhitya Saputra. Computerized Generation of Ornamental Designs by Placing Instances of Simple Shapes in Accordance With a Direction Guide, February 11 2020. US Patent 10,559,061.
- <sup>1938</sup>
- <sup>1939</sup>
- <sup>1940</sup>
- <sup>1941</sup> [Ase10] Paul Asente. Folding Avoidance in Skeletal Strokes. In *Eurographics Workshop on Sketch-Based Interfaces and Modeling*. The Eurographics Association, 2010.
- <sup>1942</sup>
- <sup>1943</sup> [ASF<sup>+</sup>13] Maya Alsheh Aliy, Johanne Seguin, Aurélie Fischer, Nathalie Mignet, Laurent Wendling, and Thomas Hurtut. Comparison of the Spatial Organization in Colorectal Tumors using Second-Order Statistics and Functional ANOVA. In *2013 8th International Symposium on Image and Signal Processing and Analysis (ISPA)*, pages 257–261, 2013.
- <sup>1944</sup>
- <sup>1945</sup>
- <sup>1946</sup>
- <sup>1947</sup>

- 1948 [BBT<sup>+</sup>06] Pascal Barla, Simon Breslav, Joëlle Thollot, François Sillion, and Lee  
 1949 Markosian. Stroke Pattern Analysis and Synthesis. In *Computer Graphics  
 1950 Forum (Proc. of Eurographics 2006)*, volume 25, 2006.
- 1951 [BEB12] Tyson Brochu, Essex Edwards, and Robert Bridson. Efficient Geometrically  
 1952 Exact Continuous Collision Detection. *ACM Trans. Graph.*, 31(4), July 2012.
- 1953 [BL15] Thomas F. Banchoff and Stephen T. Lovett. *Differential Geometry of Curves  
 1954 and Surfaces*. Chapman and Hall/CRC, 2015.
- 1955 [BML<sup>+</sup>14] Sofien Bouaziz, Sebastian Martin, Tiantian Liu, Ladislav Kavan, and Mark  
 1956 Pauly. Projective Dynamics: Fusing Constraint Projections for Fast Simula-  
 1957 tion. *ACM Trans. Graph.*, 33(4):154:1–154:11, July 2014.
- 1958 [Bri07] Robert Bridson. Fast Poisson Disk Sampling in Arbitrary Dimensions. In  
 1959 *ACM SIGGRAPH 2007 Sketches*, SIGGRAPH ’07, New York, NY, USA, 2007.  
 1960 ACM.
- 1961 [BTW09] Thomas Byholm, Martti Toivakka, and Jan Westerholm. Effective Packing of  
 1962 3-Dimensional Voxel-Based Arbitrarily Shaped Particles. *Powder Technology*,  
 1963 196(2):139–146, 2009.
- 1964 [BvMM11] B. Beneš, O. Št'ava, R. Měch, and G. Miller. Guided Procedural Modeling.  
 1965 *Computer Graphics Forum*, 30(2):325–334, 2011.
- 1966 [BW95] Armin Bruderlin and Lance Williams. Motion Signal Processing. In *Proceed-  
 1967 ings of the 22nd Annual Conference on Computer Graphics and Interactive  
 1968 Techniques*, SIGGRAPH 95, pages 97–104, New York, NY, USA, 1995. Asso-  
 1969 ciation for Computing Machinery.
- 1970 [BWL18] Xiaojun Bian, Li-Yi Wei, and Sylvain Lefebvre. Tile-Based Pattern Design  
 1971 with Topology Control. *Proc. ACM Comput. Graph. Interact. Tech.*, 1(1),  
 1972 July 2018.
- 1973 [CFH<sup>+</sup>09] M. Cui, J. Femiani, J. Hu, P. Wonka, and A. Razdan. Curve Matching for  
 1974 Open 2D Curves. *Pattern Recogn. Lett.*, 30(1):1–10, January 2009.
- 1975 [CKHL11] Myung Geol Choi, Manmyung Kim, Kyung Lyul Hyun, and Jehee Lee. De-  
 1976 formable Motion: Squeezing into Cluttered Environments. *Computer Graphics  
 1977 Forum*, 30(2):445–453, 2011.

- 1978 [Coh92] Michael F. Cohen. Interactive Spacetime Control for Animation. In *Proceedings of the 19th Annual Conference on Computer Graphics and Interactive Techniques*, SIGGRAPH 92, pages 293–302, New York, NY, USA, 1992. Association for Computing Machinery.
- 1979
- 1980
- 1981
- 1982 [CSKM13] Sung Nok Chiu, Dietrich Stoyan, Wilfrid S. Kendall, and Joseph Mecke. *Stochastic Geometry and Its Applications*. Wiley Series in Probability and Statistics. Wiley, 2013.
- 1983
- 1984
- 1985 [CSR10] N. Chernov, Yu. Stoyan, and T. Romanova. Mathematical Model and Efficient Algorithms for Object Packing Problem. *Computational Geometry*, 43(5):535 – 553, 2010.
- 1986
- 1987
- 1988 [CSY02] Jonathan Cagan, Kenji Shimada, and Sun Yin. A Survey of Computational Approaches to Three-Dimensional Layout Problems. *Computer-Aided Design*, 34(8):597–611, 2002.
- 1989
- 1990
- 1991 [CZL<sup>+</sup>15] Xuelin Chen, Hao Zhang, Jinjie Lin, Ruizhen Hu, Lin Lu, Qixing Huang, Bedrich Benes, Daniel Cohen-Or, and Baoquan Chen. Dapper: Decompose-and-Pack for 3D Printing. *ACM Trans. Graph.*, 34(6), October 2015.
- 1992
- 1993
- 1994 [CZX<sup>+</sup>16] Weikai Chen, Xiaolong Zhang, Shiqing Xin, Yang Xia, Sylvain Lefebvre, and Wenping Wang. Synthesis of Filigrees for Digital Fabrication. *ACM Trans. Graph.*, 35(4):98:1–98:13, July 2016.
- 1995
- 1996
- 1997 [DKLS06] Ketan Dalal, Allison W. Klein, Yunjun Liu, and Kaleigh Smith. A Spectral Approach to NPR Packing. In *Proceedings of the 4th International Symposium on Non-Photorealistic Animation and Rendering*, NPAR ’06, pages 71–78, New York, NY, USA, 2006. ACM.
- 1998
- 1999
- 2000
- 2001 [DRvdP15] Boris Dalstein, Rémi Ronfard, and Michiel van de Panne. Vector Graphics Animation with Time-Varying Topology. *ACM Trans. Graph.*, 34(4), July 2015.
- 2002
- 2003
- 2004 [Eri05] Christer Ericson. Chapter 5: Basic Primitive Tests. In Christer Ericson, editor, *Real-Time Collision Detection*, The Morgan Kaufmann Series in Interactive 3D Technology, pages 125–233. Morgan Kaufmann, San Francisco, 2005.
- 2005
- 2006
- 2007 [GALF17] Lena Gieseke, Paul Asente, Jingwan Lu, and Martin Fuchs. Organized Order in Ornamentation. In *Proceedings of the Symposium on Computational Aesthetics*, CAE ’17, pages 4:1–4:9, New York, NY, USA, 2017. ACM.
- 2008
- 2009

- 2010 [GBLM16] Paul Guerrero, Gilbert Bernstein, Wilmot Li, and Niloy J. Mitra. PATEX: Exploring Pattern Variations. *ACM Trans. Graph.*, 35(4):48:1–48:13, July 2016.
- 2011
- 2012
- 2013 [Gle01] Michael Gleicher. Motion Path Editing. In *Proceedings of the 2001 Symposium on Interactive 3D Graphics*, I3D ’01, pages 195–202, New York, NY, USA, 2001. Association for Computing Machinery.
- 2014
- 2015
- 2016 [Gom84] Ernst Hans Gombrich. *The Sense of Order: A Study in the Psychology of Decorative Art*. Phaidon Press Limited, 1984.
- 2017
- 2018 [GSCO07] Ran Gal, Ariel Shamir, and Daniel Cohen-Or. Pose-Oblivious Shape Signature. *IEEE Transactions on Visualization and Computer Graphics*, 13(2):261–271, March 2007.
- 2019
- 2020
- 2021 [GSP<sup>+</sup>07] Ran Gal, Olga Sorkine, Tiberiu Popa, Alla Sheffer, and Daniel Cohen-Or. 3D Collage: Expressive Non-Realistic Modeling. In *Proceedings of the 5th International Symposium on Non-Photorealistic Animation and Rendering*, NPAR ’07, pages 7–14, New York, NY, USA, 2007. ACM.
- 2022
- 2023
- 2024
- 2025 [Hau01] Alejo Hausner. Simulating Decorative Mosaics. In *Proceedings of the 28th Annual Conference on Computer Graphics and Interactive Techniques*, SIGGRAPH ’01, pages 573–580, New York, NY, USA, 2001. ACM.
- 2026
- 2027
- 2028 [HHD03] Stefan Hiller, Heino Hellwig, and Oliver Deussen. Beyond Stippling - Methods for Distributing Objects on the Plane. *Computer Graphics Forum*, 2003.
- 2029
- 2030 [HKT10] Edmond S. L. Ho, Taku Komura, and Chiew-Lan Tai. Spatial Relationship Preserving Character Motion Adaptation. *ACM Trans. Graph.*, 29(4), July 2010.
- 2031
- 2032
- 2033 [HLT<sup>+</sup>09] Thomas Hurtut, Pierre-Edouard Landes, Joëlle. Thollot, Yann Gousseau, Remy Drouillhet, and Jean-François Coeurjolly. Appearance-Guided Synthesis of Element Arrangements by Example. In *Proceedings of the 7th International Symposium on Non-Photorealistic Animation and Rendering*, NPAR ’09, pages 51–60, New York, NY, USA, 2009. ACM.
- 2034
- 2035
- 2036
- 2037
- 2038 [HLW93] S. C. Hsu, I. H. H. Lee, and N. E. Wiseman. Skeletal Strokes. In *Proceedings of the 6th Annual ACM Symposium on User Interface Software and Technology*, UIST ’93, pages 197–206, New York, NY, USA, 1993. ACM.
- 2039
- 2040

- 2041 [Hut29] Francis Hutcheson. *An Inquiry Into the Original of Our Ideas of Beauty and*  
 2042 *Virtue*. J. and J. Knapton and others, 1729.
- 2043 [HWFL14] Zhe Huang, Jiang Wang, Hongbo Fu, and Rynson W.H. Lau. Structured Me-  
 2044 chanical Collage. *IEEE Transactions on Visualization and Computer Graph-*  
 2045 *ics*, 20(7):1076–1082, 2014.
- 2046 [HWYZ20] Chen-Yuan Hsu, Li-Yi Wei, Lihua You, and Jian Jun Zhang. Autocomplete  
 2047 Element Fields. In *Proceedings of the 2020 CHI Conference on Human Factors*  
 2048 *in Computing Systems*, CHI 20, New York, NY, USA, 2020. Association for  
 2049 Computing Machinery.
- 2050 [HZZ11] Hua Huang, Lei Zhang, and Hong-Chao Zhang. Arcimboldo-Like Collage Us-  
 2051 ing Internet Images. In *Proceedings of the 2011 SIGGRAPH Asia Conference*,  
 2052 SA ’11, pages 155:1–155:8, New York, NY, USA, 2011. ACM.
- 2053 [IMH05] Takeo Igarashi, Tomer Moscovich, and John F. Hughes. As-Rigid-As-Possible  
 2054 Shape Manipulation. In *SIGGRAPH ’05*, SIGGRAPH ’05, pages 1134–1141,  
 2055 New York, NY, USA, 2005. ACM.
- 2056 [IMIM08] Takashi Ijiri, Radomír Mečh, Takeo Igarashi, and Gavin S. P. Miller. An  
 2057 Example-Based Procedural System for Element Arrangement. *Computer*  
 2058 *Graphics Forum*, 27:429–436, 2008.
- 2059 [Ise13] Tobias Isenberg. Evaluating and Validating Non-Photorealistic and Illustrative  
 2060 Rendering. In Paul Rosin and John Collomosse, editors, *Image and Video*  
 2061 *based Artistic Stylisation*, volume 42 of *Computational Imaging and Vision*,  
 2062 chapter 15, pages 311–331. Springer, London, Heidelberg, 2013.
- 2063 [JBPS11] Alec Jacobson, Ilya Baran, Jovan Popović, and Olga Sorkine. Bounded Bihar-  
 2064 monic Weights for Real-Time Deformation. *ACM Trans. Graph.*, 30(4):78:1–  
 2065 78:8, July 2011.
- 2066 [JDM19] Ali Sattari Javid, Lars Doyle, and David Mould. Irregular Pebble Mosaics with  
 2067 Sub-Pebble Detail. In *ACM/EG Expressive Symposium*. The Eurographics  
 2068 Association, 2019.
- 2069 [JL97] Bruno Jobard and Wilfrid Lefer. Creating Evenly-Spaced Streamlines of Ar-  
 2070 bitrary Density. In Wilfrid Lefer and Michel Grave, editors, *Visualization*  
 2071 *in Scientific Computing ’97: Proceedings of the Eurographics Workshop in*

- 2072                    *Boulogne-sur-Mer France, April 28–30, 1997*, pages 43–55, Vienna, 1997.  
 2073                    Springer Vienna.
- 2074     [Joh14]        Angus Johnson. Clipper—An Open Source Freeware Library for Clipping and  
 2075                    Offsetting Lines and Polygons. <http://www.angusj.com/delphi/clipper.php>, 2014.
- 2077     [Kap19]        Craig S. Kaplan. Animated Isohedral Tilings. In *Proceedings of Bridges 2019: Mathematics, Art, Music, Architecture, Education, Culture*, pages 99–106, Phoenix, Arizona, 2019. Tessellations Publishing. Available online at <http://archive.bridgesmathart.org/2019/bridges2019-99.pdf>.
- 2081     [KHHL12]      Manmyung Kim, Youngseok Hwang, Kyunglyul Hyun, and Jehee Lee. Tiling Motion Patches. In *Proceedings of the ACM SIGGRAPH/Eurographics Symposium on Computer Animation*, SCA 12, pages 117–126, Goslar, DEU, 2012. Eurographics Association.
- 2085     [KOHY11]      Dongwann Kang, Yong-Jin Ohn, Myoung-Hun Han, and Kyung-Hyun Yoon. Animation for Ancient Tile Mosaics. In *Proceedings of the ACM SIGGRAPH/Eurographics Symposium on Non-Photorealistic Animation and Rendering*, NPAR ’11, pages 157–166, New York, NY, USA, 2011. ACM.
- 2089     [KP02]        Junhwan Kim and Fabio Pellacini. Jigsaw Image Mosaics. In *Proceedings of the 29th Annual Conference on Computer Graphics and Interactive Techniques*, SIGGRAPH ’02, pages 657–664, New York, NY, USA, 2002. ACM.
- 2092     [KS00]        Craig S. Kaplan and David H. Salesin. Escherization. In *Proceedings of the 27th Annual Conference on Computer Graphics and Interactive Techniques*, SIGGRAPH ’00, pages 499–510, New York, NY, USA, 2000. ACM Press/Addison-Wesley Publishing Co.
- 2096     [KS04]        Craig S. Kaplan and David H. Salesin. Dihedral Escherization. In *Proceedings of Graphics Interface 2004*, GI ’04, pages 255–262, School of Computer Science, University of Waterloo, Waterloo, Ontario, Canada, 2004. Canadian Human-Computer Communications Society / Société canadienne du dialogue humain-machine.
- 2101     [KSH<sup>+</sup>16]    Kin Chung Kwan, Lok Tsun Sinn, Chu Han, Tien-Tsin Wong, and Chi-Wing Fu. Pyramid of Arclength Descriptor for Generating Collage of Shapes. *ACM Trans. Graph.*, 35(6):229:1–229:12, November 2016.

- 2104 [LBMH19] Yifei Li, David E. Breen, James McCann, and Jessica Hodgins. Algorithmic  
 2105 Quilting Pattern Generation for Pieced Quilts. In *Proceedings of Graphics*  
 2106 *Interface 2019*, GI 2019. Canadian Information Processing Society, 2019.
- 2107 [LBW<sup>+</sup>14] Jingwan Lu, Connelly Barnes, Connie Wan, Paul Asente, Radomír Mečh, and  
 2108 Adam Finkelstein. DecoBrush: Drawing Structured Decorative Patterns by  
 2109 Example. *ACM Trans. Graph.*, 33(4):90:1–90:9, July 2014.
- 2110 [LBZ<sup>+</sup>11] Yuanyuan Li, Fan Bao, Eugene Zhang, Yoshihiro Kobayashi, and Peter  
 2111 Wonka. Geometry Synthesis on Surfaces Using Field-Guided Shape  
 2112 Grammars. *IEEE Transactions on Visualization and Computer Graphics*,  
 2113 17(2):231243, February 2011.
- 2114 [LHVT17] Hugo Loi, Thomas Hurtut, Romain Vergne, and Joelle Thollot. Programmable  
 2115 2D Arrangements for Element Texture Design. *ACM Trans. Graph.*, 36(4),  
 2116 May 2017.
- 2117 [LJG14] Songrun Liu, Alec Jacobson, and Yotam Gingold. Skinning Cubic Bézier  
 2118 Splines and Catmull-Clark Subdivision Surfaces. *ACM Trans. Graph.*,  
 2119 33(6):190:1–190:9, November 2014.
- 2120 [LML<sup>+</sup>18] Shih-Syun Lin, Charles C. Morace, Chao-Hung Lin, Li-Fong Hsu, and Tong-  
 2121 Yee Lee. Generation of Escher Arts with Dual Perception. *IEEE Transactions*  
 2122 *on Visualization and Computer Graphics*, 24(2):1103–1113, 2018.
- 2123 [LV09] Yu Liu and Olga Veksler. Animated Classic Mosaics from Video. In *Proceed-  
 2124 ings of the 5th International Symposium on Advances in Visual Computing: Part II*, ISVC ’09, pages 1085–1096, Berlin, Heidelberg, 2009. Springer-Verlag.
- 2126 [MBS<sup>+</sup>11] Ron Maharik, Mikhail Bessmeltsev, Alla Sheffer, Ariel Shamir, and Nathan  
 2127 Carr. Digital Micrography. In *ACM SIGGRAPH 2011 Papers*, SIGGRAPH  
 2128 ’11, pages 100:1–100:12, New York, NY, USA, 2011. ACM.
- 2129 [MCHW18] Y. Ma, Z. Chen, W. Hu, and W. Wang. Packing Irregular Objects in 3D Space  
 2130 via Hybrid Optimization. *Computer Graphics Forum*, 37(5):49–59, 2018.
- 2131 [MCKM15] Matthias Müller, Nuttapong Chentanez, Tae-Yong Kim, and Miles Macklin.  
 2132 Air Meshes for Robust Collision Handling. *ACM Trans. Graph.*, 34(4):133:1–  
 2133 133:9, July 2015.

- 2134 [MHHR07] Matthias Müller, Bruno Heidelberger, Marcus Hennix, and John Ratcliff. Position Based Dynamics. *J. Vis. Comun. Image Represent.*, 18(2):109–118, April 2007.
- 2135
- 2136
- 2137 [MIA<sup>+</sup>08] Ross Maciejewski, Tobias Isenberg, William M. Andrews, David S. Ebert, Mario Costa Sousa, and Wei Chen. Measuring Stipple Aesthetics in Hand-Drawn and Computer-Generated Images. *IEEE Computer Graphics and Applications*, 28(2):62–74, March/April 2008.
- 2138
- 2139
- 2140
- 2141 [MSS<sup>+</sup>19] Jonàs Martínez, Mélina Skouras, Christian Schumacher, Samuel Hornus, Sylvain Lefebvre, and Bernhard Thomaszewski. Star-Shaped Metrics for Mechanical Metamaterial Design. *ACM Trans. Graph.*, 38(4), July 2019.
- 2142
- 2143
- 2144 [MWLT13] Chongyang Ma, Li-Yi Wei, Sylvain Lefebvre, and Xin Tong. Dynamic Element Textures. *ACM Trans. Graph.*, 32(4), July 2013.
- 2145
- 2146 [MWT11] Chongyang Ma, Li-Yi Wei, and Xin Tong. Discrete Element Textures. In *SIGGRAPH ’11*, SIGGRAPH ’11, pages 62:1–62:10, New York, NY, USA, 2011. ACM.
- 2147
- 2148
- 2149 [OS88] Stanley Osher and James A. Sethian. Fronts Propagating with Curvature-Dependent Speed: Algorithms Based on Hamilton-Jacobi Formulations. *J. Comput. Phys.*, 79(1):12–49, November 1988.
- 2150
- 2151
- 2152 [Osh17] Masaki Oshita. Lattice-Guided Human Motion Deformation for Collision Avoidance. In *Proceedings of the Tenth International Conference on Motion in Games*, MIG ’17, New York, NY, USA, 2017. Association for Computing Machinery.
- 2153
- 2154
- 2155
- 2156 [PS06] Hans Pedersen and Karan Singh. Organic Labyrinths and Mazes. In *Proceedings of the 4th International Symposium on Non-Photorealistic Animation and Rendering*, NPAR ’06, pages 79–86, New York, NY, USA, 2006. ACM.
- 2157
- 2158
- 2159 [PYW14] Chi-Han Peng, Yong-Liang Yang, and Peter Wonka. Computing Layouts with Deformable Templates. *ACM Trans. Graph.*, 33(4):99:1–99:11, July 2014.
- 2160
- 2161 [PZ07] Jonathan Palacios and Eugene Zhang. Rotational Symmetry Field Design on Surfaces. In *ACM SIGGRAPH 2007 Papers*, SIGGRAPH ’07, New York, NY, USA, 2007. ACM.
- 2162
- 2163

- 2164 [Sch90] Philip J. Schneider. An Algorithm for Automatically Fitting Digitized Curves.  
 2165 In Andrew S. Glassner, editor, *Graphics Gems*, pages 612–626. Academic Press  
 2166 Professional, Inc., San Diego, CA, USA, 1990.
- 2167 [Sec02] Adrian Secord. Weighted Voronoi Stippling. In *Proceedings of the 2nd International Symposium on Non-Photorealistic Animation and Rendering*, NPAR 02,  
 2168 page 3743, New York, NY, USA, 2002. Association for Computing Machinery.
- 2170 [SKA18] Reza Adhitya Saputra, Craig S. Kaplan, and Paul Asente. RepulsionPak:  
 2171 Deformation-Driven Element Packing with Repulsion Forces. In *Proceedings  
 2172 of the 44th Graphics Interface Conference*, GI '18. Canadian Human-Computer  
 2173 Communications Society / Société canadienne du dialogue humain-machine,  
 2174 2018.
- 2175 [SKA19] Reza Adhitya Saputra, Craig S. Kaplan, and Paul Asente. Improved  
 2176 Deformation-Driven Element Packing with RepulsionPak. *IEEE Transactions  
 2177 on Visualization and Computer Graphics*, pages 1–1, 2019.
- 2178 [SKA20] Reza Adhitya Saputra, Craig S. Kaplan, and Paul Asente. AnimationPak:  
 2179 Packing Elements with Scripted Animations. In *Proceedings of the 46th Graphics  
 2180 Interface Conference*, GI '20. Canadian Human-Computer Communications  
 2181 Society / Société canadienne du dialogue humain-machine, 2020.
- 2182 [SKAM17] Reza Adhitya Saputra, Craig S. Kaplan, Paul Asente, and Radomír Měch.  
 2183 FLOWPAK: Flow-Based Ornamental Element Packing. In *Proceedings of the  
 2184 43rd Graphics Interface Conference*, GI '17, pages 8–15. Canadian Human-  
 2185 Computer Communications Society / Société canadienne du dialogue humain-  
 2186 machine, 2017.
- 2187 [SLK05] Kaleigh Smith, Yunjun Liu, and Allison Klein. Animosaics. In *Proceedings of  
 2188 the 2005 ACM SIGGRAPH/Eurographics Symposium on Computer Animation*, SCA '05,  
 2189 pages 201–208, New York, NY, USA, 2005. ACM.
- 2190 [TRdAS07] Christian Theobalt, Christian Rössl, Edilson de Aguiar, and Hans-Peter  
 2191 Seidel. Animation Collage. In *Proceedings of the 2007 ACM SIG-  
 2192 GRAPH/Eurographics Symposium on Computer Animation*, SCA 07, page  
 2193 271280, Goslar, DEU, 2007. Eurographics Association.
- 2194 [VGB<sup>+</sup>14] J. Vanek, J. A. Garcia Galicia, B. Benes, R. Měch, N. Carr, O. Št'ava, and  
 2195 G. S. Miller. PackMerger: A 3D Print Volume Optimizer. *Computer Graphics  
 2196 Forum*, 33(6):322332, September 2014.

- 2197 [WK88] Andrew Witkin and Michael Kass. Spacetime Constraints. In *Proceedings*  
 2198 *of the 15th Annual Conference on Computer Graphics and Interactive Techniques*, SIGGRAPH 88, pages 159–168, New York, NY, USA, 1988. Association  
 2199 for Computing Machinery.
- 2201 [WP95] Andrew Witkin and Zoran Popović. Motion Warping. In *Proceedings of the*  
 2202 *22nd Annual Conference on Computer Graphics and Interactive Techniques*, SIGGRAPH 95, pages 105–108, New York, NY, USA, 1995. Association for  
 2203 Computing Machinery.
- 2205 [WZS98] Michael T. Wong, Douglas E. Zongker, and David H. Salesin. Computer-  
 2206 Generated Floral Ornament. In *Proceedings of the 25th Annual Conference*  
 2207 *on Computer Graphics and Interactive Techniques*, SIGGRAPH ’98, pages  
 2208 423–434, New York, NY, USA, 1998. ACM.
- 2209 [XK07] Jie Xu and Craig S. Kaplan. Calligraphic Packing. In *Proceedings of Graphics*  
 2210 *Interface 2007*, GI ’07, pages 43–50, New York, NY, USA, 2007. ACM.
- 2211 [XM09] Ling Xu and David Mould. Magnetic Curves: Curvature-Controlled Aesthetic  
 2212 Curves Using Magnetic Fields. In *Proceedings of the Fifth Eurographics Con-*  
 2213 *ference on Computational Aesthetics in Graphics, Visualization and Imaging*,  
 2214 Computational Aesthetics’09, pages 1–8, Aire-la-Ville, Switzerland, Switzer-  
 2215 land, 2009. Eurographics Association.
- 2216 [XM15] Ling Xu and David Mould. Procedural Tree Modeling with Guiding Vectors.  
 2217 *Computer Graphics Forum*, 34(7):47–56, October 2015.
- 2218 [YSC20] Christopher Yu, Henrik Schumacher, and Keenan Crane. Repulsive Curves,  
 2219 2020.
- 2220 [ZCR<sup>+</sup>16] Changqing Zou, Junjie Cao, Warunika Ranaweera, Ibraheem Alhashim, Ping  
 2221 Tan, Alla Sheffer, and Hao Zhang. Legible Compact Calligrams. *ACM Trans.*  
 2222 *Graph.*, 35(4):122:1–122:12, July 2016.
- 2223 [ZCT16] Jonas Zehnder, Stelian Coros, and Bernhard Thomaszewski. Designing  
 2224 Structurally-Sound Ornamental Curve Networks. *ACM Trans. Graph.*,  
 2225 35(4):99:1–99:10, July 2016.
- 2226 [ZS84] TY Zhang and Ching Y. Suen. A Fast Parallel Algorithm for Thinning Digital  
 2227 Patterns. *Communications of the ACM*, 27(3):236–239, 1984.